



OpenStack

# Virtual Machine Image Guide

havana (April 4, 2014)

BUILT FOR



openstack™  
CLOUD SOFTWARE



## OpenStack Virtual Machine Image Guide

havana (2014-04-04)

Copyright © 2013 OpenStack Foundation Some rights reserved.

This guide describes how to obtain, create, and modify virtual machine images that are compatible with OpenStack.



Except where otherwise noted, this document is licensed under  
**Creative Commons Attribution 3.0 License.**  
<http://creativecommons.org/licenses/by/3.0/legalcode>

# Table of Contents

Preface .....	6
Conventions .....	6
Document change history .....	6
1. Introduction .....	1
Disk and container formats for images .....	3
Image metadata .....	4
2. Get images .....	8
CirrOS (test) images .....	8
Official Ubuntu images .....	8
Official Red Hat Enterprise Linux images .....	8
Official Fedora images .....	9
Official openSUSE and SLES images .....	9
Official images from other Linux distributions .....	9
Rackspace Cloud Builders (multiple distros) images .....	9
Microsoft Windows images .....	9
3. OpenStack Linux image requirements .....	10
Disk partitions and resize root partition on boot (cloud-init) .....	10
No hard-coded MAC address information .....	12
Ensure ssh server runs .....	13
Disable firewall .....	13
Access instance by using ssh public key (cloud-init) .....	13
Process user data and other metadata (cloud-init) .....	14
Ensure image writes boot log to console .....	14
Paravirtualized Xen support in the kernel (Xen hypervisor only) .....	15
Manage the image cache .....	15
4. Modify images .....	17
guestfish .....	17
guestmount .....	19
virt-* tools .....	19
Loop devices, kpartx, network block devices .....	20
5. Create images manually .....	24
Verify the libvirt default network is running .....	24
Use the virt-manager X11 GUI .....	24
Use virt-install and connect by using a local VNC client .....	26
Example: CentOS image .....	27
Example: Ubuntu image .....	34
Example: Microsoft Windows image .....	40
Example: FreeBSD image .....	41
6. Tool support for image creation .....	42
Oz .....	42
VMBuilder .....	43
BoxGrinder .....	44
VeeWee .....	44
imagefactory .....	44
SUSE Studio .....	44
7. Converting between image formats .....	45
A. Community support .....	46
Documentation .....	46

---

ask.openstack.org .....	47
OpenStack mailing lists .....	47
The OpenStack wiki .....	47
The Launchpad Bugs area .....	48
The OpenStack IRC channel .....	49
Documentation feedback .....	49
OpenStack distribution packages .....	49

## List of Tables

3.1. Image cache management configuration options .....	15
7.1. qemu-img format strings .....	45

# Preface

Conventions .....	6
Document change history .....	6

## Conventions

The OpenStack documentation uses several typesetting conventions:

## Admonitions

Admonitions take three forms:



### Note

This is a note. The information in a note is usually in the form of a handy tip or reminder.



### Important

This is important. The information in an important admonition is something you must be aware of before moving on.



### Warning

This is a warning. The information in warnings is critical. Warnings provide additional information about risk of data loss or security issues.

## Command prompts

Commands prefixed with the `#` prompt are to be executed by the `root` user. These examples can also be executed using the `sudo` command, if available.

Commands prefixed with the `$` prompt can be executed by any user, including `root`.

## Document change history

This version of the guide replaces and obsoletes all previous versions. The following table describes the most recent changes:

Revision Date	Summary of Changes
October 25, 2013	<ul style="list-style-type: none"><li>• Adds information about image formats, properties.</li></ul>
October 17, 2013	<ul style="list-style-type: none"><li>• Havana release.</li></ul>
June 4, 2013	<ul style="list-style-type: none"><li>• Updated title for consistency.</li></ul>
May 28, 2013	<ul style="list-style-type: none"><li>• Initial release of this guide.</li></ul>

# 1. Introduction

Disk and container formats for images .....	3
Image metadata .....	4

An OpenStack Compute cloud is not very useful unless you have virtual machine images (which some people call "virtual appliances"). This guide describes how to obtain, create, and modify virtual machine images that are compatible with OpenStack.

To keep things brief, we'll sometimes use the term "image" instead of "virtual machine image".

What is a virtual machine image?

A virtual machine image is a single file which contains a virtual disk that has a bootable operating system installed on it.

Virtual machine images come in different formats, some of which are described below. In a later chapter, we'll describe how to convert between formats.

**Raw** The "raw" image format is the simplest one, and is natively supported by both KVM and Xen hypervisors. You can think of a raw image as being the bit-equivalent of a block device file, created as if somebody had copied, say, `/dev/sda` to a file using the **dd** command.



## Note

We don't recommend creating raw images by dd'ing block device files, we discuss how to create raw images later.

**qcow2** The [qcow2](#) (QEMU copy-on-write version 2) format is commonly used with the KVM hypervisor. It has some additional features over the raw format, such as:

- Using sparse representation, so the image size is smaller
- Support for snapshots

Because qcow2 is sparse, it's often faster to convert a raw image to qcow2 and upload it then to upload the raw file.



## Note

Because raw images don't support snapshots, OpenStack Compute will automatically convert raw image files to qcow2 as needed.

**AMI/AKI/ARI** The [AMI/AKI/ARI](#) format was the initial image format supported by Amazon EC2. The image consists of three files:

- AMI (Amazon Machine Image):

This is a virtual machine image in raw format, as described above.

- AKI (Amazon Kernel Image)

A kernel file that the hypervisor will load initially to boot the image. For a Linux machine, this would be a *vmlinux* file.

- ARI (Amazon Ramdisk Image)

An optional ramdisk file mounted at boot time. For a Linux machine, this would be an *initrd* file.

UEC tarball

A UEC (Ubuntu Enterprise Cloud) tarball is a gzipped tarfile that contains an AMI file, AKI file, and ARI file.



### Note

Ubuntu Enterprise Cloud refers to a discontinued Eucalyptus-based Ubuntu cloud solution that has been replaced by the OpenStack-based Ubuntu Cloud Infrastructure.

VMDK

VMWare's ESXi hypervisor uses the [VMDK](#) (Virtual Machine Disk) format for images.

VDI

VirtualBox uses the [VDI](#) (Virtual Disk Image) format for image files. None of the OpenStack Compute hypervisors support VDI directly, so you will need to convert these files to a different format to use them with OpenStack.

VHD

Microsoft Hyper-V uses the VHD (Virtual Hard Disk) format for images.

VHDX

The version of Hyper-V that ships with Microsoft Server 2012 uses the newer [VHDX](#) format, which has some additional features over VHD such as support for larger disk sizes and protection against data corruption during power failures.

OVF

[OVF](#) (Open Virtualization Format) is a packaging format for virtual machines, defined by the Distributed Management Task Force (DMTF) standards group. An OVF package contains one or more image files, a .ovf XML metadata file that contains information about the virtual machine, and possibly other files as well.

An OVF package can be distributed in different ways. For example, it could be distributed as a set of discrete files, or as a tar archive file with an .ova (open virtual appliance/application) extension.

OpenStack Compute does not currently have support for OVF packages, so you will need to extract the image file(s) from an OVF package if you wish to use it with OpenStack.

ISO

The [ISO](#) format is a disk image formatted with the read-only ISO 9660 (also known as ECMA-119) filesystem commonly used for CDs and DVDs.



While we don't normally think of ISO as a virtual machine image format, since ISOs contain bootable filesystems with an installed operating system, you can treat them the same as you treat other virtual machine image files.

## Disk and container formats for images

When you add an image to the Image Service, you can specify its disk and container formats.

### Disk formats

The disk format of a virtual machine image is the format of the underlying disk image. Virtual appliance vendors have different formats for laying out the information contained in a virtual machine disk image.

Set the disk format for your image to one of the following values:

- `raw`. An unstructured disk image format; if you have a file without an extension it is possibly a raw format
- `vhd`. The VHD disk format, a common disk format used by virtual machine monitors from VMWare, Xen, Microsoft, VirtualBox, and others
- `vmdk`. Common disk format supported by many common virtual machine monitors
- `vdi`. Supported by VirtualBox virtual machine monitor and the QEMU emulator
- `iso`. An archive format for the data contents of an optical disc, such as CD-ROM.
- `qcow2`. Supported by the QEMU emulator that can expand dynamically and supports Copy on Write
- `aki`. An Amazon kernel image.
- `ari`. An Amazon ramdisk image.
- `ami`. An Amazon machine image.

### Container formats

The container format indicates whether the virtual machine image is in a file format that also contains metadata about the actual virtual machine.



#### Note

The Image Service and other OpenStack projects do not currently support the container format. It is safe to specify `bare` as the container format if you are unsure.

You can set the container format for your image to one of the following values:

- `bare`. The image does not have a container or metadata envelope.
- `ovf`. The OVF container format.
- `aki`. An Amazon kernel image.
- `ari`. An Amazon ramdisk image.
- `ami`. An Amazon machine image.

## Image metadata

You can associate metadata with an image using the `--property key=value` argument to `glance image-create` or `glance image-update`. For example:

```
$ glance image-update img-uuid --property architecture=arm --property
hypervisor_type=qemu
```

If the following properties are set on an image, and the ImagePropertiesFilter scheduler filter is enabled (which it is by default), then the scheduler will only consider compute hosts that satisfy these properties:

architecture	The CPU architecture that must be supported by the hypervisor. For example, <code>x86_64</code> , <code>arm</code> , or <code>ppc64</code> . Run <code>uname -m</code> to get the architecture of a machine. We strongly recommend using the architecture data vocabulary defined by the <a href="#">libosinfo project</a> for this purpose. Recognized values for this field are:
alpha	DEC 64-bit RISC
armv7l	ARM Cortex-A7 MPCore
cris	Ethernet, Token Ring, AXis - Code Reduced Instruction Set
i686	Intel sixth-generation x86 (P6 micro architecture)
ia64	Itanium
lm32	Lattice Micro32
m68k	Motorola 68000
microblaze	Xilinx 32-bit FPGA (Big Endian)
microblazeel	Xilinx 32-bit FPGA (Little Endian)
mips	MIPS 32-bit RISC (Big Endian)
mipsel	MIPS 32-bit RISC (Little Endian)
mips64	MIPS 64-bit RISC (Big Endian)
mips64el	MIPS 64-bit RISC (Little Endian)
openrisc	OpenCores RISC
parisc	HP Precision Architecture RISC
parisc64	HP Precision Architecture 64-bit RISC
ppc	PowerPC 32-bit
ppc64	PowerPC 64-bit
ppcemb	PowerPC (Embedded 32-bit)
s390	IBM Enterprise Systems Architecture/390
s390x	S/390 64-bit
sh4	SuperH SH-4 (Little Endian)
sh4eb	SuperH SH-4 (Big Endian)
sparc	Scalable Processor Architecture, 32-bit

	sparc64 <a href="#">Scalable Processor Architecture, 64-bit</a> unicore32 <a href="#">Microprocessor Research and Development Center RISC Unicores</a> x86_64 <a href="#">64-bit extension of IA-32</a> xtensa <a href="#">Tensilica Xtensa configurable microprocessor core</a> xtensaeb <a href="#">Tensilica Xtensa configurable microprocessor core (Big Endian)</a>
hypervisor_type	The hypervisor type. Allowed values include: xen, qemu, kvm, lxc, uml, vmware, hyperv.
vm_mode	The virtual machine mode. This represents the host/guest ABI (application binary interface) used for the virtual machine. Allowed values are:  hvm    Fully virtualized. This is the mode used by QEMU and KVM.  xen    Xen 3.0 paravirtualized.  uml    User Mode Linux paravirtualized.  exe    Executables in containers. This is the mode used by LXC.

The following metadata properties are specific to the XenAPI driver:

auto_disk_config	A boolean option. If true, the root partition on the disk will be automatically resized before the instance boots. This value is only taken into account by the Compute service when using a Xen-based hypervisor with the XenAPI driver. The Compute service will only attempt to resize if there is a single partition on the image, and only if the partition is in ext3 or ext4 format.
os_type	The operating system installed on the image. For example, <code>linux</code> or <code>windows</code> . The XenAPI driver contains logic that will take different actions depending on the value of the <code>os_type</code> parameter of the image. For example, for images where <code>os_type=windows</code> , it creates a FAT32-based swap partition instead of a Linux swap partition, and it limits the injected host name to less than 16 characters.

The following metadata properties are specific to the libvirt API driver:

hw_video_model	The video image driver used. Allowed values include: vga, cirrus, vmvga, xen and qxl.
----------------	---

The following metadata properties are specific to the VMware API driver:

vmware_adaptype	Indicates the virtual SCSI or IDE controller used by the hypervisor. Allowed values: <code>lsiLogic</code> , <code>busLogic</code> , <code>ide</code>
vmware_ostype	A VMware GuestID which describes the operating system installed in the image. This will be passed to the hypervisor when creating a virtual machine. See <a href="http://thinkvirt.com">thinkvirt.com</a> for a

list of valid values. If this is not specified, it will default to `otherGuest`.

`vmware_image_version`      Currently unused, set it to 1.

To assist end-users in utilizing images, you may wish to put additional common metadata on Glance images. By community agreement, the following metadata keys may be used across Glance installations for the purposes described, as follows.

<code>instance_uuid</code>	For snapshot images, this is the UUID of the server used to create this image.
<code>kernel_id</code>	The ID of image stored in Glance that should be used as the kernel when booting an AMI-style image.
<code>ramdisk_id</code>	The ID of image stored in Glance that should be used as the ramdisk when booting an AMI-style image.
<code>os_version</code>	The operating system version as specified by the distributor.
<code>os_distro</code>	The value of this property is the common name of the operating system distribution in all-lowercase. For this purpose, we use the same data vocabulary as the <a href="#">libosinfo project</a> . Following are the recognized values for this property. In the interest of interoperability, please use only a recognized value for this field. The deprecated values are listed to assist you in searching for the recognized value. Allowed values are:
<code>arch</code>	This is: Arch Linux  Do not use: <code>archlinux</code> , or <code>org.archlinux</code>
<code>centos</code>	This is: Community Enterprise Operating System  Do not use: <code>org.centos</code> , or <code>CentOS</code>
<code>debian</code>	This is: Debian  Do not use: <code>Debian</code> , or <code>org.debian</code>
<code>fedora</code>	This is: Fedora  Do not use: <code>Fedora</code> , <code>org.fedora</code> , or <code>org.fedoraproject</code>
<code>freebsd</code>	This is: FreeBSD  Do not use: <code>org.freebsd</code> , <code>freeBSD</code> , or <code>FreeBSD</code>
<code>gentoo</code>	This is: Gentoo Linux  Do not use: <code>Gentoo</code> , or <code>org.gentoo</code>
<code>mandrake</code>	This is: Mandrakelinux (MandrakeSoft)  Do not use: <code>mandrakelinux</code> , or <code>MandrakeLinux</code>

mandriva	<p>This is: Mandriva Linux</p> <p>Do not use: mandrivalinux</p>
mes	<p>This is: Mandriva Enterprise Server</p> <p>Do not use: mandrivaent, or mandrivaES</p>
msdos	<p>This is: Microsoft Disc Operating System</p> <p>Do not use: ms-dos</p>
netbsd	<p>This is: NetBSD</p> <p>Do not use: NetBSD, or org.netbsd</p>
netware	<p>This is: Novell NetWare</p> <p>Do not use: novell, or NetWare</p>
openbsd	<p>This is: OpenBSD</p> <p>Do not use: OpenBSD, or org.openbsd</p>
opensolaris	<p>Do not use: OpenSolaris, or org.opensolaris</p>
opensuse	<p>This is: openSUSE</p> <p>Do not use: suse, SuSE, or org.opensuse</p>
rhel	<p>This is: Red Hat Enterprise Linux</p> <p>Do not use: redhat, RedHat, or com.redhat</p>
sled	<p>This is: SUSE Linux Enterprise Desktop</p> <p>Do not use: com.suse</p>
ubuntu	<p>This is: Ubuntu</p> <p>Do not use: Ubuntu, com.ubuntu, org.ubuntu, or canonical</p>
windows	<p>This is: Microsoft Windows</p> <p>Do not use: com.microsoft.server, or windoze</p>

## 2. Get images

Cirros (test) images .....	8
Official Ubuntu images .....	8
Official Red Hat Enterprise Linux images .....	8
Official Fedora images .....	9
Official openSUSE and SLES images .....	9
Official images from other Linux distributions .....	9
Rackspace Cloud Builders (multiple distros) images .....	9
Microsoft Windows images .....	9

The simplest way to obtain a virtual machine image that works with OpenStack is to download one that someone else has already created.

### Cirros (test) images

Cirros is a minimal Linux distribution that was designed for use as a test image on clouds such as OpenStack Compute. You can download a Cirros image in various formats from the [Cirros Launchpad download page](#).

If your deployment uses QEMU or KVM, we recommend using the images in qcow2 format. The most recent 64-bit qcow2 image as of this writing is [cirros-0.3.1-x86\\_64-disk.img](#)



#### Note

In a Cirros image, the login account is `cirros`. The password is `cubswin:`

### Official Ubuntu images

Canonical maintains an [official set of Ubuntu-based images](#).

Images are arranged by Ubuntu release, and by image release date, with "current" being the most recent. For example, the page that contains the most recently built image for Ubuntu 12.04 "Precise Pangolin" is <http://cloud-images.ubuntu.com/precise/current/>. Scroll to the bottom of the page for links to images that can be downloaded directly.

If your deployment uses QEMU or KVM, we recommend using the images in qcow2 format. The most recent version of the 64-bit QCOW2 image for Ubuntu 12.04 is [precise-server-cloudimg-amd64-disk1.img](#).



#### Note

In an Ubuntu cloud image, the login account is `ubuntu`.

### Official Red Hat Enterprise Linux images

Red Hat maintains an official Red Hat Enterprise Linux 6 cloud image. A valid Red Hat Enterprise Linux 6 subscription is required to download the [KVM Guest Image](#).

A [boot.iso](#) image is also available. This is the same as the `boot.iso` file provided on the Red Hat Enterprise Linux installation DVD. It is a minimal image suitable for kickstarting operating system installation over a network.

## Official Fedora images

The Fedora project maintains a list of official cloud images at <http://cloud.fedoraproject.org/>. The images include the `cloud-init` utility to support key and user data injection. The default user name is `fedora`.

## Official openSUSE and SLES images

SUSE does not provide openSUSE or SUSE Linux Enterprise Server (SLES) images for direct download. Instead, they provide a web-based tool called [SUSE Studio](#) that you can use to build openSUSE and SLES images.

For example, Christian Berendt used [openSUSE](#) to create [a test openSUSE 12.3 image](#).

## Official images from other Linux distributions

As of this writing, we are not aware of other distributions that provide images for download.

## Rackspace Cloud Builders (multiple distros) images

Rackspace Cloud Builders maintains a list of pre-built images from various distributions (Red Hat, CentOS, Fedora, Ubuntu). Links to these images can be found at [rackerjoe/oz-image-build on Github](#).

## Microsoft Windows images

Cloudbase Solutions hosts an [OpenStack Windows Server 2012 Standard Evaluation image](#) that runs on Hyper-V, KVM, and XenServer/XCP.

## 3. OpenStack Linux image requirements

Disk partitions and resize root partition on boot (cloud-init) .....	10
No hard-coded MAC address information .....	12
Ensure ssh server runs .....	13
Disable firewall .....	13
Access instance by using ssh public key (cloud-init) .....	13
Process user data and other metadata (cloud-init) .....	14
Ensure image writes boot log to console .....	14
Paravirtualized Xen support in the kernel (Xen hypervisor only) .....	15
Manage the image cache .....	15

For a Linux-based image to have full functionality in an OpenStack Compute cloud, there are a few requirements. For some of these, you can fulfill the requirement by installing the [cloud-init](#) package. Read this section before you create your own image to be sure that the image supports the OpenStack features that you plan to use.

- Disk partitions and resize root partition on boot (cloud-init)
- No hard-coded MAC address information
- SSH server running
- Disable firewall
- Access instance using ssh public key (cloud-init)
- Process user data and other metadata (cloud-init)
- Paravirtualized Xen support in Linux kernel (Xen hypervisor only with Linux kernel version < 3.0)

### Disk partitions and resize root partition on boot (cloud-init)

When you create a Linux image, you must decide how to partition the disks. The choice of partition method can affect the resizing functionality, as described in the following sections.

The size of the disk in a virtual machine image is determined when you initially create the image. However, OpenStack lets you launch instances with different size drives by specifying different flavors. For example, if your image was created with a 5 GB disk, and you launch an instance with a flavor of `m1.small`. The resulting virtual machine instance has, by default, a primary disk size of 10 GB. When the disk for an instance is resized up, zeros are just added to the end.

Your image must be able to resize its partitions on boot to match the size requested by the user. Otherwise, after the instance boots, you must manually resize the partitions to access the additional storage to which you have access when the disk size associated with the flavor exceeds the disk size with which your image was created.



## Xen: 1 ext3/ext4 partition (no LVM, no /boot, no swap)

If you use the OpenStack XenAPI driver, the Compute service automatically adjusts the partition and file system for your instance on boot. Automatic resize occurs if the following conditions are all true:

- `auto_disk_config=True` is set as a property on the image in the image registry.
- The disk on the image has only one partition.
- The file system on the one partition is ext3 or ext4.

Therefore, if you use Xen, we recommend that when you create your images, you create a single ext3 or ext4 partition (not managed by LVM). Otherwise, read on.

## Non-Xen with cloud-init/cloud-tools: One ext3/ext4 partition (no LVM, no /boot, no swap)

You must configure these items for your image:

- The partition table for the image describes the original size of the image
- The file system for the image fills the original size of the image

Then, during the boot process, you must:

- Modify the partition table to make it aware of the additional space:
  - If you do not use LVM, you must modify the table to extend the existing root partition to encompass this additional space.
  - If you use LVM, you can add a new LVM entry to the partition table, create a new LVM physical volume, add it to the volume group, and extend the logical partition with the root volume.
- Resize the root volume file system.

The simplest way to support this in your image is to install the [cloud-utils](#) package (contains the **growpart** tool for extending partitions), the [cloud-initramfs-tools](#) package (which supports resizing root partition on the first boot), and the [cloud-init](#) package into your image. With these installed, the image performs the root partition resize on boot. For example, in the `/etc/rc.local` file. These packages are in the Ubuntu and Debian package repository, as well as the EPEL repository (for Fedora/RHEL/CentOS/Scientific Linux guests).

If you cannot install `cloud-initramfs-tools`, Robert Plestenjak has a [github project](#) called [linux-rootfs-resize](#) that contains scripts that update a ramdisk by using **growpart** so that the image resizes properly on boot.

If you can install the `cloud-utils` and `cloud-init` packages, we recommend that when you create your images, you create a single ext3 or ext4 partition (not managed by LVM).

## Non-Xen without cloud-init/cloud-tools: LVM

If you cannot install cloud-init and cloud-tools inside of your guest, and you want to support resize, you must write a script that your image runs on boot to modify the partition table. In this case, we recommend using LVM to manage your partitions. Due to a limitation in the Linux kernel (as of this writing), you cannot modify a partition table of a raw disk that has partitions currently mounted, but you can do this for LVM.

Your script must do something like the following:

1. Detect if any additional space is available on the disk. For example, parse the output of `parted /dev/sda --script "print free"`.
2. Create a new LVM partition with the additional space. For example, `parted /dev/sda --script "mkpart lvm ..."`.
3. Create a new physical volume. For example, `pvccreate /dev/sda6`.
4. Extend the volume group with this physical partition. For example, `vgextend vg00 /dev/sda6`.
5. Extend the logical volume contained the root partition by the amount of space. For example, `lvextend /dev/mapper/node-root /dev/sda6`.
6. Resize the root file system. For example, `resize2fs /dev/mapper/node-root`.

You do not need a `/boot` partition unless your image is an older Linux distribution that requires that `/boot` is not managed by LVM.

## No hard-coded MAC address information

You must remove the network persistence rules in the image because they cause the network interface in the instance to come up as an interface other than `eth0`. This is because your image has a record of the MAC address of the network interface card when it was first installed, and this MAC address is different each time that the instance boots. You should alter the following files:

- Replace `/etc/udev/rules.d/70-persistent-net.rules` with an empty file (contains network persistence rules, including MAC address)
- Replace `/lib/udev/rules.d/75-persistent-net-generator.rules` with an empty file (this generates the file above)
- Remove the `HWADDR` line from `/etc/sysconfig/network-scripts/ifcfg-eth0` on Fedora-based images



### Note

If you delete the network persistent rules files, you may get a udev kernel warning at boot time, which is why we recommend replacing them with empty files instead.

## Ensure ssh server runs

You must install an ssh server into the image and ensure that it starts up on boot, or you cannot connect to your instance by using ssh when it boots inside of OpenStack. This package is typically called `openssh-server`.

## Disable firewall

In general, we recommend that you disable any firewalls inside of your image and use OpenStack security groups to restrict access to instances. The reason is that having a firewall installed on your instance can make it more difficult to troubleshoot networking issues if you cannot connect to your instance.

## Access instance by using ssh public key (cloud-init)

The typical way that users access virtual machines running on OpenStack is to ssh using public key authentication. For this to work, your virtual machine image must be configured to download the ssh public key from the OpenStack metadata service or config drive, at boot time.

## Use cloud-init to fetch the public key

The cloud-init package automatically fetches the public key from the metadata server and places the key in an account. The account varies by distribution. On Ubuntu-based virtual machines, the account is called `ubuntu`. On Fedora-based virtual machines, the account is called `ec2-user`.

You can change the name of the account used by cloud-init by editing the `/etc/cloud/cloud.cfg` file and adding a line with a different user. For example, to configure cloud-init to put the key in an account named `admin`, edit the configuration file so it has the line:

```
user: admin
```

## Write a custom script to fetch the public key

If you are unable or unwilling to install cloud-init inside the guest, you can write a custom script to fetch the public key and add it to a user account.

To fetch the ssh public key and add it to the root account, edit the `/etc/rc.local` file and add the following lines before the line `"touch /var/lock/subsys/local"`. This code fragment is taken from the [rackerjoe oz-image-build CentOS 6 template](#).

```
if [ ! -d /root/.ssh ]; then
    mkdir -p /root/.ssh
    chmod 700 /root/.ssh
fi

# Fetch public key using HTTP
ATTEMPTS=30
```

```
FAILED=0
while [ ! -f /root/.ssh/authorized_keys ]; do
  curl -f http://169.254.169.254/latest/meta-data/public-keys/0/openssh-key
  > /tmp/metadata-key 2>/dev/null
  if [ $? -eq 0 ]; then
    cat /tmp/metadata-key >> /root/.ssh/authorized_keys
    chmod 0600 /root/.ssh/authorized_keys
    restorecon /root/.ssh/authorized_keys
    rm -f /tmp/metadata-key
    echo "Successfully retrieved public key from instance metadata"
    echo "*****"
    echo "AUTHORIZED KEYS"
    echo "*****"
    cat /root/.ssh/authorized_keys
    echo "*****"
  else
    FAILED=`expr $FAILED + 1`
    if [ $FAILED -ge $ATTEMPTS ]; then
      echo "Failed to retrieve public key from instance metadata after $FAILED
attempts, quitting"
      break
    fi
    echo "Could not retrieve public key from instance metadata (attempt #
$FAILED/$ATTEMPTS), retrying in 5 seconds..."
    sleep 5
  fi
done
```



### Note

Some VNC clients replace : (colon) with ; (semicolon) and \_ (underscore) with - (hyphen). If editing a file over a VNC session, make sure it's http: not http; and authorized\_keys not authorized-keys.

## Process user data and other metadata (cloud-init)

In addition to the ssh public key, an image might need additional information from OpenStack, such as [user data](#) that the user submitted when requesting the image. For example, you might want to set the host name of the instance when it is booted. Or, you might wish to configure your image so that it executes user data content as a script on boot.

This information is accessible through the metadata service or the [config drive](#). As the OpenStack metadata service is compatible with version 2009-04-04 of the Amazon EC2 metadata service, consult the Amazon EC2 documentation on [Using Instance Metadata](#) for details on how to retrieve user data.

The easiest way to support this type of functionality is to install the cloud-init package into your image, which is configured by default to treat user data as an executable script, and sets the host name.

## Ensure image writes boot log to console

You must configure the image so that the kernel writes the boot log to the `ttys0` device. In particular, the `console=ttys0` argument must be passed to the kernel on boot.

If your image uses grub2 as the boot loader, there should be a line in the grub configuration file. For example, `/boot/grub/grub.cfg`, which looks something like this:

```
linux /boot/vmlinuz-3.2.0-49-virtual root=UUID=6d2231e4-0975-4f35-
a94f-56738c1a8150 ro console=ttyS0
```

If `console=ttyS0` does not appear, you must modify your grub configuration. In general, you should not update the `grub.cfg` directly, since it is automatically generated. Instead, you should edit `/etc/default/grub` and modify the value of the `GRUB_CMDLINE_LINUX_DEFAULT` variable:

```
GRUB_CMDLINE_LINUX_DEFAULT="console=ttyS0"
```

Next, update the grub configuration. On Debian-based operating-systems such as Ubuntu, run this command:

```
# update-grub
```

On Fedora-based systems, such as RHEL and CentOS, and on openSUSE, run this command:

```
# grub2-mkconfig -o /boot/grub2/grub.cfg
```

## Paravirtualized Xen support in the kernel (Xen hypervisor only)

Prior to Linux kernel version 3.0, the mainline branch of the Linux kernel did not have support paravirtualized Xen virtual machine instances (what Xen calls DomU guests). If you are running the Xen hypervisor with paravirtualization, and you want to create an image for an older Linux distribution that has a pre 3.0 kernel, you must ensure that the image boots a kernel that has been compiled with Xen support.

## Manage the image cache

Use options in `nova.conf` to control whether, and for how long, unused base images are stored in `/var/lib/nova/instances/_base/`. If you have configured live migration of instances, all your compute nodes share one common `/var/lib/nova/instances/` directory.

For information about libvirt images in OpenStack, see [The life of an OpenStack libvirt image from Pádraig Brady](#).

**Table 3.1. Image cache management configuration options**

Configuration option=Default value	(Type) Description
<code>preallocate_images=none</code>	(StrOpt) VM image preallocation mode: <ul style="list-style-type: none"> <li><code>none</code>. No storage provisioning occurs up front.</li> <li><code>space</code>. Storage is fully allocated at instance start. The <code>\$instance_dir/</code> images are <b>allocated</b> to immediately determine if enough space is available, and to possibly improve VM I/O performance due to ongoing allocation avoidance, and better locality of block allocations.</li> </ul>

Configuration option=Default value	(Type) Description
remove_unused_base_images=True	(BoolOpt) Should unused base images be removed? When set to True, the interval at which base images are removed are set with the following two settings. If set to False base images are never removed by Compute.
remove_unused_original_minimum_age_seconds=86400	(IntOpt) Unused unresized base images younger than this are not removed. Default is 86400 seconds, or 24 hours.
remove_unused_resized_minimum_age_seconds=3600	(IntOpt) Unused resized base images younger than this are not removed. Default is 3600 seconds, or one hour.

To see how the settings affect the deletion of a running instance, check the directory where the images are stored:

```
# ls -lash /var/lib/nova/instances/_base/
```

In the `/var/log/compute/compute.log` file, look for the identifier:

```
2012-02-18 04:24:17 41389 WARNING nova.virt.libvirt.imagecache [-] Unknown
base file: /var/lib/nova/instances/_base/06a057b9c7b0b27e3b496f53d1e88810
a0d1d5d3_20
2012-02-18 04:24:17 41389 INFO nova.virt.libvirt.imagecache [-] Removable base
files: /var/lib/nova/instances/_base/06a057b9c7b0b27e3b496f53d1e88810
a0d1d5d3 /var/lib/nova/instances/_base/
06a057b9c7b0b27e3b496f53d1e88810a0d1d5d3_20
2012-02-18 04:24:17 41389 INFO nova.virt.libvirt.imagecache [-] Removing base
file: /var/lib/nova/instances/_base/06a057b9c7b0b27e3b496f53d1e88810a0d1d5d3
```

Because 86400 seconds (24 hours) is the default time for `remove_unused_original_minimum_age_seconds`, you can either wait for that time interval to see the base image removed, or set the value to a shorter time period in `nova.conf`. Restart all nova services after changing a setting in `nova.conf`.

## 4. Modify images

guestfish .....	17
guestmount .....	19
virt-* tools .....	19
Loop devices, kpartx, network block devices .....	20

Once you have obtained a virtual machine image, you may want to make some changes to it before uploading it to the OpenStack Image Service. Here we describe several tools available that allow you to modify images.



### Warning

Do not attempt to use these tools to modify an image that is attached to a running virtual machine. These tools are designed to only modify images that are not currently running.

## guestfish

The **guestfish** program is a tool from the [libguestfs](#) project that allows you to modify the files inside of a virtual machine image.



### Note

**guestfish** does not mount the image directly into the local file system. Instead, it provides you with a shell interface that enables you to view, edit, and delete files. Many of **guestfish** commands, such as **touch**, **chmod**, and **rm**, resemble traditional bash commands.

## Example guestfish session

Sometimes, you must modify a virtual machine image to remove any traces of the MAC address that was assigned to the virtual network interface card when the image was first created, because the MAC address will be different when it boots the next time. This example shows how to use **guestfish** to remove references to the old MAC address by deleting the `/etc/udev/rules.d/70-persistent-net.rules` file and removing the `HWADDR` line from the `/etc/sysconfig/network-scripts/ifcfg-eth0` file.

Assume that you have a CentOS qcow2 image called `centos63_desktop.img`. Mount the image in read-write mode as root, as follows:

```
# guestfish --rw -a centos63_desktop.img

Welcome to guestfish, the libguestfs filesystem interactive shell for
editing virtual machine filesystems.

Type: 'help' for help on commands
      'man' to read the manual
      'quit' to quit the shell
```

```
><fs>
```

This starts a guestfish session. Note that the guestfish prompt looks like a fish: `> <fs>`.

We must first use the **run** command at the guestfish prompt before we can do anything else. This will launch a virtual machine, which will be used to perform all of the file manipulations.

```
><fs> run
```

We can now view the file systems in the image using the **list-file systems** command:

```
><fs> list-file systems
/dev/vda1: ext4
/dev/vg_centosbase/lv_root: ext4
/dev/vg_centosbase/lv_swap: swap
```

We need to mount the logical volume that contains the root partition:

```
><fs> mount /dev/vg_centosbase/lv_root /
```

Next, we want to delete a file. We can use the **rm** guestfish command, which works the same way it does in a traditional shell.

```
><fs> rm /etc/udev/rules.d/70-persistent-net.rules
```

We want to edit the `ifcfg-eth0` file to remove the `HWADDR` line. The **edit** command will copy the file to the host, invoke your editor, and then copy the file back.

```
><fs> edit /etc/sysconfig/network-scripts/ifcfg-eth0
```

If you want to modify this image to load the 8021q kernel at boot time, you must create an executable script in the `/etc/sysconfig/modules/` directory. You can use the **touch** guestfish command to create an empty file, the **edit** command to edit it, and the **chmod** command to make it executable.

```
><fs> touch /etc/sysconfig/modules/8021q.modules
><fs> edit /etc/sysconfig/modules/8021q.modules
```

We add the following line to the file and save it:

```
modprobe 8021q
```

Then we set to executable:

```
><fs> chmod 0755 /etc/sysconfig/modules/8021q.modules
```

We're done, so we can exit using the **exit** command:

```
><fs> exit
```

## Go further with guestfish

There is an enormous amount of functionality in guestfish and a full treatment is beyond the scope of this document. Instead, we recommend that you read the [guestfs-recipes](#) documentation page for a sense of what is possible with these tools.



## guestmount

For some types of changes, you may find it easier to mount the image's file system directly in the guest. The **guestmount** program, also from the libguestfs project, allows you to do so.

For example, to mount the root partition from our `centos63_desktop.qcow2` image to `/mnt`, we can do:

```
# guestmount -a centos63_desktop.qcow2 -m /dev/vg_centosbase/lv_root --rw /mnt
```

If we didn't know in advance what the mount point is in the guest, we could use the `-i` (inspect) flag to tell guestmount to automatically determine what mount point to use:

```
# guestmount -a centos63_desktop.qcow2 -i --rw /mnt
```

Once mounted, we could do things like list the installed packages using `rpm`:

```
# rpm -qa --dbpath /mnt/var/lib/rpm
```

Once done, we unmount:

```
# umount /mnt
```

## virt-\* tools

The [libguestfs](#) project has a number of other useful tools, including:

- [virt-edit](#) for editing a file inside of an image.
- [virt-df](#) for displaying free space inside of an image.
- [virt-resize](#) for resizing an image.
- [virt-sysprep](#) for preparing an image for distribution (for example, delete SSH host keys, remove MAC address info, or remove user accounts).
- [virt-sparsify](#) for making an image sparse
- [virt-p2v](#) for converting a physical machine to an image that runs on KVM
- [virt-v2v](#) for converting Xen and VMWare images to KVM images

## Modify a single file inside of an image

This example shows how to use **virt-edit** to modify a file. The command can take either a filename as an argument with the `-a` flag, or a domain name as an argument with the `-d` flag. The following examples shows how to use this to modify the `/etc/shadow` file in instance with libvirt domain name `instance-000000e1` that is currently running:

```
# virsh shutdown instance-000000e1
# virt-edit -d instance-000000e1 /etc/shadow
```

```
# virsh start instance-000000e1
```

## Resize an image

Here's a simple example of how to use **virt-resize** to resize an image. Assume we have a 16 GB Windows image in qcow2 format that we want to resize to 50 GB. First, we use **virt-filesystems** to identify the partitions:

```
# virt-filesystems --long --parts --blkdevs -h -a /data/images/win2012.qcow2
Name      Type      MBR  Size  Parent
/dev/sda1  partition  07   350M  /dev/sda
/dev/sda2  partition  07   16G   /dev/sda
/dev/sda   device    -    16G   -
```

In this case, it's the `/dev/sda2` partition that we want to resize. We create a new qcow2 image and use the **virt-resize** command to write a resized copy of the original into the new image:

```
# qemu-img create -f qcow2 /data/images/win2012-50gb.qcow2 50G
# virt-resize --expand /dev/sda2 /data/images/win2012.qcow2 \
  /data/images/win2012-50gb.qcow2
Examining /data/images/win2012.qcow2 ...
*****

Summary of changes:

/dev/sda1: This partition will be left alone.

/dev/sda2: This partition will be resized from 15.7G to 49.7G.  The
           filesystem ntfs on /dev/sda2 will be expanded using the
           'ntfsresize' method.

*****
Setting up initial partition table on /data/images/win2012-50gb.qcow2 ...
Copying /dev/sda1 ...
 100% #####
00:00
Copying /dev/sda2 ...
 100% #####
00:00
Expanding /dev/sda2 using the 'ntfsresize' method ...

Resize operation completed with no errors.  Before deleting the old
disk, carefully check that the resized disk boots and works correctly.
```

## Loop devices, kpartx, network block devices

If you don't have access to `libguestfs`, you can mount image file systems directly in the host using loop devices, `kpartx`, and network block devices.



### Warning

Mounting untrusted guest images using the tools described in this section is a security risk, always use `libguestfs` tools such as `guestfish` and `guestmount` if you have access to them. See [A reminder why you should never mount guest disk images on the host OS](#) by Daniel Berrangé for more details.

## Mount a raw image (without LVM)

If you have a raw virtual machine image that is not using LVM to manage its partitions. First, use the **losetup** command to find an unused loop device.

```
# losetup -f
/dev/loop0
```

In this example, `/dev/loop0` is free. Associate a loop device with the raw image:

```
# losetup /dev/loop0 fedora17.img
```

If the image only has a single partition, you can mount the loop device directly:

```
# mount /dev/loop0 /mnt
```

If the image has multiple partitions, use **kpartx** to expose the partitions as separate devices (for example, `/dev/mapper/loop0p1`), then mount the partition that corresponds to the root file system:

```
# kpartx -av /dev/loop0
```

If the image has, say three partitions (`/boot`, `/`, `/swap`), there should be one new device created per partition:

```
$ ls -l /dev/mapper/loop0p*
brw-rw---- 1 root disk 43, 49 2012-03-05 15:32 /dev/mapper/loop0p1
brw-rw---- 1 root disk 43, 50 2012-03-05 15:32 /dev/mapper/loop0p2
brw-rw---- 1 root disk 43, 51 2012-03-05 15:32 /dev/mapper/loop0p3
```

To mount the second partition, as root:

```
# mkdir /mnt/image
# mount /dev/mapper/loop0p2 /mnt
```

Once you're done, to clean up:

```
# umount /mnt
# kpartx -d /dev/loop0
# losetup -d /dev/loop0
```

## Mount a raw image (with LVM)

If your partitions are managed with LVM, use **losetup** and **kpartx** as in the previous example to expose the partitions to the host:

```
# losetup -f
/dev/loop0
# losetup /dev/loop0 rhel62.img
# kpartx -av /dev/loop0
```

Next, you need to use the **vgscan** command to identify the LVM volume groups and then **vgchange** to expose the volumes as devices:

```
# vgscan
Reading all physical volumes. This may take a while...
Found volume group "vg_rhel62x8664" using metadata type lvm2
# vgchange -ay
```

```
2 logical volume(s) in volume group "vg_rhel62x8664" now active
# mount /dev/vg_rhel62x8664/lv_root /mnt
```

Clean up when you're done:

```
# umount /mnt
# vgchange -an vg_rhel62x8664
# kpartx -d /dev/loop0
# losetup -d /dev/loop0
```

## Mount a qcow2 image (without LVM)

You need the `nbdk` (network block device) kernel module loaded to mount qcow2 images. This will load it with support for 16 block devices, which is fine for our purposes. As root:

```
# modprobe nbdk max_part=16
```

Assuming the first block device (`/dev/nbd0`) is not currently in use, we can expose the disk partitions using the `qemu-nbd` and `partprobe` commands. As root:

```
# qemu-nbd -c /dev/nbd0 image.qcow2
# partprobe /dev/nbd0
```

If the image has, say three partitions (`/boot`, `/`, `/swap`), there should be one new device created for each partition:

```
$ ls -l /dev/nbd3*
brw-rw---- 1 root disk 43, 48 2012-03-05 15:32 /dev/nbd0
brw-rw---- 1 root disk 43, 49 2012-03-05 15:32 /dev/nbd0p1
brw-rw---- 1 root disk 43, 50 2012-03-05 15:32 /dev/nbd0p2
brw-rw---- 1 root disk 43, 51 2012-03-05 15:32 /dev/nbd0p3
```



### Note

If the network block device you selected was already in use, the initial `qemu-nbd` command will fail silently, and the `/dev/nbd3p{1,2,3}` device files will not be created.

If the image partitions are not managed with LVM, they can be mounted directly:

```
# mkdir /mnt/image
# mount /dev/nbd3p2 /mnt
```

When you're done, clean up:

```
# umount /mnt
# qemu-nbd -d /dev/g nbd0
```

## Mount a qcow2 image (with LVM)

If the image partitions are managed with LVM, after you use `qemu-nbd` and `partprobe`, you must use `vgscan` and `vgchange -ay` in order to expose the LVM partitions as devices that can be mounted:

```
# modprobe nbdk max_part=16
# qemu-nbd -c /dev/nbd0 image.qcow2
# partprobe /dev/nbd0# vgscan
```

```
Reading all physical volumes. This may take a while...
Found volume group "vg_rhel62x8664" using metadata type lvm2
# vgchange -ay
  2 logical volume(s) in volume group "vg_rhel62x8664" now active
# mount /dev/vg_rhel62x8664/lv_root /mnt
```

When you're done, clean up:

```
# umount /mnt
# vgchange -an vg_rhel62x8664
# qemu-nbd -d /dev/nbd0
```

## 5. Create images manually

Verify the libvirt default network is running .....	24
Use the virt-manager X11 GUI .....	24
Use virt-install and connect by using a local VNC client .....	26
Example: CentOS image .....	27
Example: Ubuntu image .....	34
Example: Microsoft Windows image .....	40
Example: FreeBSD image .....	41

Creating a new image is a step done outside of your OpenStack installation. You create the new image manually on your own system and then upload the image to your cloud.

To create a new image, you will need the installation CD or DVD ISO file for the guest operating system. You'll also need access to a virtualization tool. You can use KVM for this. Or, if you have a GUI desktop virtualization tool (such as, VMWare Fusion and VirtualBox), you can use that instead and just convert the file to raw once you're done.

When you create a new virtual machine image, you will need to connect to the graphical console of the hypervisor, which acts as the virtual machine's display and allows you to interact with the guest operating system's installer using your keyboard and mouse. KVM can expose the graphical console using the [VNC](#) (Virtual Network Computing) protocol or the newer [SPICE](#) protocol. We'll use the VNC protocol here, since you're more likely to be able to find a VNC client that works on your local desktop.

### Verify the libvirt default network is running

Before starting a virtual machine with libvirt, verify that the libvirt "default" network has been started. This network must be active for your virtual machine to be able to connect out to the network. Starting this network will create a Linux bridge (usually called `virbr0`), iptables rules, and a dnsmasq process that will serve as a DHCP server.

To verify that the libvirt "default" network is enabled, use the **virsh net-list** command and verify that the "default" network is active:

```
# virsh net-list
Name                State      Autostart
-----
default             active     yes
```

If the network is not active, start it by doing:

```
# virsh net-start default
```

### Use the virt-manager X11 GUI

If you plan to create a virtual machine image on a machine that can run X11 applications, the simplest way to do so is to use the **virt-manager** GUI, which is installable as the `virt-manager` package on both Fedora-based and Debian-based systems. This GUI has an embedded VNC client in it that will let you view and interact with the guest's graphical console.

If you are building the image on a headless server, and you have an X server on your local machine, you can launch **virt-manager** using ssh X11 forwarding to access the GUI. Since virt-manager interacts directly with libvirt, you typically need to be root to access it. If you can ssh directly in as root (or with a user that has permissions to interact with libvirt), do:

```
$ ssh -X root@server virt-manager
```

If the account you use to ssh into your server does not have permissions to run libvirt, but has sudo privileges, do:


```
$ ssh -X root@server  
$ sudo virt-manager
```



### Note

The `-X` flag passed to ssh will enable X11 forwarding over ssh. If this does not work, try replacing it with the `-Y` flag.

Click the "New" button at the top-left and step through the instructions.



You will be shown a series of dialog boxes that will allow you to specify information about the virtual machine.



### Note

When using qcow2 format images you should check the option 'customize before install', go to disk properties and explicitly select the qcow2 format. This ensures the virtual machine disk size will be correct.

## Use virt-install and connect by using a local VNC client

If you do not wish to use virt-manager (for example, you do not want to install the dependencies on your server, you don't have an X server running locally, the X11 forwarding over SSH isn't working), you can use the **virt-install** tool to boot the virtual machine through libvirt and connect to the graphical console from a VNC client installed on your local machine.

Because VNC is a standard protocol, there are multiple clients available that implement the VNC spec, including [TigerVNC](#) (multiple platforms), [TightVNC](#) (multiple platforms), [RealVNC](#) (multiple platforms), [Chicken](#) (Mac OS X), [Krde](#) (KDE), and [Vinagre](#) (GNOME).

The following example shows how to use the **qemu-img** command to create an empty image file **virt-install** command to start up a virtual machine using that image file. As root:

```
# qemu-img create -f qcow2 /data/centos-6.4.qcow2 10G
# virt-install --virt-type kvm --name centos-6.4 --ram 1024 \
--cdrom=/data/CentOS-6.4-x86_64-netinstall.iso \
--disk path=/data/centos-6.4.qcow2,size=10,format=qcow2 \
--network network=default\
--graphics vnc,listen=0.0.0.0 --noautoconsole \
--os-type=linux --os-variant=rhel6
```

```
Starting install...
Creating domain...
```

```
0 B      00:00
```

```
Domain installation still in progress. You can reconnect to
the console to complete the installation process.
```

The KVM hypervisor starts the virtual machine with the libvirt name, `centos-6.4`, with 1024 MB of RAM. The virtual machine also has a virtual CD-ROM drive associated with the `/data/CentOS-6.4-x86_64-netinstall.iso` file and a local 10 GB hard disk in `qcow2` format that is stored in the host at `/data/centos-6.4.qcow2`. It configures networking to use libvirt's default network. There is a VNC server that is listening on all interfaces, and libvirt will not attempt to launch a VNC client automatically nor try to display the text console (`--no-autoconsole`). Finally, libvirt will attempt to optimize the configuration for a Linux guest running a RHEL 6.x distribution.



### Note

When using the libvirt default network, libvirt will connect the virtual machine's interface to a bridge called `virbr0`. There is a `dnsmasq` process managed by libvirt that will hand out an IP address on the `192.168.122.0/24` subnet, and libvirt has iptables rules for doing NAT for IP addresses on this subnet.

Run the **virt-install --os-variant list** command to see a range of allowed `--os-variant` options.

Use the **virsh vncdisplay *vm-name*** command to get the VNC port number.

```
# virsh vncdisplay centos-6.4
```



```
:1
```

In the example above, the guest `centos-6.4` uses VNC display `:1`, which corresponds to tcp port 5901. You should be able to connect to a VNC client running on your local machine to display `:1` on the remote machine and step through the installation process.

## Example: CentOS image

We'll run through an example of installing a CentOS image. This will focus mainly on CentOS 6.4. Because the CentOS installation process may change across versions, if you are using a different version of CentOS the installer steps may differ.

## Download a CentOS install ISO

1. Navigate to the [CentOS mirrors](#) page.
2. Click one of the HTTP links in the right-hand column next to one of the mirrors.
3. Click the folder link of the CentOS version you want to use. For example, `6.4/`.
4. Click the `isos/` folder link.
5. Click the `x86_64/` folder link for 64-bit images.
6. Click the ISO image you want to download. The `netinstall` ISO. For example, `CentOS-6.4-x86_64-netinstall.iso` is a good choice since it's a smaller image that will download missing packages from the Internet during the install process.

## Start the install process

Start the installation process using either **virt-manager** or **virt-install** as described in the previous section. If using **virt-install**, don't forget to connect your VNC client to the virtual machine.

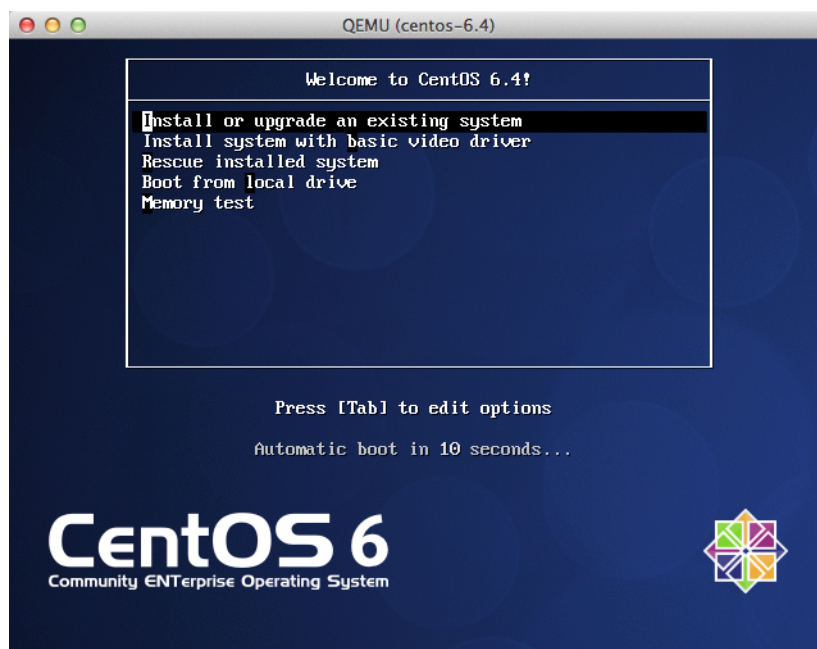
We will assume the name of your virtual machine image is `centos-6.4`, which we need to know when using **virsh** commands to manipulate the state of the image.

If you're using **virt-manager**, the commands should look something like this:

```
# qemu-img create -f qcow2 /tmp/centos-6.4.qcow2 10G
# virt-install --virt-type kvm --name centos-6.4 --ram 1024 \
--cdrom=/data/isos/CentOS-6.4-x86_64-netinstall.iso \
--disk /tmp/centos-6.4.qcow2,format=qcow2 \
--network network=default \
--graphics vnc,listen=0.0.0.0 --noautoconsole \
--os-type=linux --os-variant=rhel6
```

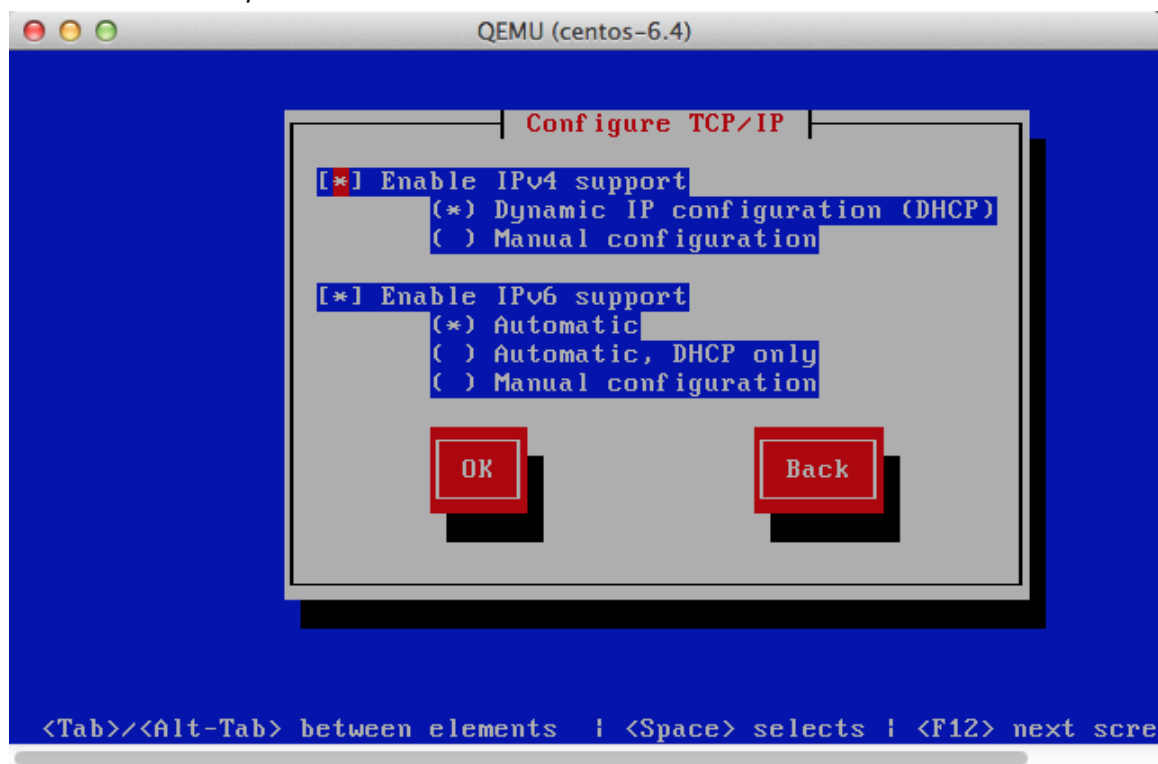
## Step through the install

At the initial Installer boot menu, choose the "Install or upgrade an existing system" option. Step through the install prompts, the defaults should be fine.



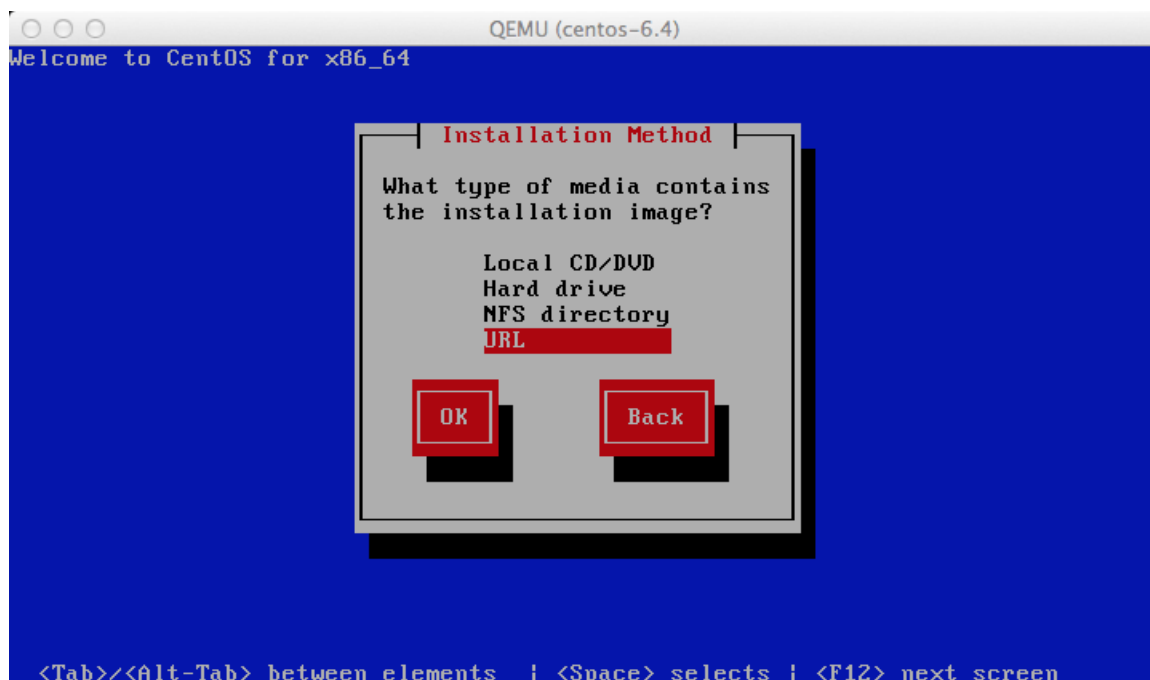
## Configure TCP/IP

The default TCP/IP settings are fine. In particular, ensure that Enable IPv4 support is enabled with DHCP, which is the default.



## Point the installer to a CentOS web server

Choose URL as the installation method.

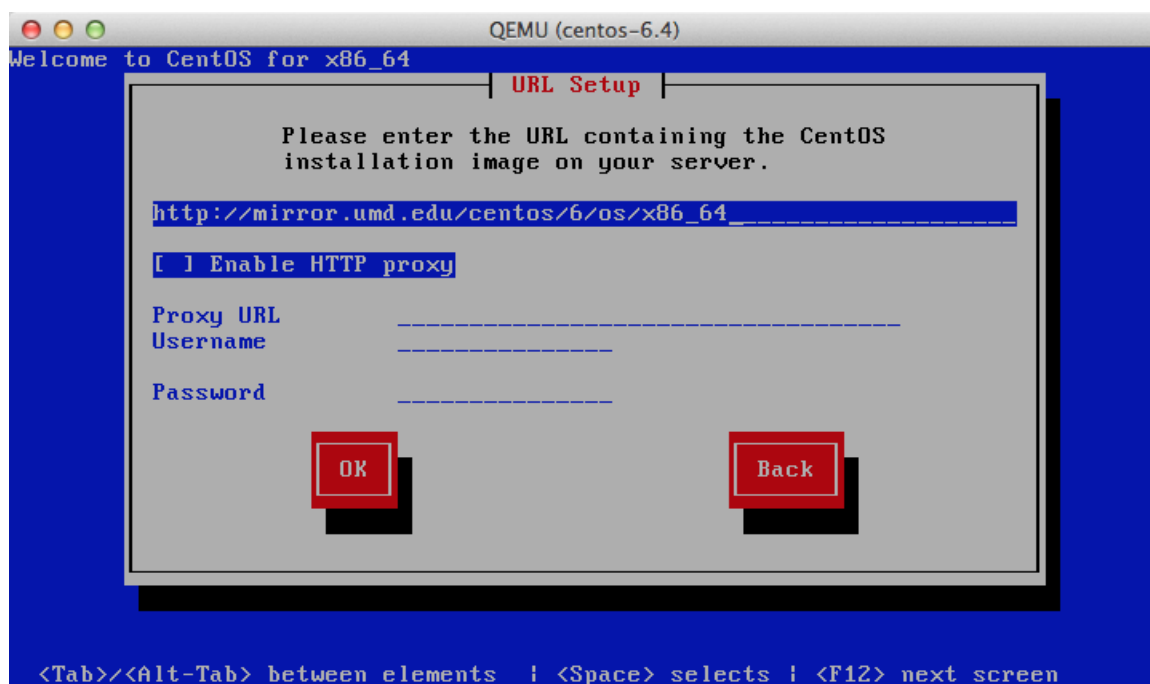


Depending on the version of CentOS, the net installer requires that the user specify either a URL, or the web site and a CentOS directory that corresponds to one of the CentOS mirrors. If the installer asks for a single URL, an example of a valid URL would be: `http://mirror.umd.edu/centos/6/os/x86_64`.



### Note

Consider using other mirrors as an alternative to `mirror.umd.edu`.



If the installer asks for web site name and CentOS directory separately, an example would be:

- Web site name: `mirror.umd.edu`
- CentOS directory: `centos/6/os/x86_64`

See [CentOS mirror page](#) to get a full list of mirrors, click on the "HTTP" link of a mirror to retrieve the web site name of a mirror.

## Storage devices

If asked about what type of devices your installation involves, choose "Basic Storage Devices".

## Hostname

The installer may ask you to choose a hostname. The default (`localhost.localdomain`) is fine. We will install the cloud-init package later, which will set the hostname on boot when a new instance is provisioned using this image.

## Partition the disks

There are different options for partitioning the disks. The default installation will use LVM partitions, and will create three partitions (`/boot`, `/`, `swap`), and this will work fine. Alternatively, you may wish to create a single ext4 partition, mounted to `/`, should also work fine.

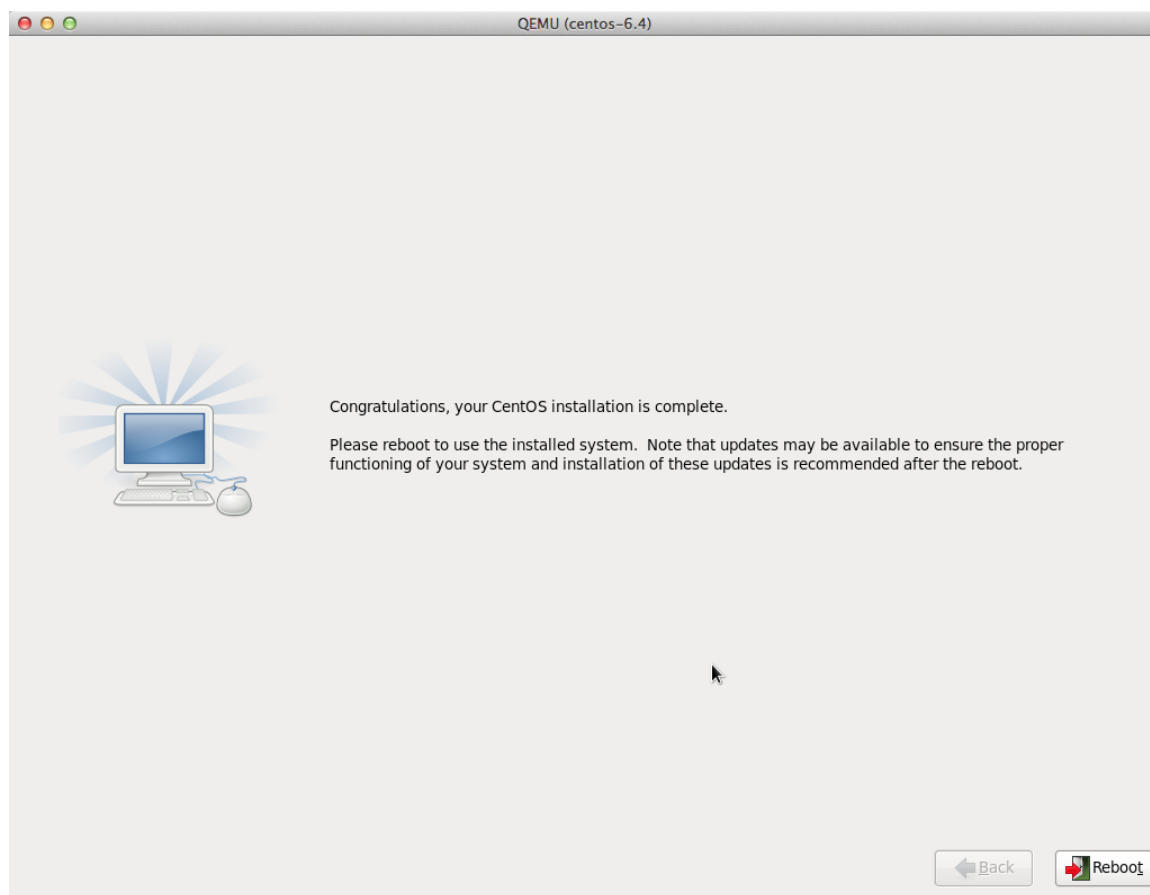
If unsure, we recommend you use the installer's default partition scheme, since there is no clear advantage to one scheme or another.

## Step through the install

Step through the install, using the default options. The simplest thing to do is to choose the "Basic Server" install (may be called "Server" install on older versions of CentOS), which will install an SSH server.

## Detach the CD-ROM and reboot

Once the install completes, you will see the screen "Congratulations, your CentOS installation is complete".



To eject a disk using **virsh**, libvirt requires that you attach an empty disk at the same target that the CDROM was previously attached, which should be `hdc`. You can confirm the appropriate target using the `dom dumpxml vm-image` command.

```
# virsh dumpxml centos-6.4
<domain type='kvm'>
  <name>centos-6.4</name>
  ...
  <disk type='block' device='cdrom'>
    <driver name='qemu' type='raw' />
    <target dev='hdc' bus='ide' />
    <readonly />
    <address type='drive' controller='0' bus='1' target='0' unit='0' />
  </disk>
  ...
</domain>
```

Run the following commands from the host to eject the disk and reboot using **virsh**, as root. If you are using **virt-manager**, the commands below will work, but you can also use the GUI to detach and reboot it by manually stopping and starting.

```
# virsh attach-disk --type cdrom --mode readonly centos-6.4 "" hdc
# virsh destroy centos-6.4
# virsh start centos-6.4
```



## Note

In theory, the **virsh reboot centos-6.4** command can be used instead of using **destroy** and **start** commands. However, in our testing we were unable to reboot successfully using the **virsh reboot** command.

## Log in to newly created image

When you boot for the first time after install, it may ask you about authentication tools, you can just choose "Exit". Then, log in as root using the root password you specified.

## Configure to fetch metadata

An instance must perform several steps on start up by interacting with the metadata service. For example, retrieve ssh public key and execute user data script. There are several ways to implement this functionality, including:

- Install a cloud-init RPM, which is a port of the Ubuntu [cloud-init](#) package. This is the recommended approach.
- Modify `/etc/rc.local` to fetch desired information from the metadata service, as described below.

## Use cloud-init to fetch the public key

The cloud-init package will automatically fetch the public key from the metadata server and place the key in an account. You can install cloud-init inside the CentOS guest by adding the EPEL repo:

```
# rpm -Uvh http://download.fedoraproject.org/pub/epel/6/x86_64/epel-release-6-8.noarch.rpm
# yum install cloud-init
```

The account varies by distribution. On Ubuntu-based virtual machines, the account is called "ubuntu". On Fedora-based virtual machines, the account is called "ec2-user".

You can change the name of the account used by cloud-init by editing the `/etc/cloud/cloud.cfg` file and adding a line with a different user. For example, to configure cloud-init to put the key in an account named `admin`, edit the configuration file so it has the line:

```
user: admin
```

## Write a script to fetch the public key (if no cloud-init)

If you are not able to install the cloud-init package in your image, to fetch the ssh public key and add it to the root account, edit the `/etc/rc.local` file and add the following lines before the line `"touch /var/lock/subsys/local"`

```
if [ ! -d /root/.ssh ]; then
    mkdir -p /root/.ssh
    chmod 700 /root/.ssh
fi

# Fetch public key using HTTP
```

```
ATTEMPTS=30
FAILED=0
while [ ! -f /root/.ssh/authorized_keys ]; do
  curl -f http://169.254.169.254/latest/meta-data/public-keys/0/openssh-key \
    > /tmp/metadata-key 2>/dev/null
  if [ $? -eq 0 ]; then
    cat /tmp/metadata-key >> /root/.ssh/authorized_keys
    chmod 0600 /root/.ssh/authorized_keys
    restorecon /root/.ssh/authorized_keys
    rm -f /tmp/metadata-key
    echo "Successfully retrieved public key from instance metadata"
    echo "*****"
    echo "AUTHORIZED KEYS"
    echo "*****"
    cat /root/.ssh/authorized_keys
    echo "*****"
  fi
done
```



### Note

Some VNC clients replace : (colon) with ; (semicolon) and \_ (underscore) with - (hyphen). Make sure it's http: not http; and authorized\_keys not authorized-keys.



### Note

The above script only retrieves the ssh public key from the metadata server. It does not retrieve *user data*, which is optional data that can be passed by the user when requesting a new instance. User data is often used for running a custom script when an instance comes up.

As the OpenStack metadata service is compatible with version 2009-04-04 of the Amazon EC2 metadata service, consult the Amazon EC2 documentation on [Using Instance Metadata](#) for details on how to retrieve user data.

## Disable the zeroconf route

In order for the instance to access the metadata service, disable the default zeroconf route:

```
# echo "NOZEROCONF=yes" >> /etc/sysconfig/network
```

## Configure console

In order for **nova console-log** to work properly on CentOS 6.x, guests you may need to add the following lines to `/boot/grub/menu.lst`

```
serial --unit=0 --speed=115200
terminal --timeout=10 console serial
# Edit the kernel line to add the console entries
kernel ... console=tty0 console=ttyS0,115200n8
```

## Shut down the instance

From inside the instance, as root:

```
# /sbin/shutdown -h now
```

## Clean up (remove MAC address details)

The operating system records the MAC address of the virtual ethernet card in locations such as `/etc/sysconfig/network-scripts/ifcfg-eth0` and `/etc/udev/rules.d/70-persistent-net.rules` during the instance process. However, each time the image boots up, the virtual ethernet card will have a different MAC address, so this information must be deleted from the configuration file.

There is a utility called **virt-sysprep**, that performs various cleanup tasks such as removing the MAC address references. It will clean up a virtual machine image in place:

```
# virt-sysprep -d centos-6.4
```

## Undefine the libvirt domain

Now that the image is ready to be uploaded to the Image Service, you no longer need to have this virtual machine image managed by libvirt. Use the **virsh undefine vm-image** command to inform libvirt.

```
# virsh undefine centos-6.4
```

## Image is complete

The underlying image file you created with **qemu-img create**. For example, `/tmp/centos-6.4.qcow2` is now ready for uploading to the OpenStack Image Service.

## Example: Ubuntu image

We'll run through an example of installing an Ubuntu image. This will focus mainly on Ubuntu 12.04 (Precise Pangolin) server. Because the Ubuntu installation process may change across versions, if you are using a different version of Ubuntu the installer steps may differ.

## Download an Ubuntu install ISO

In this example, we'll use the network installation ISO, since it's a smaller image. The 64-bit 12.04 network installer ISO is at <http://archive.ubuntu.com/ubuntu/dists/precise/main/installer-amd64/current/images/netboot/mini.iso>

## Start the install process

Start the installation process using either **virt-manager** or **virt-install** as described in the previous section. If using **virt-install**, don't forget to connect your VNC client to the virtual machine.

We will assume the name of your virtual machine image is `ubuntu-12.04`, which we need to know when using **virsh** commands to manipulate the state of the image.

If you're using **virt-manager**, the commands should look something like this:


```
# qemu-img create -f qcow2 /tmp/precise.qcow2 10G
# virt-install --virt-type kvm --name precise --ram 1024 \
--cdrom=/data/isos/precise-64-mini.iso \
```



```
--disk /tmp/precise.qcow2,format=qcow2 \  
--network network=default \  
--graphics vnc,listen=0.0.0.0 --noautoconsole \  
--os-type=linux --os-variant=ubuntuprecise
```

## Step through the install

At the initial Installer boot menu, choose the "Install" option. Step through the install prompts, the defaults should be fine.



## Hostname

The installer may ask you to choose a hostname. The default (ubuntu) is fine. We will install the cloud-init package later, which will set the hostname on boot when a new instance is provisioned using this image.

## Select a mirror

The default mirror proposed by the installer should be fine.

## Step through the install

Step through the install, using the default options. When prompted for a username, the default (ubuntu) is fine.

## Partition the disks

There are different options for partitioning the disks. The default installation will use LVM partitions, and will create three partitions (`/boot`, `/`, `swap`), and this will work fine. Alternatively, you may wish to create a single ext4 partition, mounted to `/`, should also work fine.

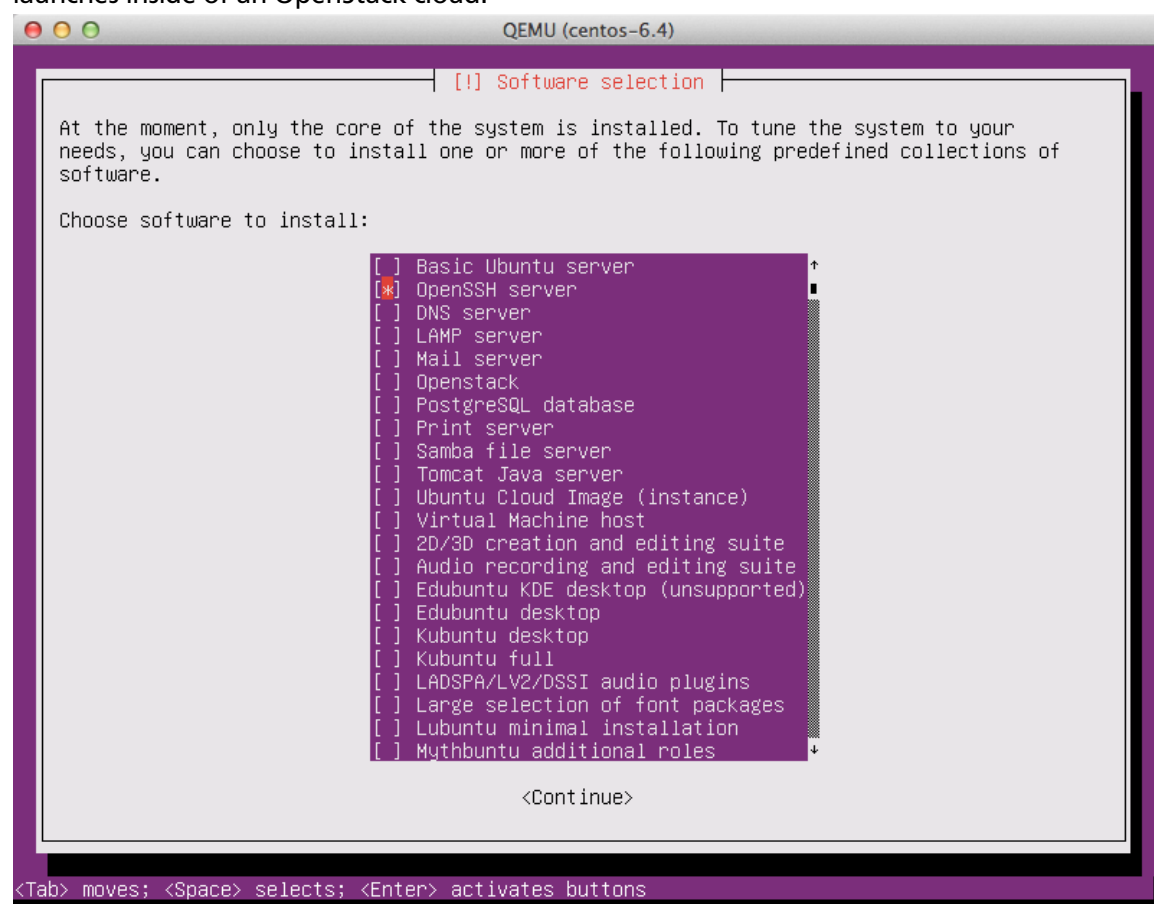
If unsure, we recommend you use the installer's default partition scheme, since there is no clear advantage to one scheme or another.

## Automatic updates

The Ubuntu installer will ask how you want to manage upgrades on your system. This option depends on your specific use case. If your virtual machine instances will be connected to the internet, we recommend "Install security updates automatically".

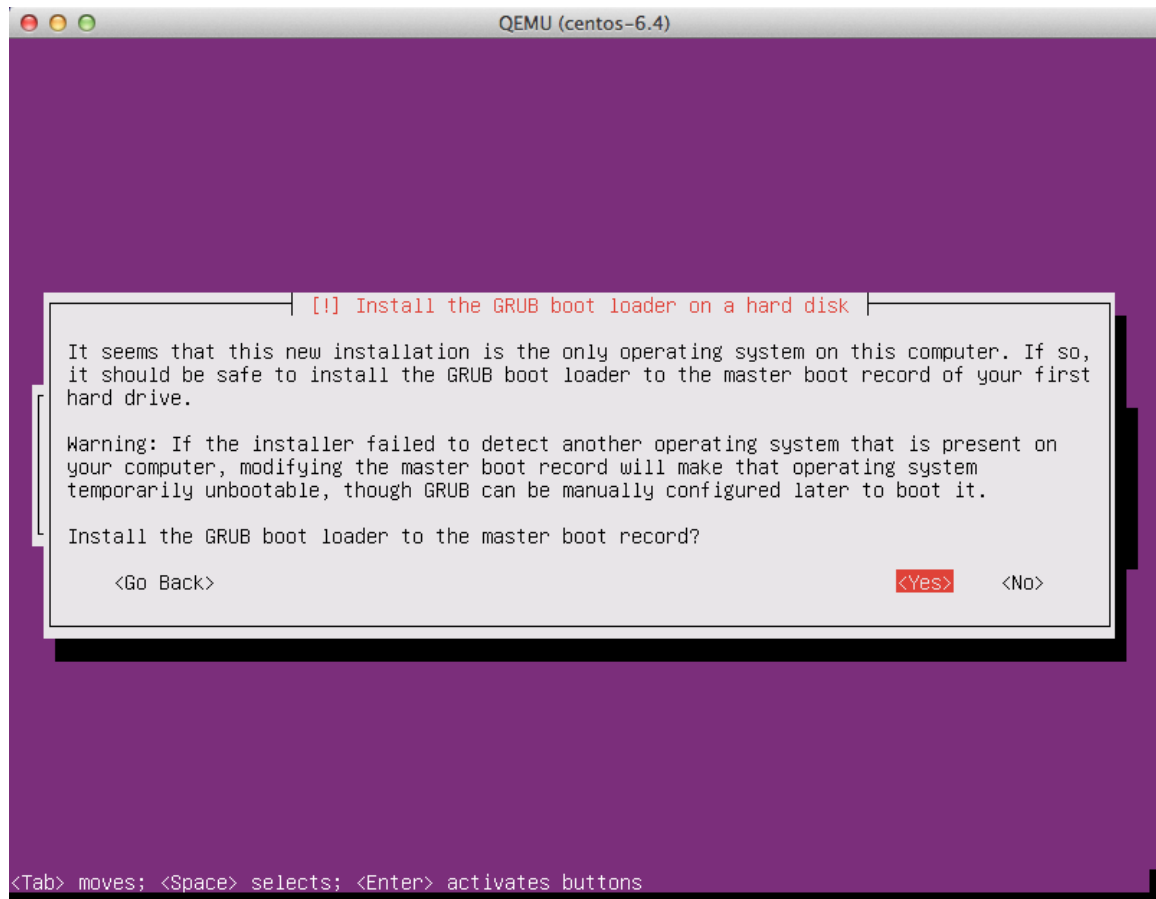
## Software selection: OpenSSH server

Choose "OpenSSH server" so that you will be able to SSH into the virtual machine when it launches inside of an OpenStack cloud.



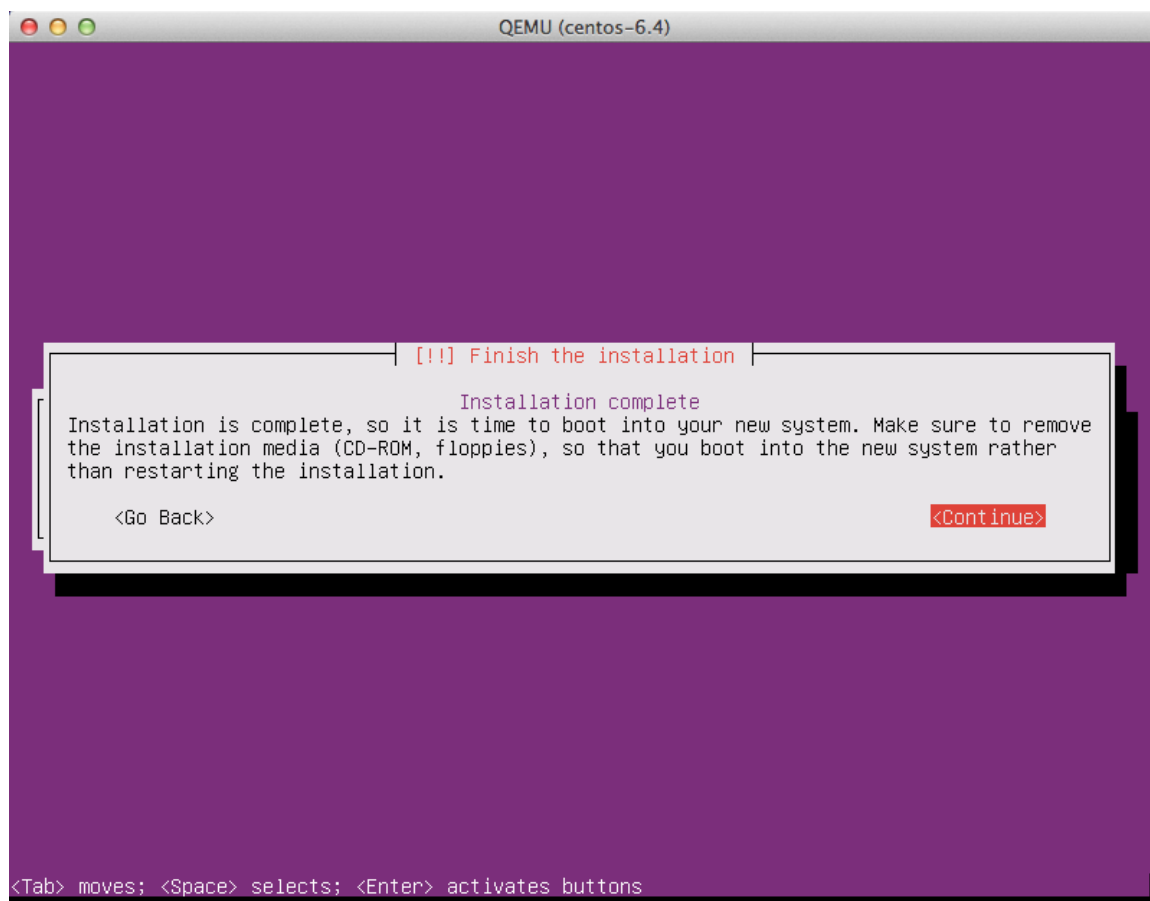
## Install GRUB boot loader

Select "Yes" when asked about installing the GRUB boot loader to the master boot record.



## Detach the CD-ROM and reboot

Select the defaults for all of the remaining options. When the installation is complete, you will be prompted to remove the CD-ROM.



## Note

When you hit "Continue" the virtual machine will shut down, even though it says it will reboot.

To eject a disk using **virsh**, libvirt requires that you attach an empty disk at the same target that the CDROM was previously attached, which should be `hdc`. You can confirm the appropriate target using the **dom dumpxml vm-image** command.

```
# virsh dumpxml precise
<domain type='kvm'>
  <name>precise</name>
  ...
  <disk type='block' device='cdrom'>
    <driver name='qemu' type='raw'/>
    <target dev='hdc' bus='ide'/>
    <readonly/>
    <address type='drive' controller='0' bus='1' target='0' unit='0'/>
  </disk>
  ...
</domain>
```

Run the following commands in the host as root to start up the machine again as paused, eject the disk and resume. If you are using virt-manager, you may use the GUI instead.

```
# virsh start precise --paused
```

```
# virsh attach-disk --type cdrom --mode readonly precise "" hdc
# virsh resume precise
```



## Note

In the example above, we start the instance paused, eject the disk, and then unpause. In theory, we could have ejected the disk at the "Installation complete" screen. However, our testing indicates that the Ubuntu installer locks the drive so that it cannot be ejected at that point.

## Log in to newly created image

When you boot for the first time after install, it may ask you about authentication tools, you can just choose 'Exit'. Then, log in as root using the root password you specified.

## Install cloud-init

The **cloud-init** script starts on instance boot and will search for a metadata provider to fetch a public key from. The public key will be placed in the default user account for the image.

Install the cloud-init package:

```
# apt-get install cloud-init
```

When building Ubuntu images **cloud-init** must be explicitly configured for the metadata source in use. The OpenStack metadata server emulates the EC2 metadata service used by images in Amazon EC2.

To set the metadata source to be used by the image run the **dpkg-reconfigure** command against the cloud-init package. When prompted select the EC2 data source:

```
# dpkg-reconfigure cloud-init
```

The account varies by distribution. On Ubuntu-based virtual machines, the account is called "ubuntu". On Fedora-based virtual machines, the account is called "ec2-user".

You can change the name of the account used by cloud-init by editing the `/etc/cloud/cloud.cfg` file and adding a line with a different user. For example, to configure cloud-init to put the key in an account named "admin", edit the config file so it has the line:

```
user: admin
```

## Shut down the instance

From inside the instance, as root:

```
# /sbin/shutdown -h now
```

## Clean up (remove MAC address details)

The operating system records the MAC address of the virtual ethernet card in locations such as `/etc/udev/rules.d/70-persistent-net.rules` during the instance process.

However, each time the image boots up, the virtual ethernet card will have a different MAC address, so this information must be deleted from the configuration file.

There is a utility called **virt-sysprep**, that performs various cleanup tasks such as removing the MAC address references. It will clean up a virtual machine image in place:

```
# virt-sysprep -d precise
```

## Undefine the libvirt domain

Now that the image is ready to be uploaded to the Image Service, we no longer need to have this virtual machine image managed by libvirt. Use the **virsh undefine vm-image** command to inform libvirt

```
# virsh undefine precise
```

## Image is complete

The underlying image file that you created with **qemu-img create**, such as `/tmp/precise.qcow2`, is now ready for uploading to the OpenStack Image Service.

## Example: Microsoft Windows image

We do not yet have a fully documented example of how to create a Microsoft Windows image. You can use libvirt to install Windows from an installation DVD using the same approach as with the CentOS and Ubuntu examples. Once the initial install is done, you will need to perform some Windows-specific customizations.

## Install VirtIO drivers

Installing the [VirtIO paravirtualization drivers for Windows](#) will improve virtual machine performance when using KVM as the hypervisor to run Windows.

## Sysprep

Microsoft has a special tool called [Sysprep](#) that must be run inside of a Windows guest to prepare it for use as a virtual machine image. On Windows 2012 server, invoke sysprep, as follows:

```
c:\windows\system32\sysprep\sysprep.exe
```

## cloudbase-init

[cloudbase-init](#) is a Windows port of cloud-init that should be installed inside of the guest. The [source code](#) is available on github.

## Jordan Rinke's OpenStack Windows resources

Jordan Rinke maintains [a collection of resources](#) for managing OpenStack Windows virtual machine guests.

## Windows instances require the timezone to be "localtime"

When uploading Windows images, you must set the `os_type` parameter to `windows`, which causes instances to use the local time zone. Failing to set this will cause the clock timezone to be UTC, which has unexpected side effects in Windows.

## Example: FreeBSD image

We do not yet have a fully documented example of how to create a FreeBSD image.

See the [bsd-cloudinit](#) project for information on how to build a FreeBSD VM image that works with OpenStack.

## 6. Tool support for image creation

Oz .....	42
VMBuilder .....	43
BoxGrinder .....	44
VeeWee .....	44
imagefactory .....	44
SUSE Studio .....	44

There are several tools that are designed to automate image creation.

### Oz

**Oz** is a command-line tool that automates the process of creating a virtual machine image file. Oz is a Python app that interacts with KVM to step through the process of installing a virtual machine. It uses a predefined set of kickstart (Red Hat-based systems) and preseeds files (Debian-based systems) for operating systems that it supports, and it can also be used to create Microsoft Windows images. On Fedora, install Oz with yum:

```
# yum install oz
```



#### Note

As of this writing, there are no Oz packages for Ubuntu, so you will need to either install from source or build your own .deb file.

A full treatment of Oz is beyond the scope of this document, but we will provide an example. You can find additional examples of Oz template files on github at [rackerjoe/oz-image-build/templates](https://github.com/rackerjoe/oz-image-build/templates). Here's how you would create a CentOS 6.4 image with Oz.

Create a template file (we'll call it `centos64.tdl`) with the following contents. The only entry you will need to change is the `<rootpw>` contents.

```
<template>
  <name>centos64</name>
  <os>
    <name>CentOS-6</name>
    <version>4</version>
    <arch>x86_64</arch>
    <install type='iso'>
      <iso>http://mirror.rackspace.com/CentOS/6/isos/x86_64/CentOS-6.4-x86_64-
bin-DVD1.iso</iso>
    </install>
    <rootpw>CHANGE THIS TO YOUR ROOT PASSWORD</rootpw>
  </os>
  <description>CentOS 6.4 x86_64</description>
  <repositories>
    <repository name='epel-6'>
      <url>http://download.fedoraproject.org/pub/epel/6/$basearch</url>
      <signed>no</signed>
    </repository>
  </repositories>
  <packages>
```



```
<package name='epel-release' />
<package name='cloud-utils' />
<package name='cloud-init' />
</packages>
<commands>
  <command name='update'>
yum update
yum clean all
sed -i '^HWADDR/d' /etc/sysconfig/network-scripts/ifcfg-eth0
echo -n > /etc/udev/rules.d/70-persistent-net.rules
echo -n > /lib/udev/rules.d/75-persistent-net-generator.rules
  </command>
</commands>
</template>
```

This Oz template specifies where to download the Centos 6.4 install ISO. Oz will use the version information to identify which kickstart file to use. In this case, it will be [RHEL6.auto](#). It adds EPEL as a repository and install the `epel-release`, `cloud-utils`, and `cloud-init` packages, as specified in the `packages` section of the file.

After Oz does the initial OS install using the kickstart file, it customizes the image by doing an update. It also removes any reference to the `eth0` device that `libvirt` creates while Oz does the customizing, as specified in the `command` section of the XML file.

To run this, do, as root:

```
# oz-install -d3 -u centos64.tdl -x centos64-libvirt.xml
```

- The `-d3` flag tells Oz to show status information as it runs.
- The `-u` tells Oz to do the customization (install extra packages, run the commands) once it does the initial install.
- The `-x <filename>` flag tells Oz what filename to use to write out a `libvirt` XML file (otherwise it will default to something like `centos64Apr_03_2013-12:39:42`).

If you leave out the `-u` flag, or you want to edit the file to do additional customizations, you can use the **oz-customize** command, using the `libvirt` XML file that **oz-install** creates. For example:

```
# oz-customize -d3 centos64.tdl centos64-libvirt.xml
```

Oz will invoke `libvirt` to boot the image inside of `KVM`, then Oz will `ssh` into the instance and perform the customizations.

## VMBuilder

**VMBuilder** (Virtual Machine Builder) is a command-line tool that creates virtual machine images for different hypervisors. The version of **VMBuilder** that ships with Ubuntu can only create Ubuntu virtual machine guests. The version of **VMBuilder** that ships with Debian can create Ubuntu and Debian virtual machine guests.

The [Ubuntu Server Guide](#) has documentation on how to use **VMBuilder** to create an Ubuntu image.

## BoxGrinder

[BoxGrinder](#) is another tool for creating virtual machine images, which it calls appliances. BoxGrinder can create Fedora, Red Hat Enterprise Linux, or CentOS images. BoxGrinder is currently only supported on Fedora.

## VeeWee

[VeeWee](#) is often used to build [Vagrant](#) boxes, but it can also be used to build KVM images.

## imagefactory

[imagefactory](#) is a newer tool designed to automate the building, converting, and uploading images to different cloud providers. It uses Oz as its back-end and includes support for OpenStack-based clouds.

## SUSE Studio

[SUSE Studio](#) is a web application for building and testing software applications in a web browser. It supports the creation of physical, virtual or cloud-based applications and includes support for building images for OpenStack based clouds using SUSE Linux Enterprise and openSUSE as distributions.

## 7. Converting between image formats

Converting images from one format to another is generally straightforward.

### qemu-img convert: raw, qcow2, VDI, VMDK

The **qemu-img convert** command can do conversion between multiple formats, including raw, qcow2, VDI (VirtualBox), VMDK (VMWare) and VHD (Hyper-V).

**Table 7.1. qemu-img format strings**

Image format	Argument to qemu-img
raw	raw
qcow2	qcow2
VDI (VirtualBox)	vdi
VMDK (VMWare)	vmdk
VHD (Hyper-V)	vpc

This example will convert a raw image file named centos63.dsk to a qcow2 image file.

```
$ qemu-img convert -f raw -O qcow2 centos64.dsk centos64.qcow2
```

To convert from vmdk to raw, you would do:

```
$ qemu-img convert -f vmdk -O raw centos64.vmdk centos64.img
```



#### Note

The `-f format` flag is optional. If omitted, qemu-img will try to infer the image format.

### VBoxManage: VDI (VirtualBox) to raw

If you've created a VDI image using VirtualBox, you can convert it to raw format using the **VBoxManage** command-line tool that ships with VirtualBox. On Mac OS X, VirtualBox stores images by default in the `~/VirtualBox VMs/` directory. The following example creates a raw image in the current directory from a VirtualBox VDI image.

```
$ VBoxManage clonehd ~/VirtualBox\ VMs/fedora18.vdi fedora18.img --format raw
```

# Appendix A. Community support

## Table of Contents

Documentation .....	46
ask.openstack.org .....	47
OpenStack mailing lists .....	47
The OpenStack wiki .....	47
The Launchpad Bugs area .....	48
The OpenStack IRC channel .....	49
Documentation feedback .....	49
OpenStack distribution packages .....	49

Many resources are available to help you run and use OpenStack. Members of the OpenStack community can answer questions and help with bug suspicions. We are constantly improving and adding to the main features of OpenStack, but if you have any problems, do not hesitate to ask. Use the following resources to get OpenStack support and troubleshoot your existing installations.

## Documentation

For the available OpenStack documentation, see [docs.openstack.org](https://docs.openstack.org).

To provide feedback on documentation, join and use the `<openstack-docs@lists.openstack.org>` mailing list at [OpenStack Documentation Mailing List](#), or [report a bug](#).

The following books explain how to install an OpenStack cloud and its associated components:

- [Installation Guide for Debian 7.0](#)
- [Installation Guide for openSUSE and SUSE Linux Enterprise Server](#)
- [Installation Guide for Red Hat Enterprise Linux, CentOS, and Fedora](#)
- [Installation Guide for Ubuntu 12.04 \(LTS\)](#)

The following books explain how to configure and run an OpenStack cloud:

- [Cloud Administrator Guide](#)
- [Configuration Reference](#)
- [Operations Guide](#)
- [High Availability Guide](#)
- [Security Guide](#)
- [Virtual Machine Image Guide](#)

The following books explain how to use the OpenStack dashboard and command-line clients:

- [API Quick Start](#)
- [End User Guide](#)
- [Admin User Guide](#)
- [Command-Line Interface Reference](#)

The following documentation provides reference and guidance information for the OpenStack APIs:

- [OpenStack API Reference](#)
- [OpenStack Block Storage Service API v2 Reference](#)
- [OpenStack Compute API v2 and Extensions Reference](#)
- [OpenStack Identity Service API v2.0 Reference](#)
- [OpenStack Image Service API v2 Reference](#)
- [OpenStack Networking API v2.0 Reference](#)
- [OpenStack Object Storage API v1 Reference](#)

The [Training Guides](#) offer software training for cloud administration and management.

## ask.openstack.org

During the set up or testing of OpenStack, you might have questions about how a specific task is completed or be in a situation where a feature does not work correctly. Use the [ask.openstack.org](http://ask.openstack.org) site to ask questions and get answers. When you visit the <http://ask.openstack.org> site, scan the recently asked questions to see whether your question has already been answered. If not, ask a new question. Be sure to give a clear, concise summary in the title and provide as much detail as possible in the description. Paste in your command output or stack traces, links to screen shots, and so on.

## OpenStack mailing lists

A great way to get answers and insights is to post your question or problematic scenario to the OpenStack mailing list. You can learn from and help others who might have similar issues. To subscribe or view the archives, go to <http://lists.openstack.org/cgi-bin/mailman/listinfo/openstack>. You might be interested in the other mailing lists for specific projects or development, which you can find [on the wiki](#). A description of all mailing lists is available at <http://wiki.openstack.org/MailingLists>.

## The OpenStack wiki

The [OpenStack wiki](#) contains a broad range of topics but some of the information can be difficult to find or is a few pages deep. Fortunately, the wiki search feature enables you to

search by title or content. If you search for specific information, such as about networking or nova, you can find lots of relevant material. More is being added all the time, so be sure to check back often. You can find the search box in the upper right corner of any OpenStack wiki page.

## The Launchpad Bugs area

The OpenStack community values your set up and testing efforts and wants your feedback. To log a bug, you must sign up for a Launchpad account at <https://launchpad.net/+login>. You can view existing bugs and report bugs in the Launchpad Bugs area. Use the search feature to determine whether the bug has already been reported or even better, already fixed. If it still seems like your bug is unreported, fill out a bug report.

Some tips:

- Give a clear, concise summary!
- Provide as much detail as possible in the description. Paste in your command output or stack traces, links to screen shots, and so on.
- Be sure to include the software and package versions that you are using, especially if you are using a development branch, such as, "Grizzly release" vs git commit `bc79c3ecc55929bac585d04a03475b72e06a3208`.
- Any deployment specific information is helpful, such as Ubuntu 12.04 or multi-node install.

The Launchpad Bugs areas are available here:

- [Bugs: OpenStack Block Storage \(cinder\)](#)
- [Bugs: OpenStack Compute \(nova\)](#)
- [Bugs: OpenStack Dashboard \(horizon\)](#)
- [Bugs: OpenStack Identity \(keystone\)](#)
- [Bugs: OpenStack Image Service \(glance\)](#)
- [Bugs: OpenStack Networking \(neutron\)](#)
- [Bugs: OpenStack Object Storage \(swift\)](#)
- [Bugs: Bare Metal \(ironic\)](#)
- [Bugs: Data Processing Service \(sahara\)](#)
- [Bugs: Database Service \(trove\)](#)
- [Bugs: Orchestration \(heat\)](#)
- [Bugs: Telemetry \(ceilometer\)](#)
- [Bugs: Queue Service \(marconi\)](#)

- [Bugs: OpenStack API Documentation \(api.openstack.org\)](http://api.openstack.org)
- [Bugs: OpenStack Documentation \(docs.openstack.org\)](http://docs.openstack.org)

## The OpenStack IRC channel

The OpenStack community lives and breathes in the #openstack IRC channel on the Freenode network. You can hang out, ask questions, or get immediate feedback for urgent and pressing issues. To install an IRC client or use a browser-based client, go to <http://webchat.freenode.net/>. You can also use Colloquy (Mac OS X, <http://colloquy.info/>), mIRC (Windows, <http://www.mirc.com/>), or XChat (Linux). When you are in the IRC channel and want to share code or command output, the generally accepted method is to use a Paste Bin. The OpenStack project has one at <http://paste.openstack.org>. Just paste your longer amounts of text or logs in the web form and you get a URL you can paste into the channel. The OpenStack IRC channel is: #openstack on irc.freenode.net. You can find a list of all OpenStack-related IRC channels at <https://wiki.openstack.org/wiki/IRC>.

## Documentation feedback

To provide feedback on documentation, join and use the <openstack-docs@lists.openstack.org> mailing list at [OpenStack Documentation Mailing List](#), or [report a bug](#).

## OpenStack distribution packages

The following Linux distributions provide community-supported packages for OpenStack:

- **Debian:** <http://wiki.debian.org/OpenStack>
- **CentOS, Fedora, and Red Hat Enterprise Linux:** <http://openstack.redhat.com/>
- **openSUSE and SUSE Linux Enterprise Server:** <http://en.opensuse.org/Portal:OpenStack>
- **Ubuntu:** <https://wiki.ubuntu.com/ServerTeam/CloudArchive>