



Cisco Nexus 7000 Hardware Architecture

BRKARC-3470



Session Goal


- To provide you with a thorough understanding of the Cisco Nexus™ 7000 switching architecture, supervisor, fabric, and I/O module design, packet flows, and key forwarding engine functions
- This session will **not** examine Unified I/O, DCB, FCoE, NX-OS software architecture, or other Nexus platforms
- Related sessions:
BRKARC-3471: Cisco NXOS Software Architecture




Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Nexus 7010 Chassis



Nexus 7018 Chassis




Agenda

- Chassis Architecture
- **Supervisor Engine Architecture**
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Supervisor Engine

- Performs control plane and management functions
- Dual-core 1.66GHz Intel Xeon processor with 4GB DRAM
- 2MB NVRAM, 2GB internal bootdisk, compact flash slots
- Out-of-band 10/100/1000 management interface
- Always-on Connectivity Management Processor (CMP) for lights-out management
- Console and auxiliary serial ports
- USB ports for file transfer

N7K-SUP1




Management Interfaces

Management Ethernet

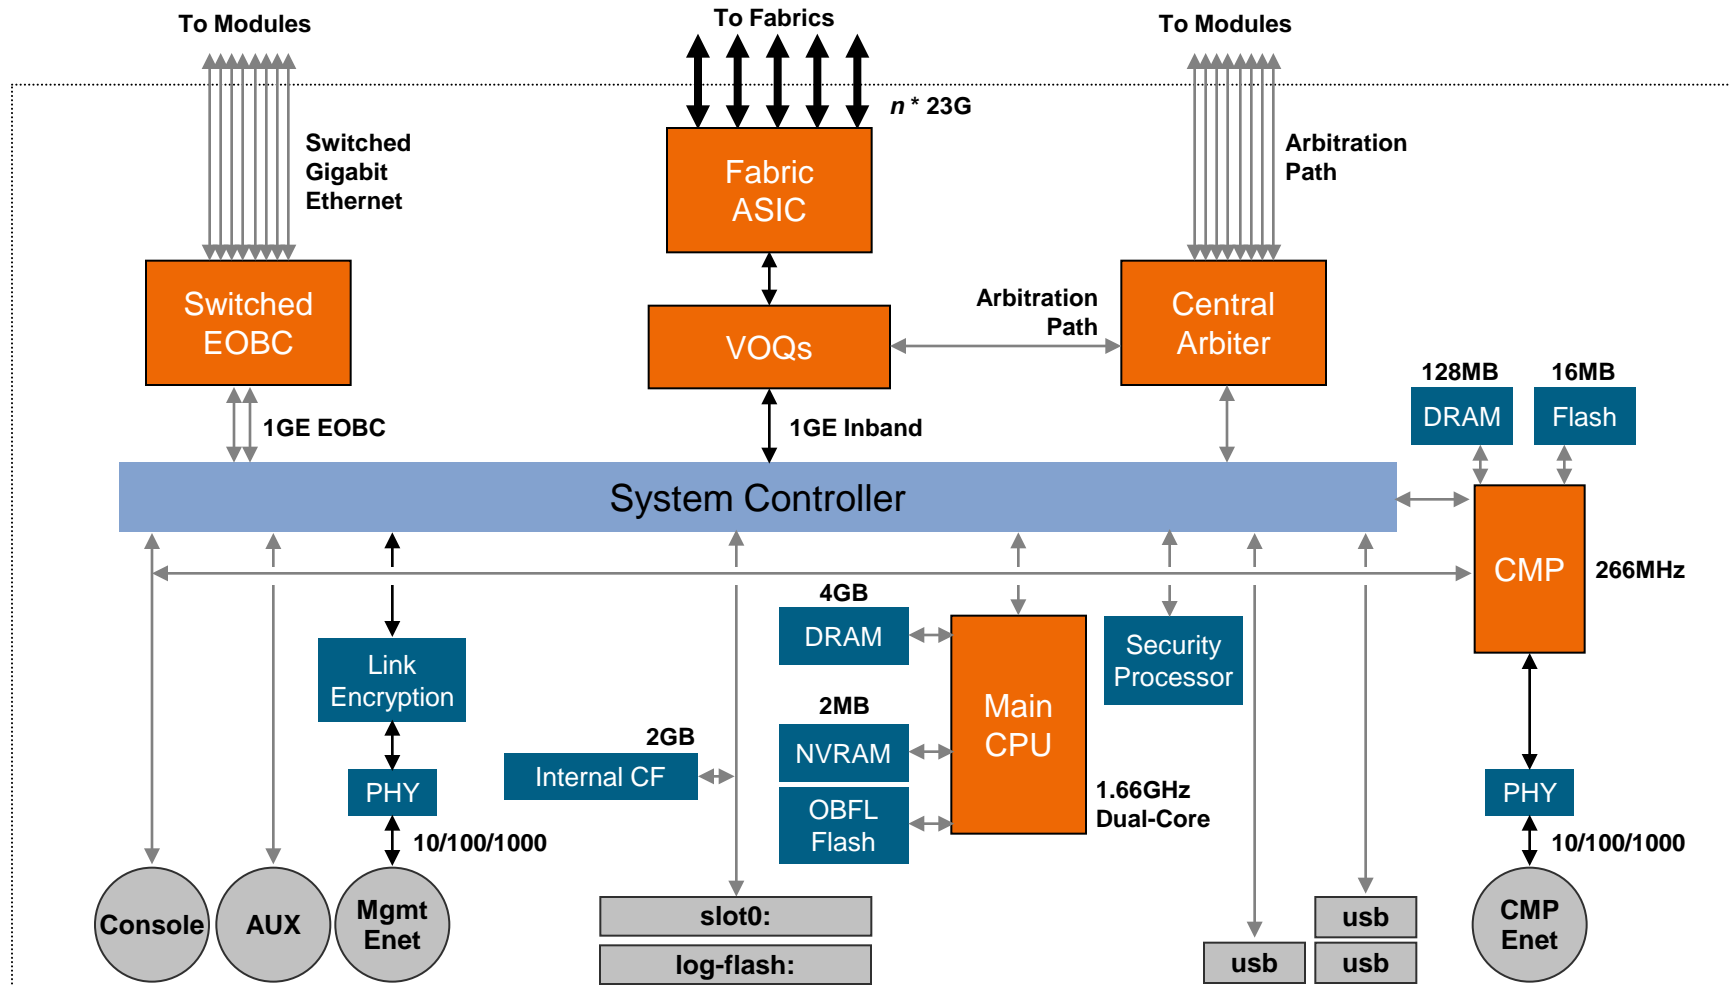
- 10/100/1000 interface used exclusively for system management
- Belongs to dedicated “management” VRF
 - Prevents data plane traffic from entering/exiting from mgmt0 interface
 - Cannot move mgmt0 interface to another VRF
 - Cannot assign other system ports to management VRF

Connectivity Management Processor (CMP) Ethernet

- Connects to standalone, always-on microprocessor on supervisor engine
 - Runs lightweight software with network stack
 - Completely independent of NX-OS on main CPU
- Provides ‘lights out’ remote management and disaster recovery via 10/100/1000 interface
 - Removes need for terminal servers



Supervisor Engine Architecture



Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- **I/O Module Architecture**
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

8-Port 10GE I/O Module

N7K-M108X2-12L

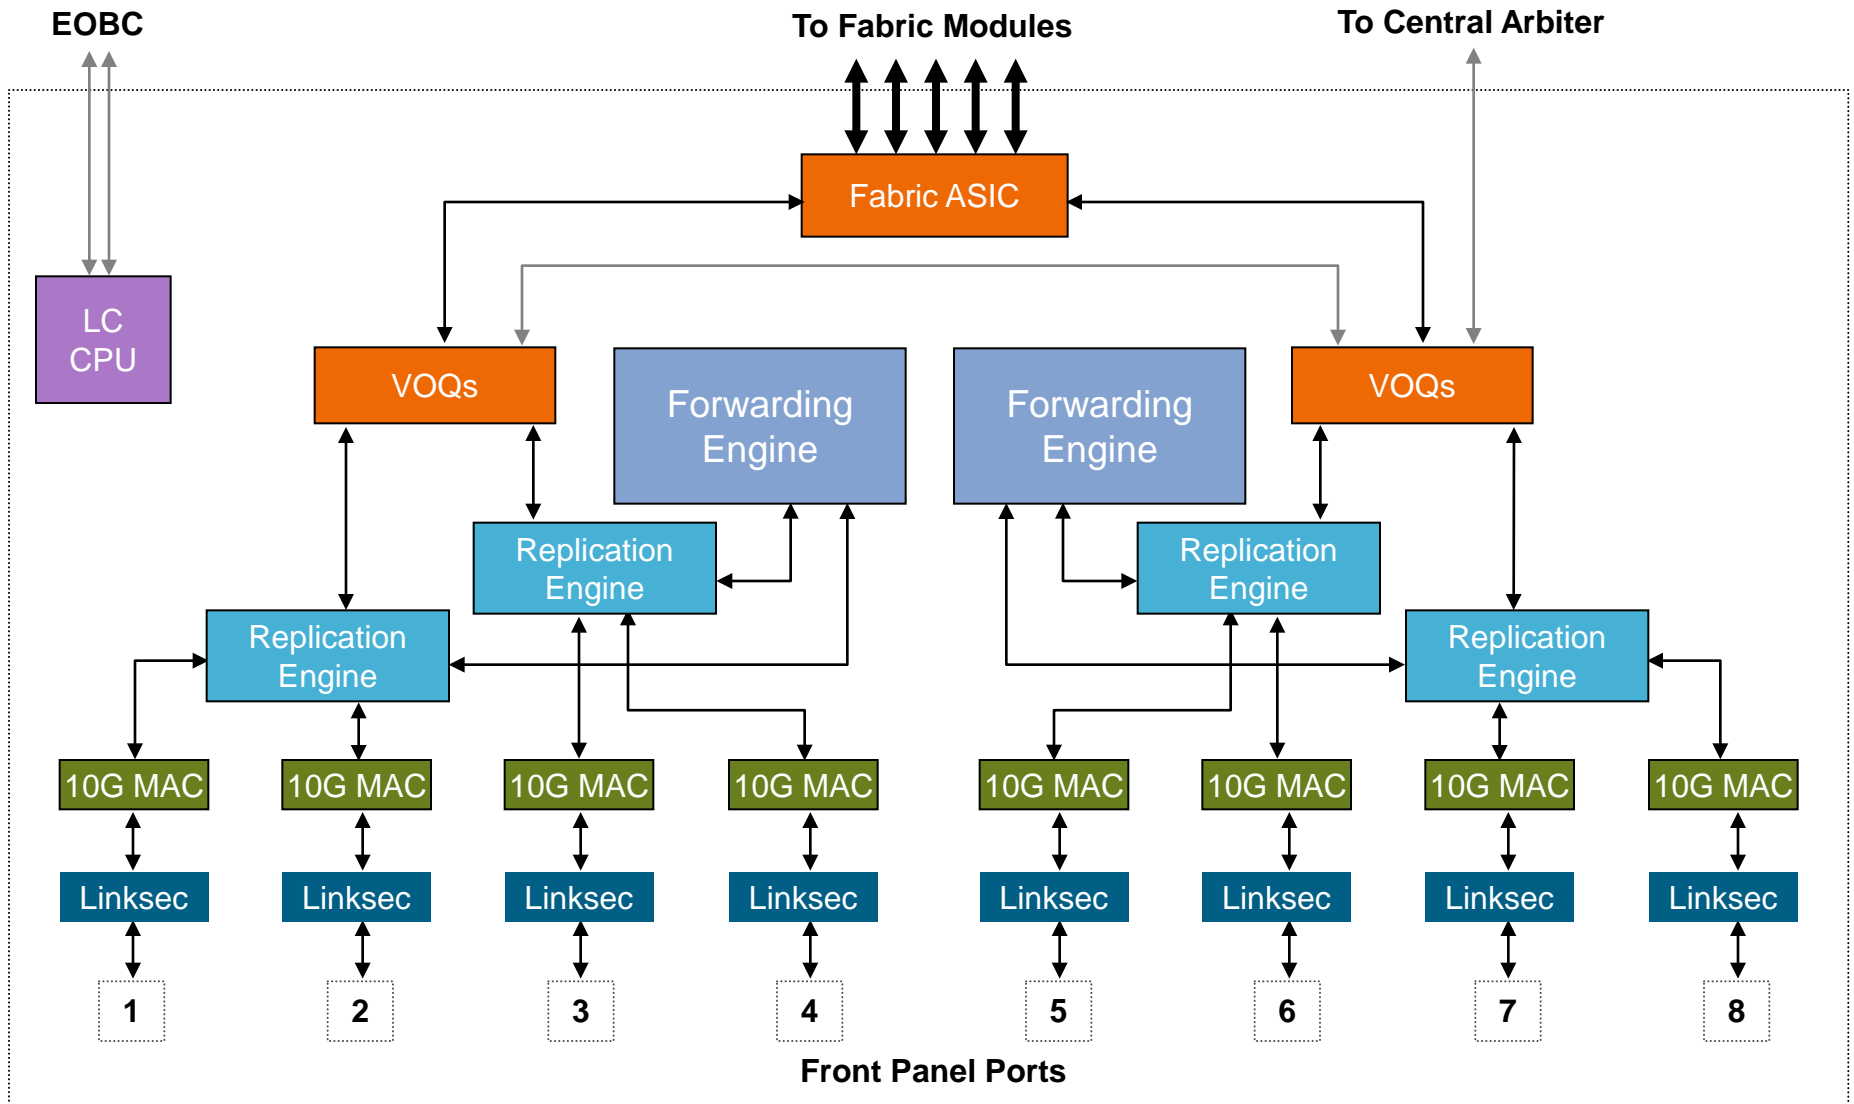
- 8-port 10G with X2 transceivers
- 80G full-duplex fabric connectivity
- Two integrated forwarding engines (120Mpps)
Support for “XL” forwarding tables (licensed feature)
- 8 ports wire-rate L3 multicast replication
- 802.1AE LinkSec

N7K-M108X2-12L



8-Port 10G XL I/O Module Architecture


N7K-M108X2-12L



32-Port 10GE I/O Module


N7K-M132XP-12

- 32-port 10G with SFP+ transceivers
- 80G full-duplex fabric connectivity
- Integrated 60Mpps forwarding engine
- Oversubscription option for higher density (up to 4:1)
- 8 ports wire-rate L3 multicast replication
- 802.1AE LinkSec



Shared vs. Dedicated Mode

rate-mode shared
(default)




Shared mode

- Four interfaces in port group share 10G bandwidth

“Port group”—group of contiguous even or odd ports that share 10G of bandwidth (e.g., ports 1,3,5,7)

rate-mode dedicated

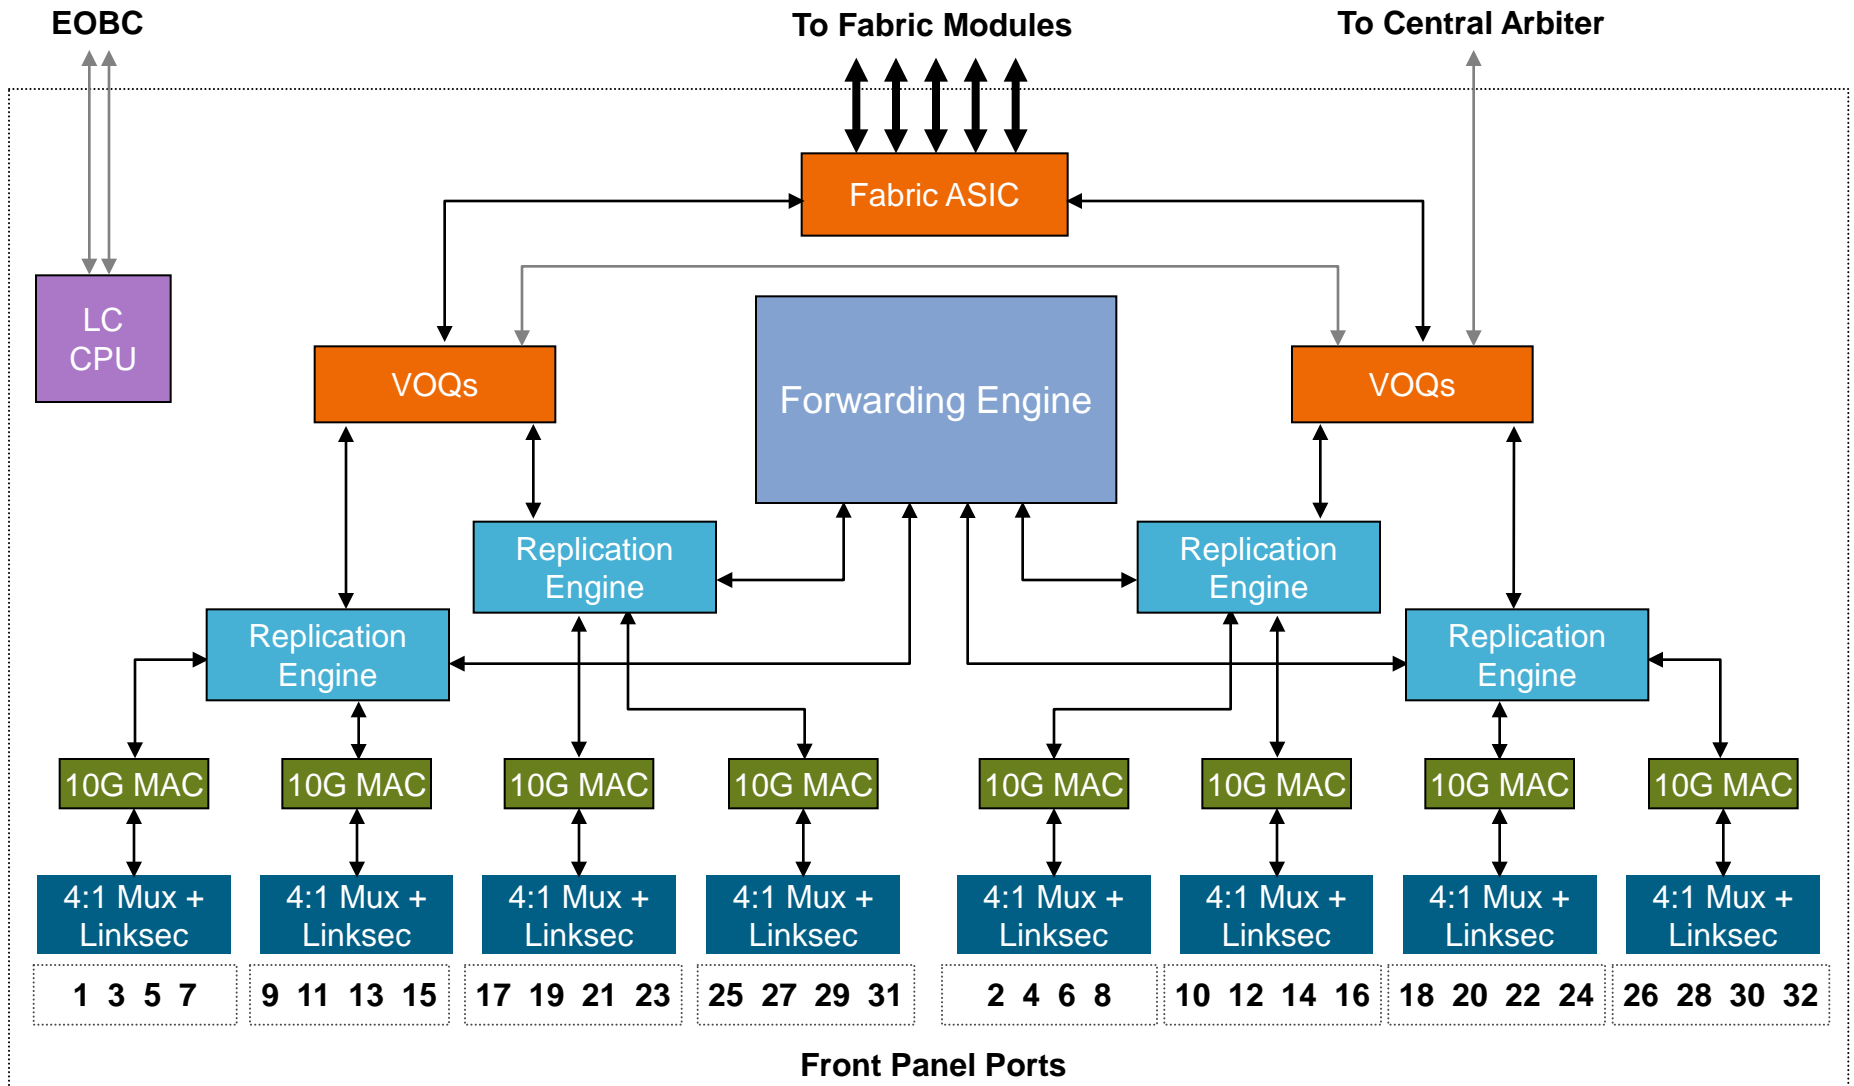


Dedicated mode

- First interface in port group gets 10G bandwidth
- Other three interfaces in port group disabled

32-Port 10G I/O Module Architecture

N7K-M132XP-12



48-Port 1G I/O Modules

N7K-M148GT-11, N7K-M148GS-11, N7K-M148GS-11L

- Three 1G I/O module options:
 - 48 10/100/1000 RJ-45 ports (N7K-M148GT-11)
 - 48 1G SFP ports (N7K-M148GS-11)
 - 48 1G SFP ports with XL forwarding engine (N7K-M148GS-11L)
- Integrated 60Mpps forwarding engine
- 46G full duplex fabric connectivity
 - Line rate on 48-ports with some local switching
- 48 ports wire-rate L3 multicast replication
- 802.1AE LinkSec



N7K-M148GT-11
Release 4.0(1) and later



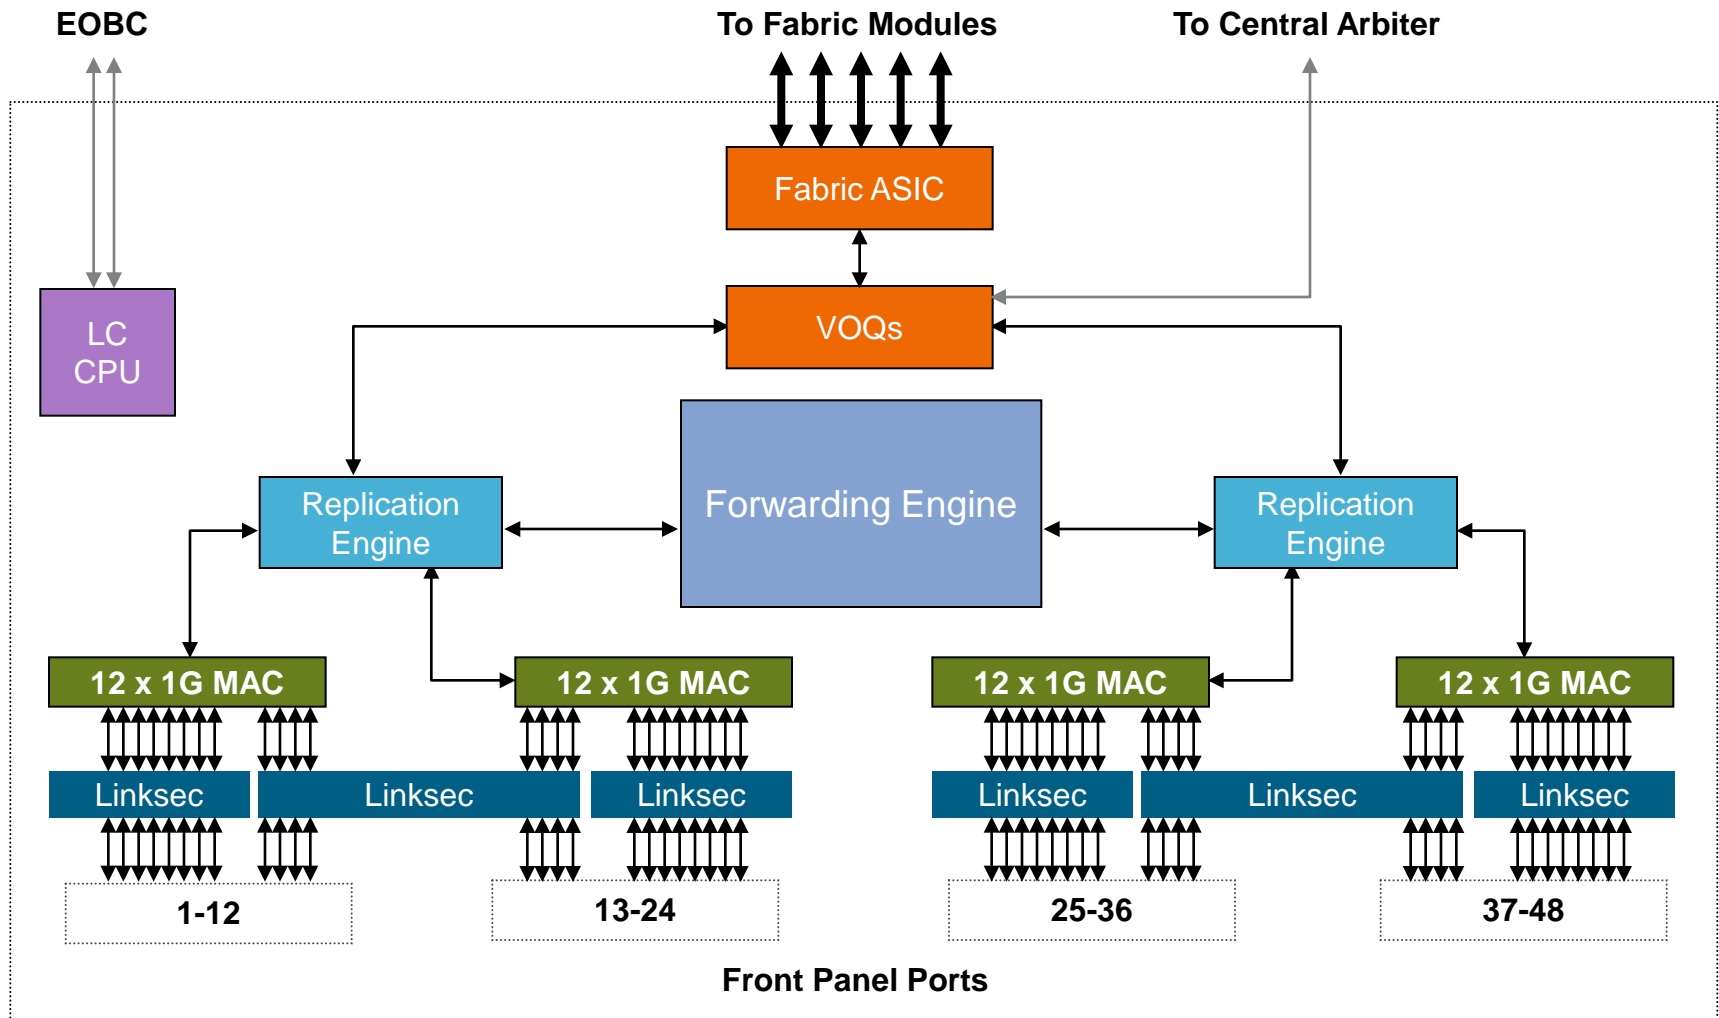
N7K-M148GS-11
Release 4.1(2) and later



N7K-M148GS-11L
Release 5.0(2) and later

48-Port 1G I/O Modules Architecture

N7K-M148GT-11, N7K-M148GS-11, N7K-M148GS-11L



Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- **Forwarding Engine Architecture**
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Forwarding Engine Hardware

- Hardware forwarding engine(s) integrated on every I/O module
- 60Mpps per forwarding engine Layer 2 bridging with hardware MAC learning
- 60Mpps per forwarding engine IPv4 and 30Mpps IPv6 unicast
- IPv4 and IPv6 multicast support (SM, SSM, bidir)
- RACL/VACL/PACLs
- Policy-based routing (PBR)
- Unicast RPF check and IP source guard
- QoS remarking and policing policies
- Ingress and egress NetFlow (full and sampled)

Hardware Table	M1 Modules	M1-XL Modules without License	M1-XL Modules with License
FIB TCAM	128K	128K	900K
Classification TCAM (ACL/QoS)	64K	64K	128K
MAC Address Table	128K	128K	128K
NetFlow Table	512K	512K	512K

“Scalable Services” License

- Forwarding engines on M1-XL I/O modules always have “XL” capacity
- Access to additional capacity controlled by presence of “Scaleable Services” license

License applies to entire system (per-chassis)

```
N7K# show license usage
```

Feature	Ins	Lic Count	Status	Expiry	Date	Comments
---------	-----	--------------	--------	--------	------	----------

SCALABLE_SERVICES_PKG	Yes	-	In use	Never		-
LAN_ADVANCED_SERVICES_PKG	Yes	-	In use	Never		-
LAN_ENTERPRISE_SERVICES_PKG	Yes	-	In use	Never		-

```
N7K#
```


Forwarding Engine Architecture

Forwarding engine chipset consists of two ASICs:

- Layer 2 Engine

- Ingress and egress SMAC/DMAC lookups

- Hardware MAC learning

- IGMP snooping and IP-based Layer 2 multicast constraint

- Layer 3 Engine

- IPv4/IPv6 Layer 3 lookups

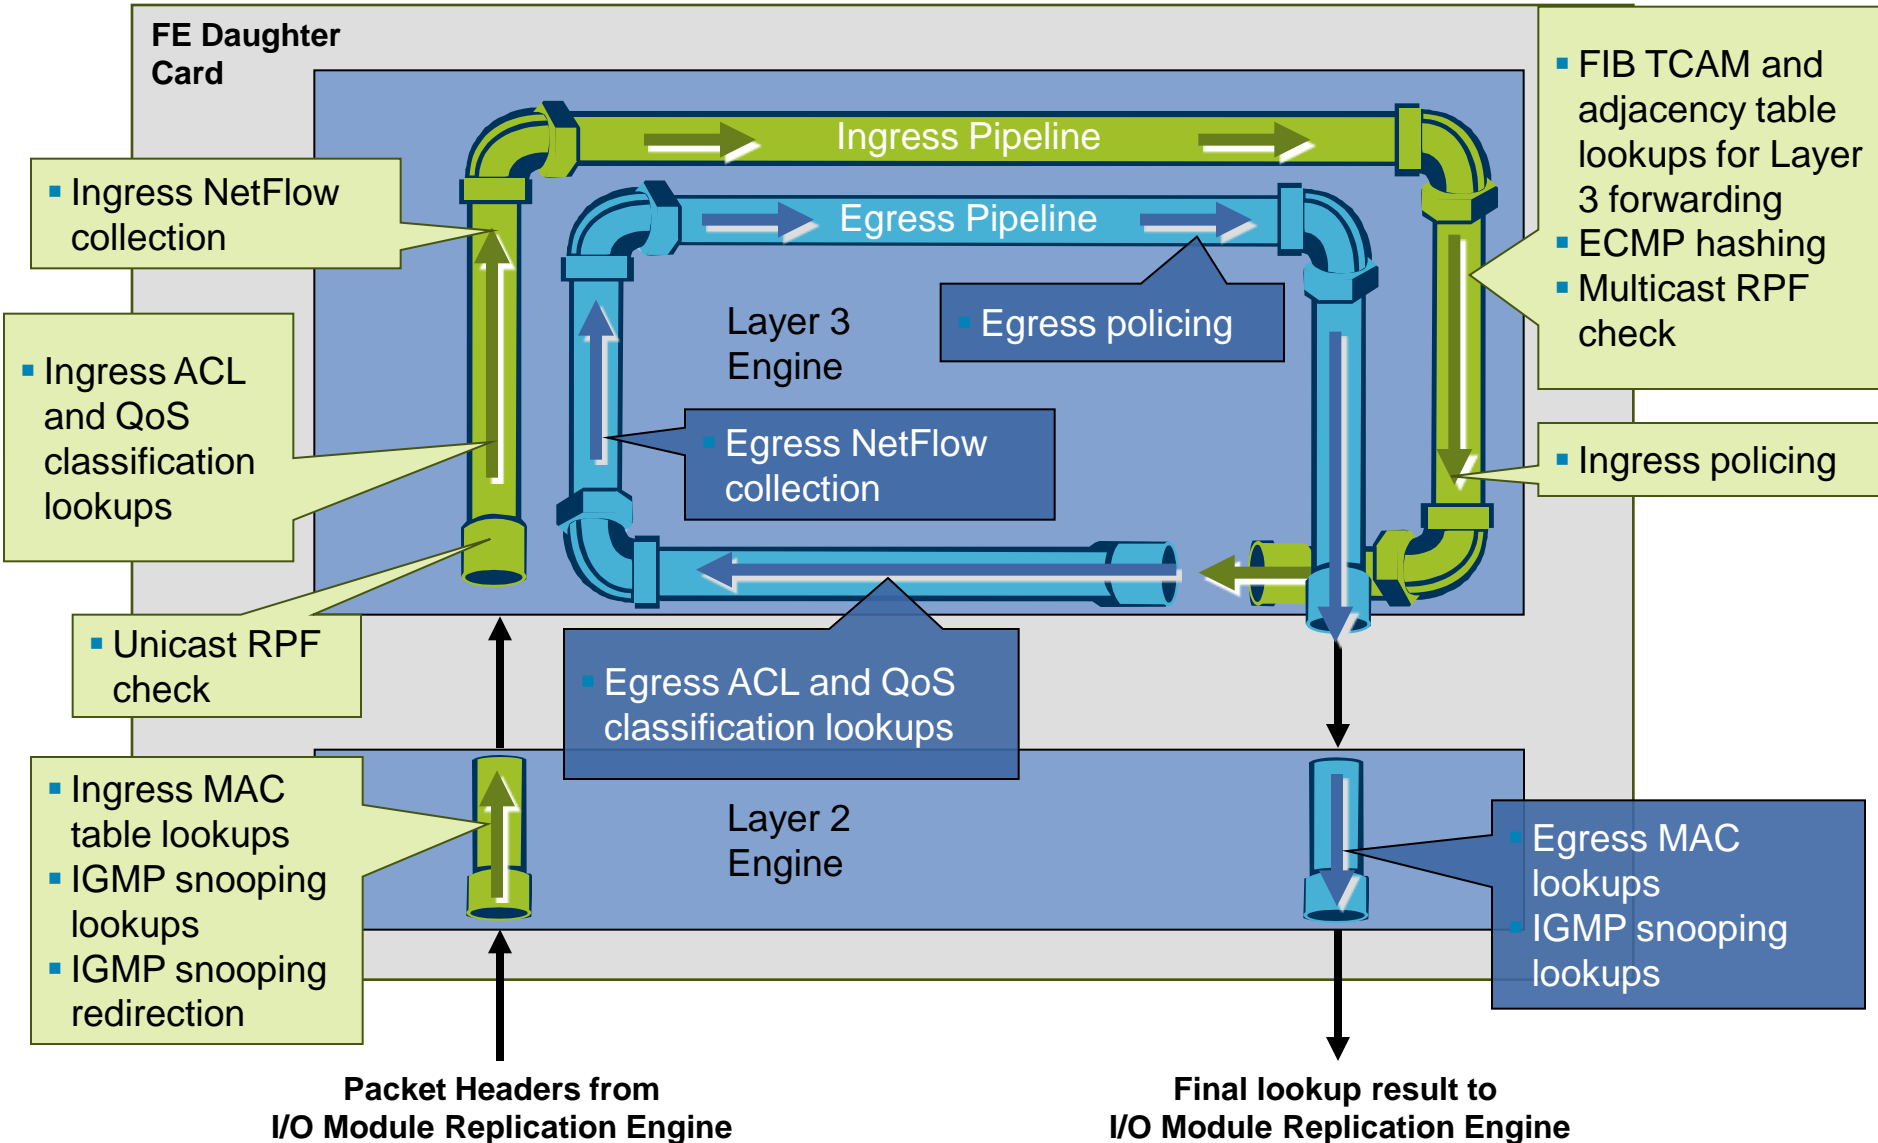
- ACL, QoS, NetFlow and other processing

- Linear, pipelined architecture—every packet subjected to both ingress and egress pipeline

- Enabling features does not affect forwarding engine performance



Forwarding Engine Pipelined Architecture




Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- **Fabric Architecture**
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Crossbar Switch Fabric Module

- Each fabric module provides 46Gbps per I/O module slot
Up to 230Gbps per slot with 5 fabric modules
- Currently shipping I/O modules do not leverage full fabric bandwidth
Maximum 80G per slot with 10G module
Future modules leverage additional available fabric bandwidth
- Access to fabric controlled using QoS-aware central arbitration with VOQ

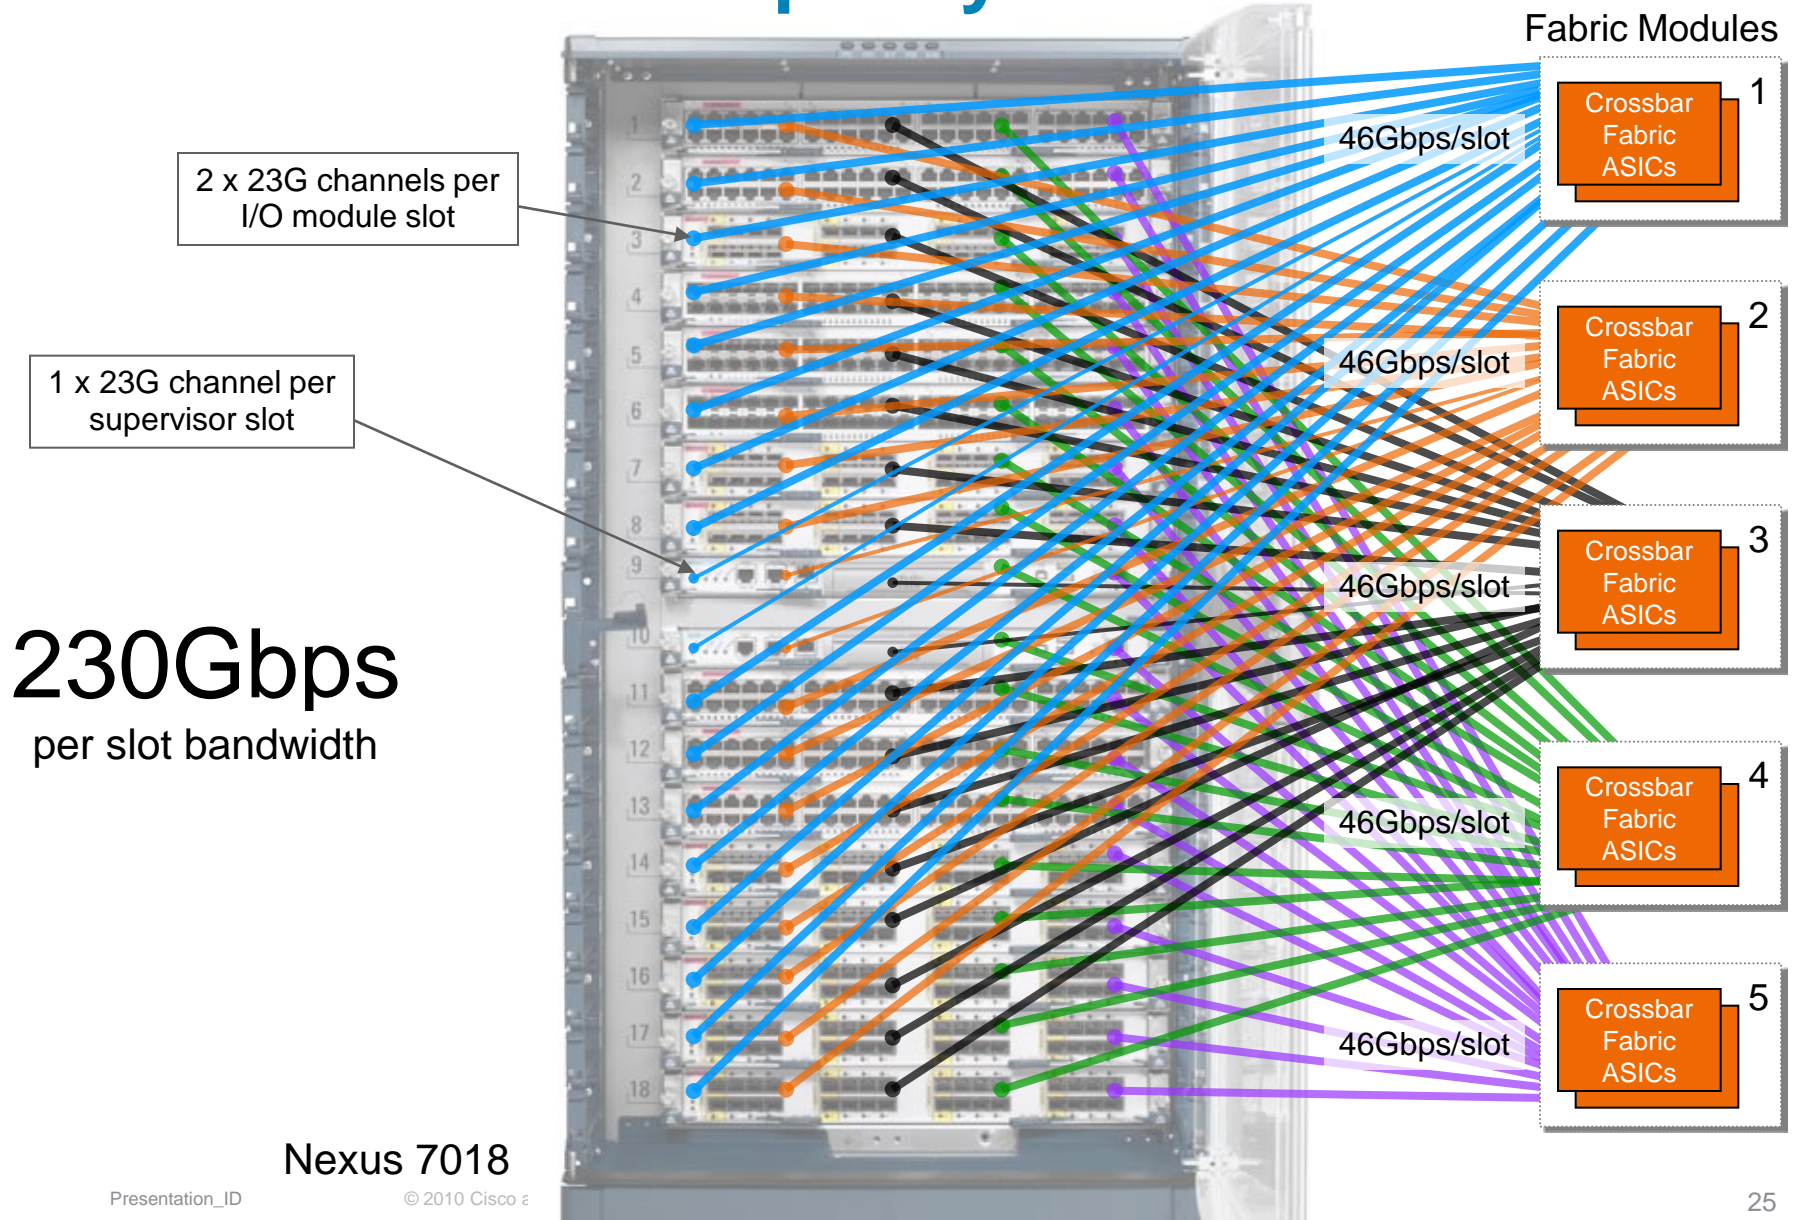
N7K-C7010-FAB-1



N7K-C7018-FAB-1



Fabric Module Capacity



I/O Module Capacity

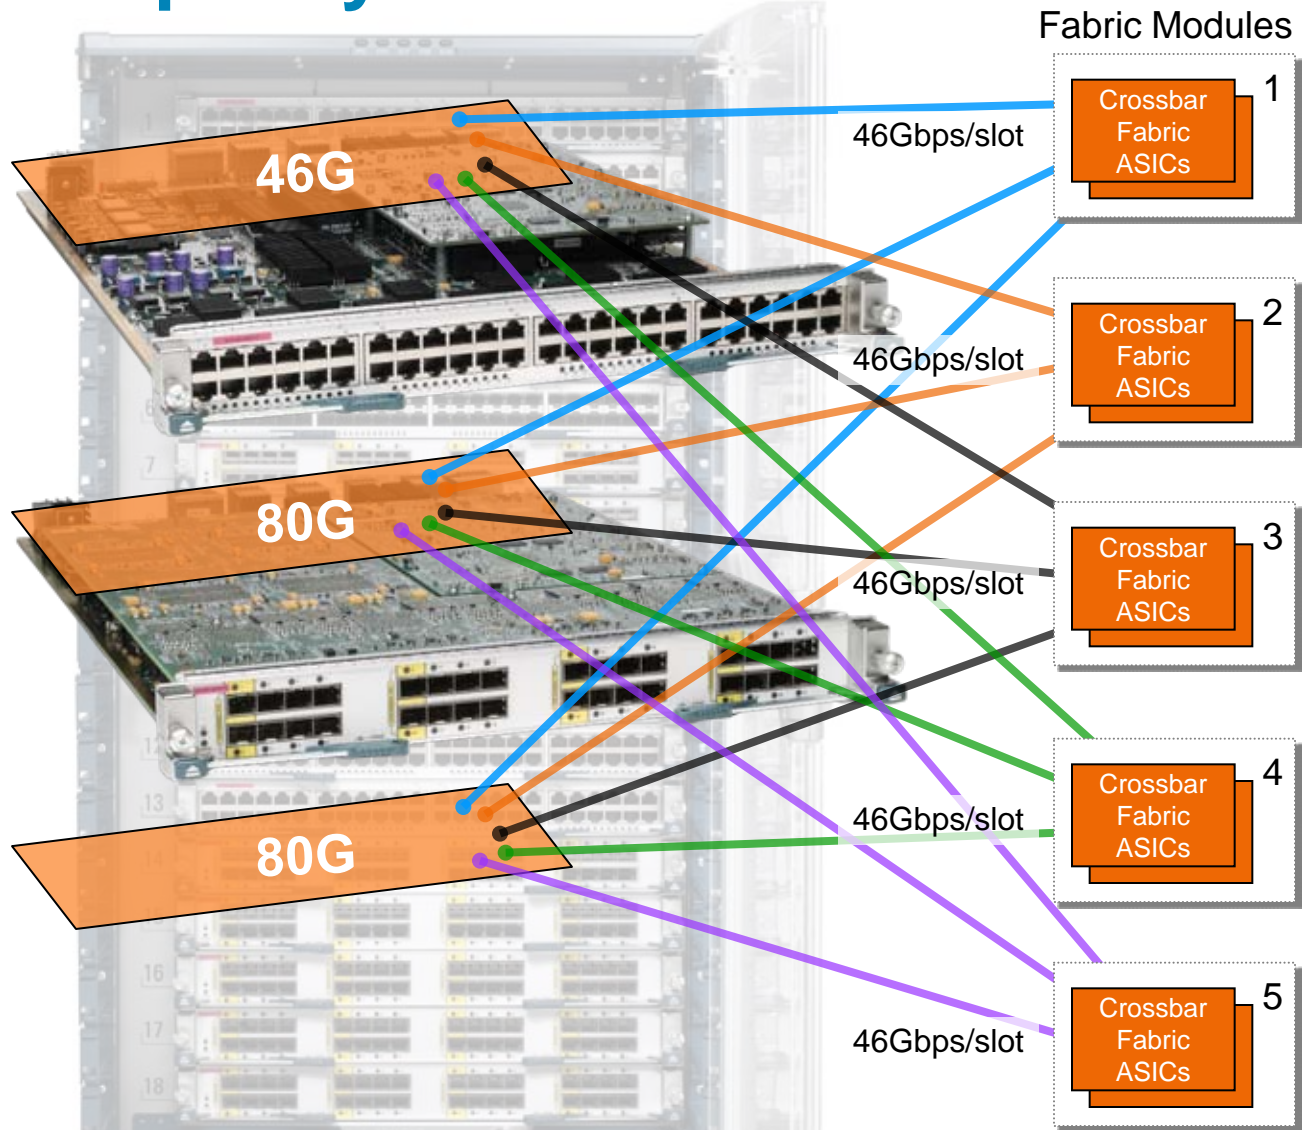
1G modules

- Require 1 fabric for full bandwidth
- Require 2 fabrics for N+1 redundancy

230Gbps

per slot bandwidth

- 4th and 5th fabric modules provide additional redundancy and future-proofing



10G modules

- Require 2 fabrics for full bandwidth
- Require 3 fabrics for N+1 redundancy

Access to Fabric Bandwidth



- Access to fabric controlled using central arbitration
 - Arbiter ASIC on supervisor engine provides fabric arbitration
- Bandwidth capacity on egress modules represented by Virtual Output Queues (VOQs) at ingress to fabric
 - I/O modules interface with arbiter to gain access to VOQs

What Are VOQs?

- Virtual Output Queues (VOQs) on ingress modules represent bandwidth capacity on egress modules
- If VOQ available on ingress to fabric, capacity exists at egress module to receive traffic from fabric
 - Central arbiter determines whether VOQ is available for a given packet
 - Bandwidth capacity represented by “credits”
 - Credits are requested by I/O modules and granted by arbiter
- VOQ is “virtual” because it represents EGRESS capacity but resides on INGRESS modules
 - It is still PHYSICAL buffer where packets are stored
- Note: VOQ is **not** equivalent to ingress or egress port buffer or queues
 - Relates ONLY to ASICs at ingress and egress to fabric

Benefits of Central Arbitration with VOQ

- Ensures priority traffic takes precedence over best-effort traffic across fabric

Four levels of priority for each VOQ destination

- Ensures fair access to bandwidth for multiple ingress ports transmitting to one egress port

Central arbiter ensures all traffic sources get appropriate access to fabric bandwidth, even with traffic sources on different modules

- Prevents congested egress ports from blocking ingress traffic destined to other ports

Mitigates head-of-line blocking by providing independent queues for individual destinations across the fabric

- In future, will provide lossless service for FCoE traffic across the fabric

Can provide strict priority and backpressure (blocking instead of dropping) for certain traffic classes, such as SAN traffic



Agenda

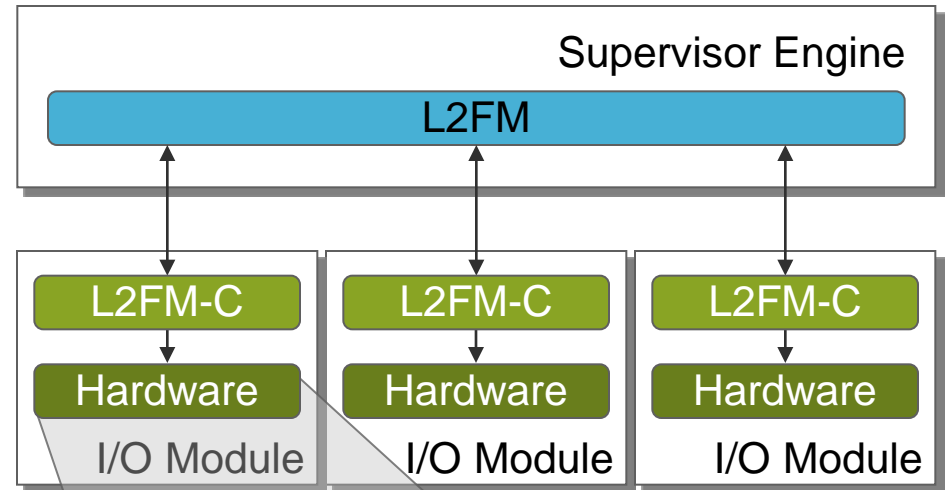
- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Layer 2 Forwarding

- MAC table is 128K entries (115K effective)
- Hardware MAC learning
 - CPU not directly involved in learning
- All modules have copy of MAC table
 - New learns communicated to other modules via hardware “flood to fabric” mechanism
 - Software process ensures continuous MAC table sync
- Spanning tree (PVRST or MST) or Virtual Port Channel (VPC) ensures loop-free Layer 2 topology

Layer 2 Forwarding Architecture

- Layer 2 Forwarding Manager (L2FM) maintains central database of MAC tables
- L2FM keeps MAC table on all forwarding engines in sync
- L2FM-Client process on I/O modules interfaces between L2FM and hardware MAC table



Hardware MAC Learning

```
n7010# sh processes cpu | egrep PID|l2fm
PID    Runtime(ms)  Invoked    uSecs  lSec  Process
 3848         1106   743970580      0     0    12fm
n7010# attach mod 9
Attaching to module 9 ...
To exit type 'exit', to abort type '$.'
Last login: Mon Apr 21 15:58:12 2009 from sup02 on
pts/0
Linux lc9 2.6.10_mvl401-pc_target #1 Fri Mar 21
23:26:28 PDT 2009 ppc GNU/Linux
module-9# sh processes cpu | egrep l2fm
1544         6396   388173      16    0.0   12fmc
module-9#
```


Hardware Layer 2 Forwarding Process

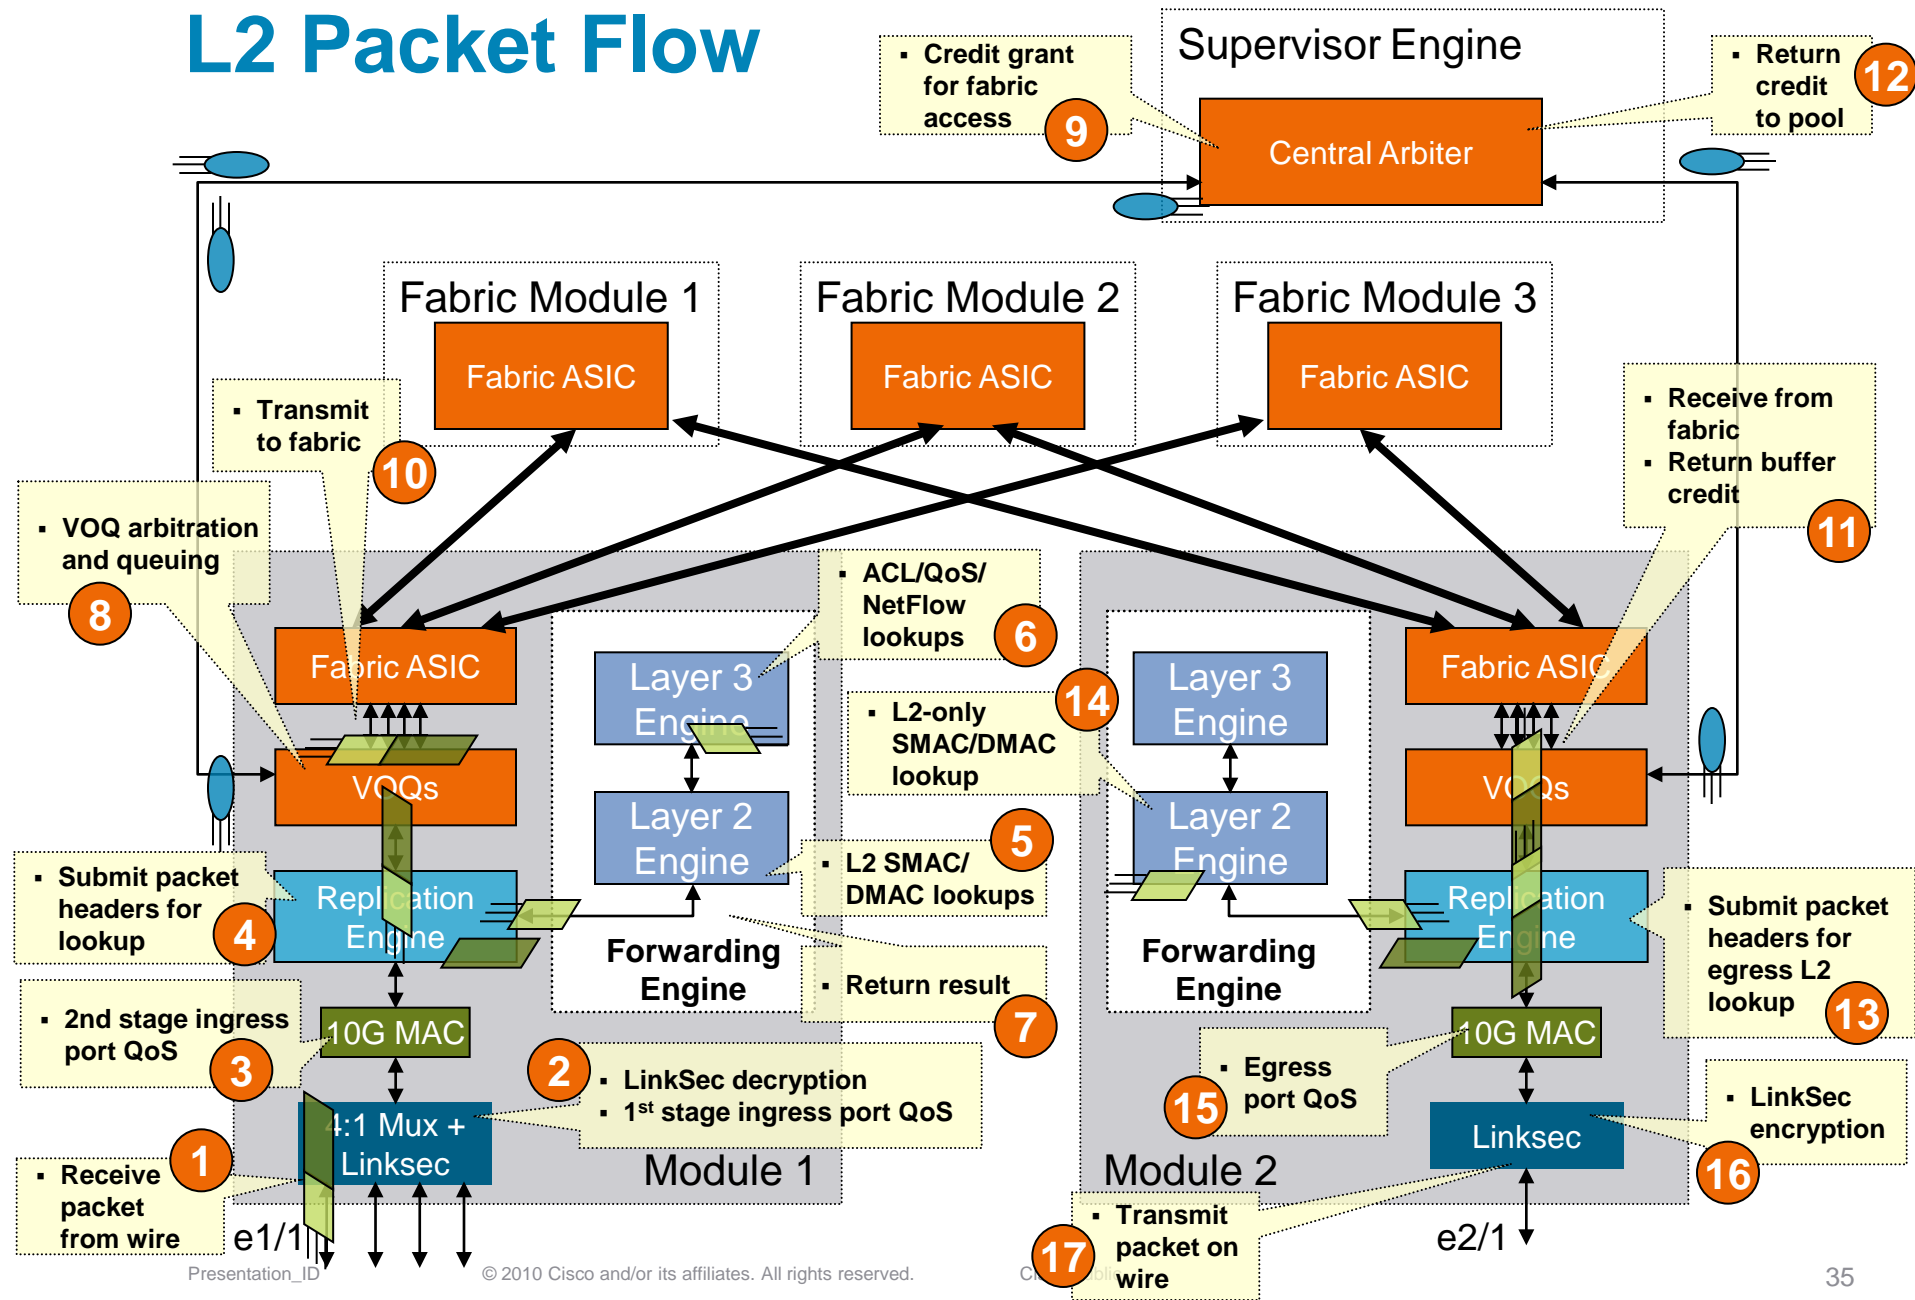
- MAC table lookup in Layer 2 Engine based on {VLAN,MAC} pairs
- Source MAC and destination MAC lookups performed for each frame

Source MAC lookup drives new learns and refreshes aging timers

Destination MAC lookup dictates outgoing switchport



L2 Packet Flow



Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- **IP Forwarding**
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

IP Forwarding

- Nexus 7000 decouples control plane and data plane
- Forwarding tables built on control plane using routing protocols or static configuration
OSPF, EIGRP, IS-IS, RIP, BGP for dynamic routing
- Tables downloaded to forwarding engine hardware for data plane forwarding



IP Forwarding Architecture

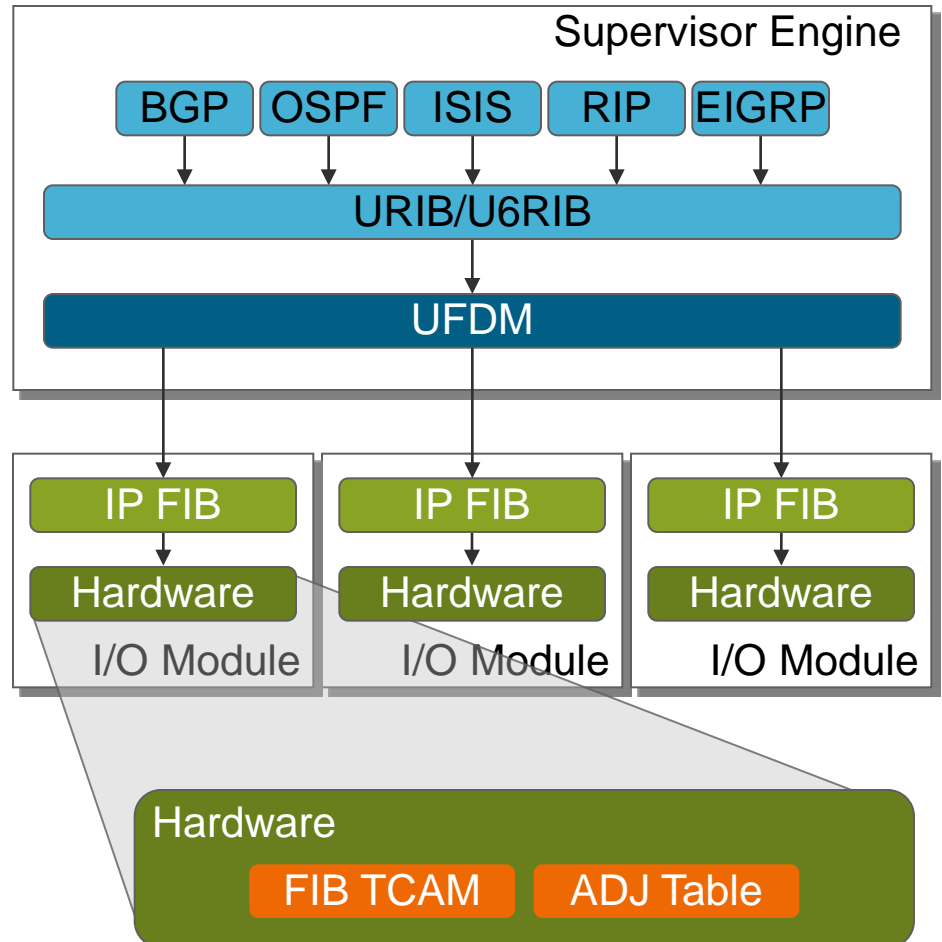
- Routing protocol processes learn routing information from neighbors
- IPv4 and IPv6 unicast RIBs calculate routing/next-hop information
- Unicast Forwarding Distribution Manager (UFDm) interfaces between URIBs on supervisor and IP FIB on I/O modules
- IP FIB process programs forwarding engine hardware on I/O modules

FIB TCAM contains IP prefixes

Adjacency table contains next-hop information

```
n7010# sh processes cpu | egrep ospf|PID
PID    Runtime(ms)  Invoked  uSecs  1Sec   Process
20944      93    33386880      0      0    ospf
n7010# sh processes cpu | egrep u.?rib
3573      117    44722390      0      0    u6rib
3574      150    34200830      0      0    urib
n7010# sh processes cpu | egrep ufdm
3836     1272    743933460      0      0    ufdm

module-9# sh processes cpu | egrep fib
1534     80042    330725    242    0.0    ipfib
module-9#
```

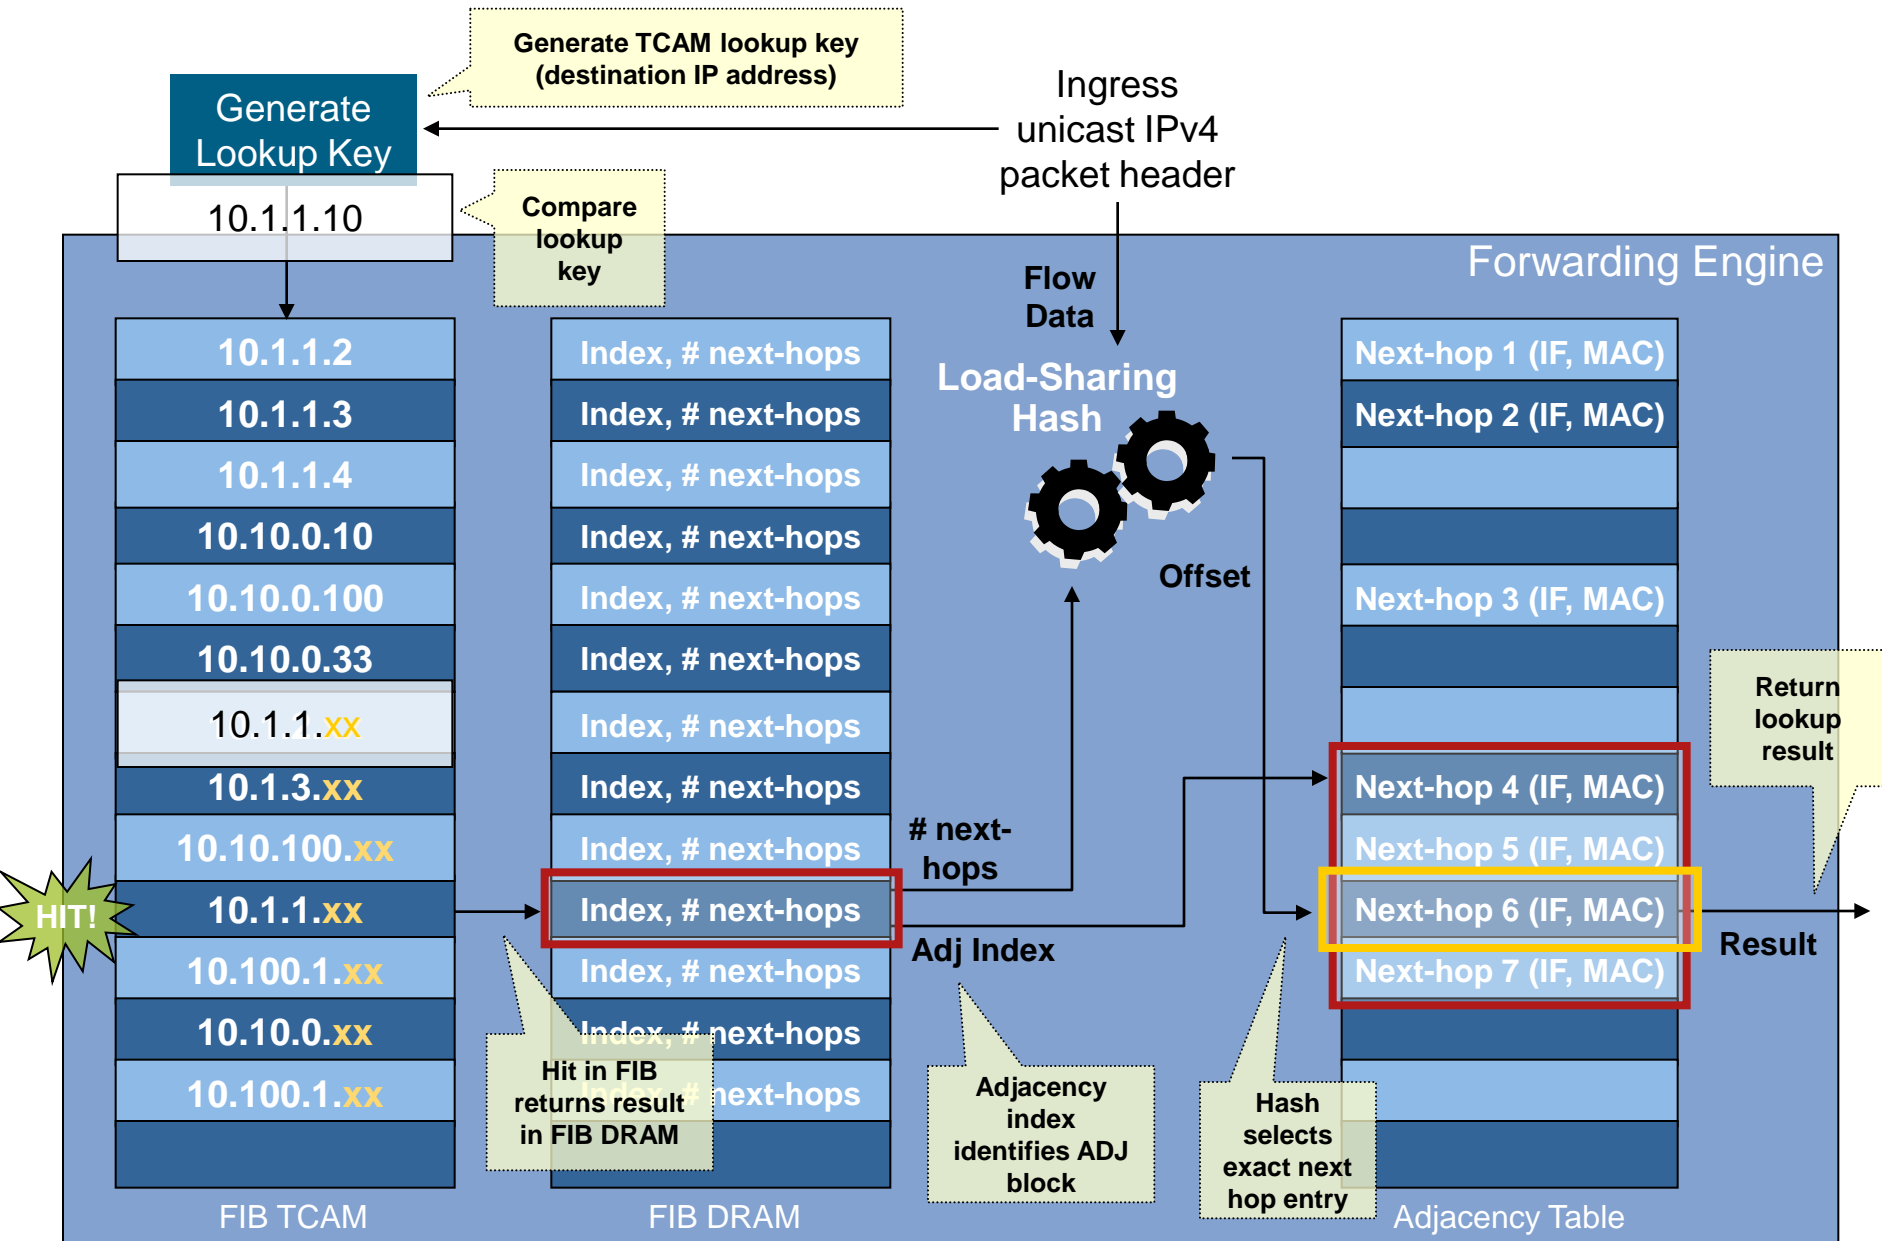


Hardware IP Forwarding Process

- FIB TCAM lookup based on destination prefix (longest-match)
- FIB “hit” returns adjacency, adjacency contains rewrite information (next-hop)
- Pipelined forwarding engine architecture also performs ACL, QoS, and NetFlow lookups, affecting final forwarding result



IPv4 FIB TCAM Lookup



“Routing” vs. “Forwarding”

- “Routing” information refers to unicast RIB contents in supervisor control plane
- “Forwarding” information refers to FIB contents at I/O module

Displaying Routing and Forwarding Information

- `show routing [ipv4|ipv6] [<prefix>] [vrf <vrf>]`

Displays software routing (URIB) information

Can also use traditional `show ip route` command

- `show forwarding [ipv4|ipv6] route module <mod> [vrf <vrf>]`

Displays hardware forwarding (FIB) information on per-module basis

- `show forwarding adjacency module <mod>`

Displays hardware adjacency table information on per-module basis

Displaying Routing and Forwarding Information (Cont.)

```
n7010# sh routing ipv4 10.100.7.0/24
```

```
IP Route Table for VRF "default"
```

```
10.100.7.0/24, 1 ucast next-hops, 0 mcast next-hops
```

```
    *via 10.1.2.2, Ethernet9/2, [110/5], 00:02:30, ospf-1, type-1
```

```
n7010# show forwarding ipv4 route 10.100.7.0/24 module 9
```

```
IPv4 routes for table default/base
```

```
-----+-----+-----
Prefix          | Next-hop          | Interface
-----+-----+-----
10.100.7.0/24   | 10.1.2.2          | Ethernet9/2
```

```
n7010# show forwarding adjacency 10.1.2.2 module 9
```


```
IPv4 adjacency information, adjacency count 1
```

```
next-hop          rewrite info      interface
-----
```



ECMP Load Sharing

- Up to 16 hardware load-sharing paths per prefix
- Use **maximum-paths** command in routing protocols to control number of load-sharing paths
- Load-sharing is per-IP flow or per-packet
 - Use caution with per-packet load-balancing!
- Configure load-sharing hash options with global **ip load-sharing** command:
 - Source and Destination IP addresses
 - Source and Destination IP addresses plus L4 ports (default)
 - Destination IP address and L4 port
- Additional randomized number added to hash prevents polarization
 - Automatically generated or user configurable value
- Configure per-packet load-sharing with interface **ip load-sharing per-packet** command
 - Ingress interface determines if load-sharing is per-flow or per-packet!



ECMP Prefix Entry Example



```
n7010# sh routing ipv4 10.200.0.0
```

```
IP Route Table for VRF "default"
```

```
10.200.0.0/16, 2 ucast next-hops, 0 mcast next-hops
```

```
    *via 10.1.1.2, Ethernet9/1, [110/5], 00:03:33, ospf-1, inter
```

```
    *via 10.1.2.2, Ethernet9/2, [110/5], 00:00:13, ospf-1, inter
```

```
n7010# sh forwarding ipv4 route 10.200.0.0 module 9
```

```
IPv4 routes for table default/base
```

Prefix	Next-hop	Interface
10.200.0.0/16	10.1.1.2	Ethernet9/1
	10.1.2.2	Ethernet9/2

```
n7010#
```

Identifying the ECMP Path for a Flow

```
show routing [ipv4|ipv6] hash <sip>  
<dip> [<sport> <dport>] [vrf <vrf>]
```

```
n7010# sh routing hash 192.168.44.12 10.200.71.188
```

```
Load-share parameters used for software forwarding:
```

```
load-share type: 1
```

```
Randomizing seed (network order): 0xebae8b9a
```

```
Hash for VRF "default"
```

```
Hashing to path *10.1.2.2 (hash: 0x29), for route:
```

```
10.200.0.0/16, 2 ucast next-hops, 0 mcast next-hops
```

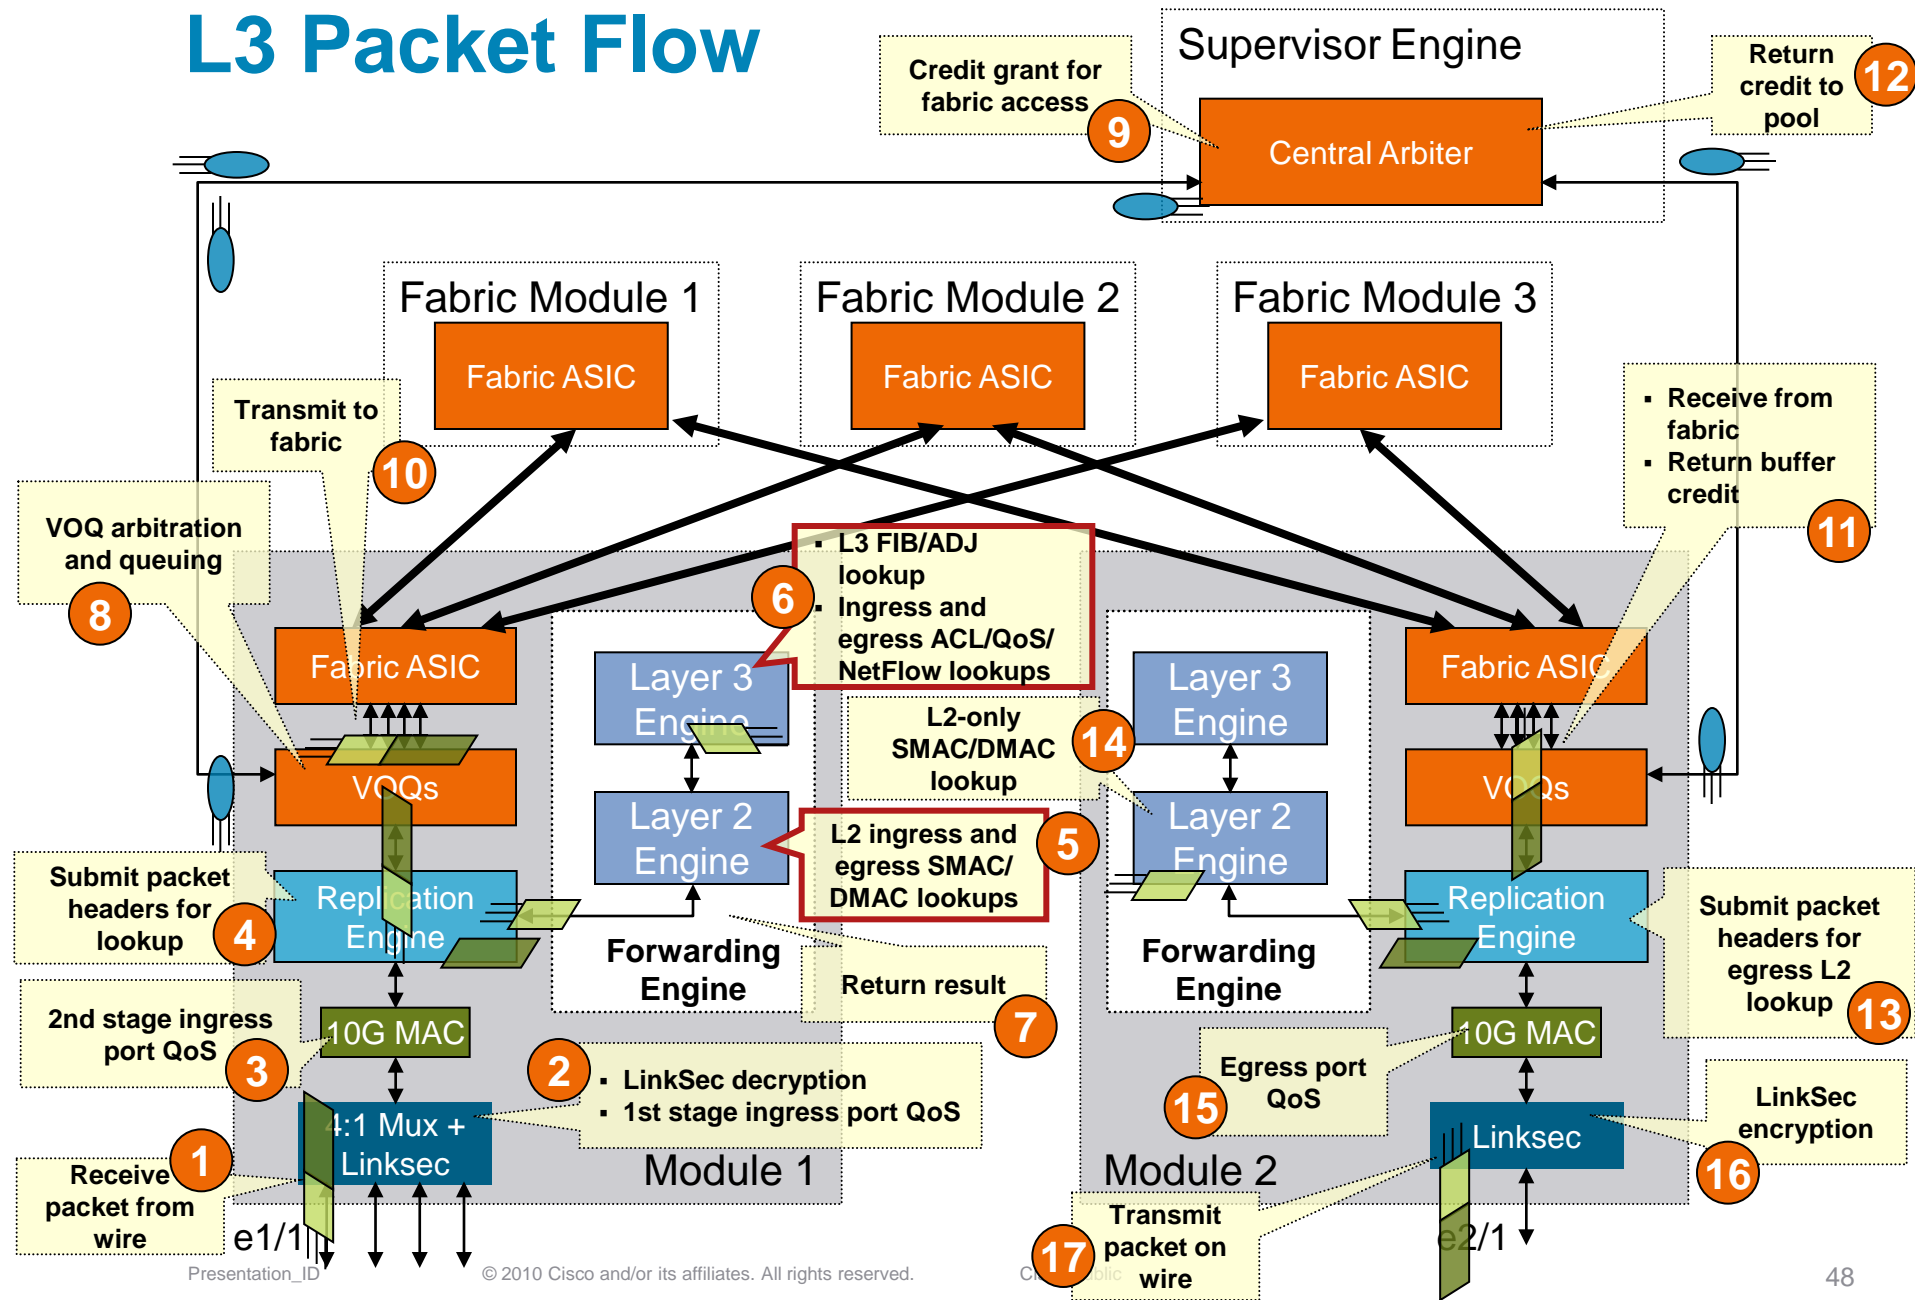
```
*via 10.1.1.2, Ethernet9/1, [110/5], 00:14:18, ospf-1, inter
```

```
*via 10.1.2.2, Ethernet9/2, [110/5], 00:10:58, ospf-1, inter
```

```
n7010#
```

Same hash algorithm applies to both hardware and software forwarding

L3 Packet Flow



Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- **IP Multicast Forwarding**
- ACLs
- QoS
- NetFlow

IP Multicast Forwarding

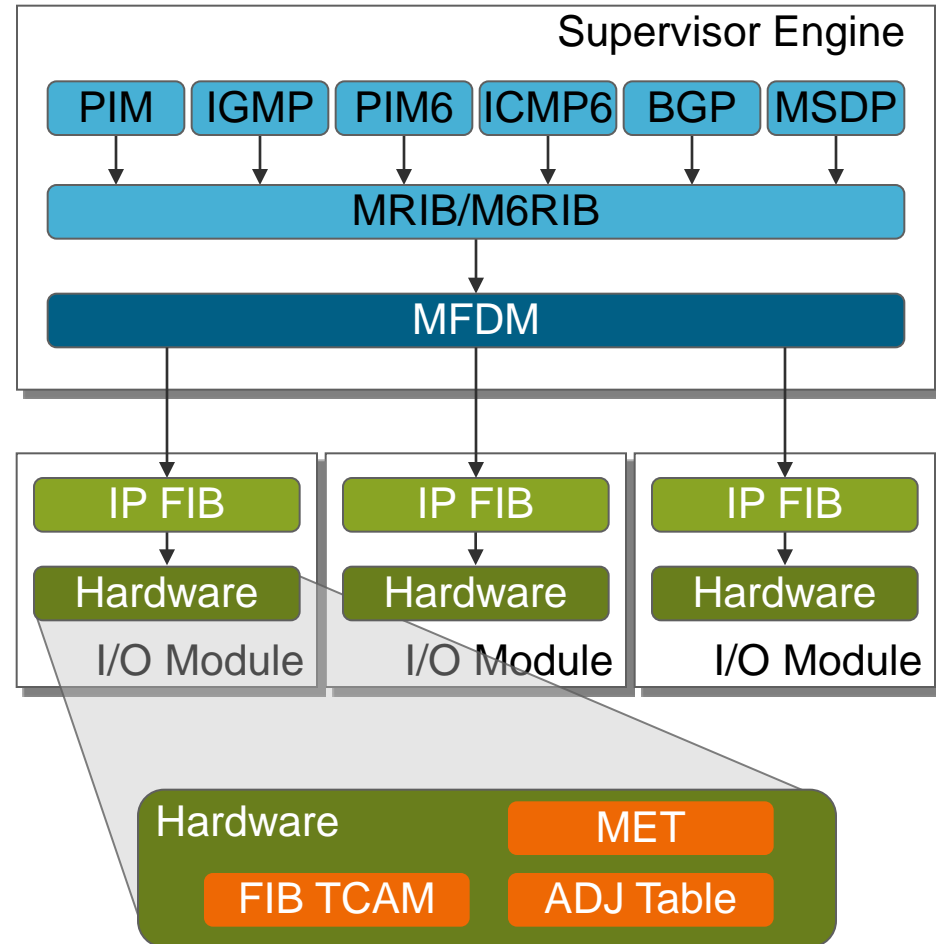
- Forwarding tables built on control plane using multicast protocols
PIM-SM, PIM-SSM, PIM-Bidir, IGMP, MLD
- Tables downloaded to:
Forwarding engine hardware for data plane forwarding
Replication engines for data plane packet replication

IP Multicast Forwarding Architecture

- Multicast routing processes learn routing information from neighbors/hosts
- IPv4 and IPv6 multicast RIBs calculate multicast routing/RP/RPF/OIL information
- Multicast Forwarding Distribution Manager (MFDM) interfaces between MRIBs on supervisor and IP FIB on I/O modules
- IP FIB process programs hardware:
 - FIB TCAM
 - Adjacency table
 - Multicast Expansion Table (MET)

```
n7010# sh processes cpu | egrep pim|igmp|PID
PID      Runtime(ms)   Invoked    uSecs   lSec   Process
  3842          109    32911620         0      0    pim
  3850          133    33279940         0      0    igmp
n7010# sh processes cpu | egrep m.?rib
  3843          177    33436550         0      0    mrib
  3847          115    47169180         0      0    m6rib
n7010# sh processes cpu | egrep mfdm
  3846         2442    743581240         0      0    mfdm
```

```
module-9# sh processes cpu | egrep fib
 1534         80153    330725         242    0.0    ipfib
module-9#
```



Hardware Programming

IP FIB process on I/O modules programs hardware:

- FIB TCAM

Part of Layer 3 Engine ASIC on forwarding engine

Consists of (S,G) and (*,G) entries as well as RPF interface

- Adjacency Table (ADJ)


Part of Layer 3 Engine ASIC on forwarding engine

Contains MET indexes

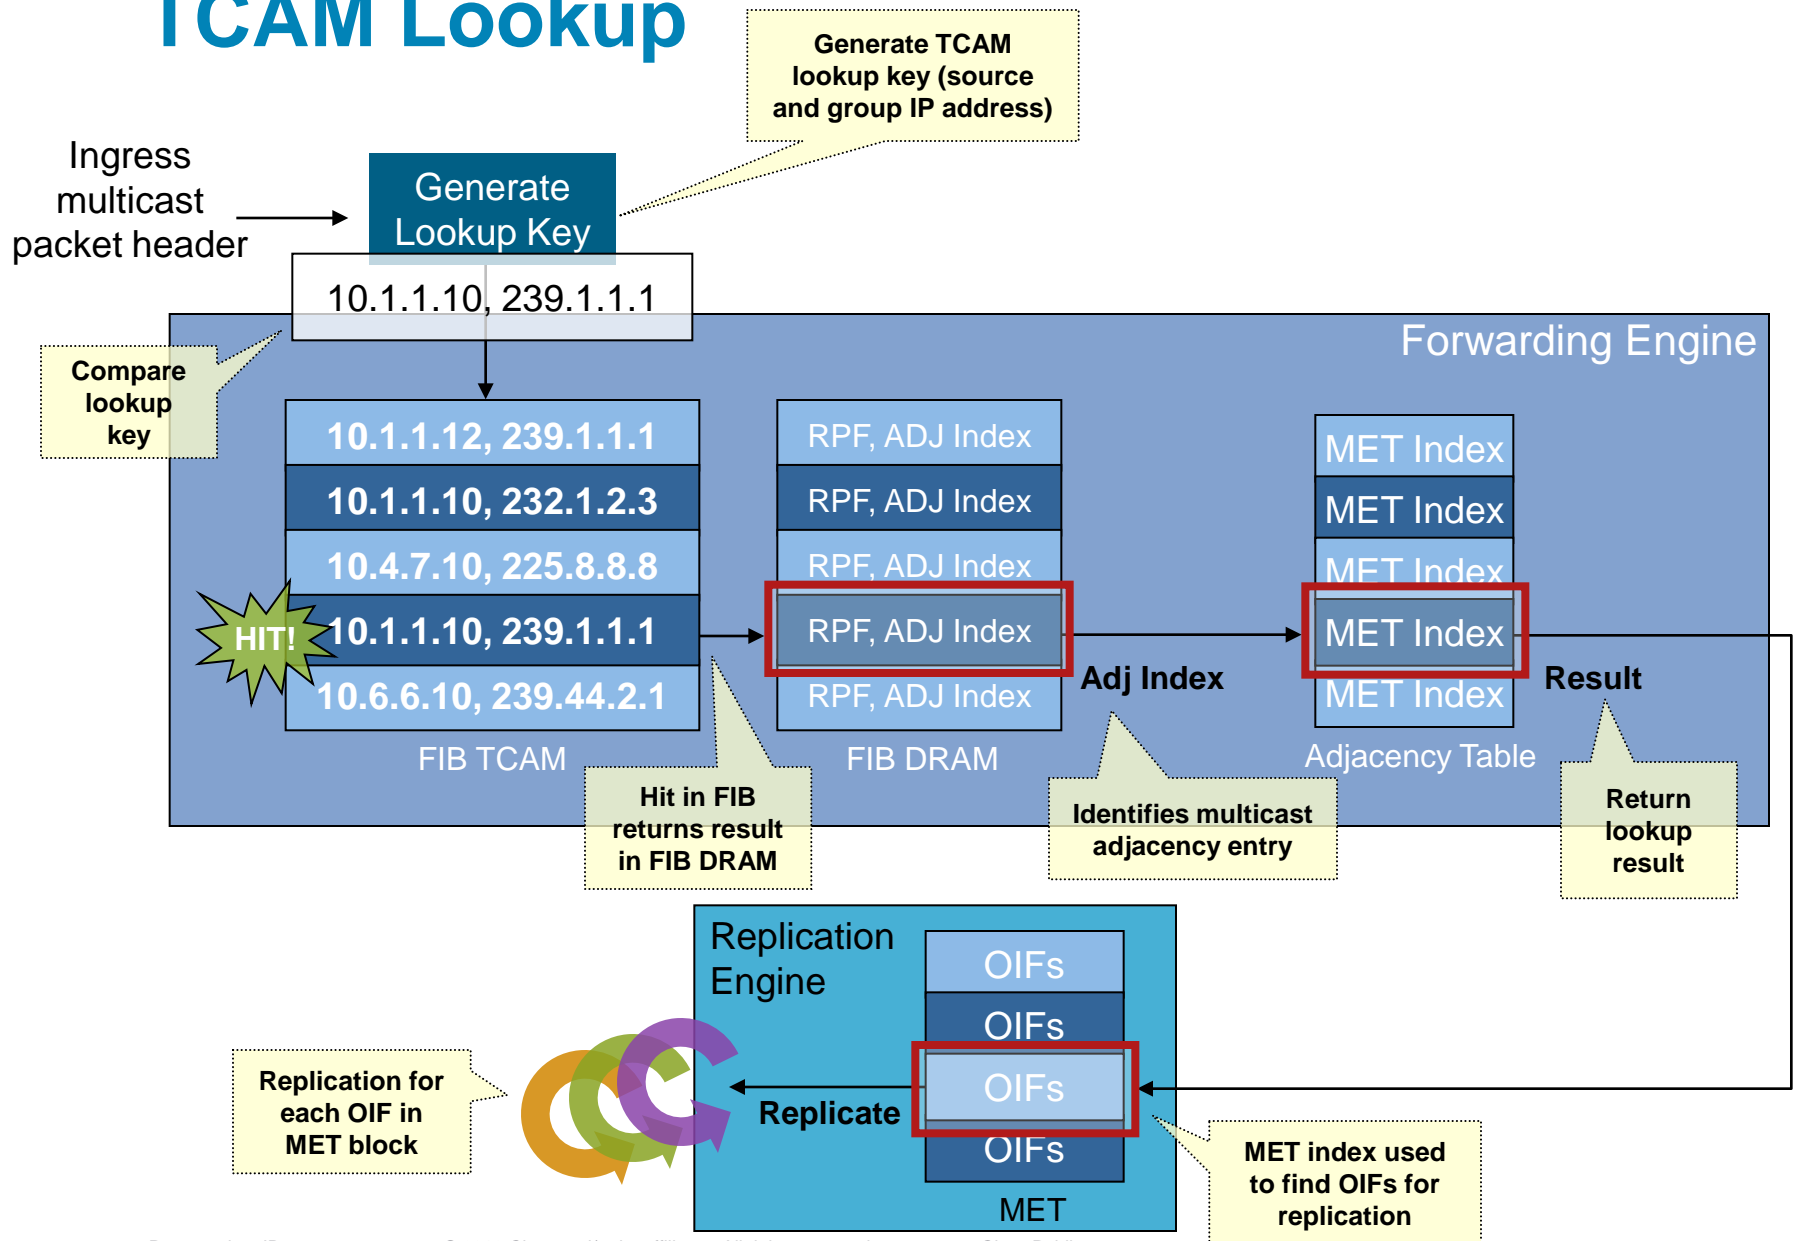
- Multicast Expansion Table (MET)

Part of replication engine ASIC on I/O modules

Contains output interface lists (OILs), i.e., lists of interfaces requiring replication



Multicast FIB TCAM Lookup



Displaying Multicast Routing and Forwarding Information

- `show routing [ipv4|ipv6] multicast [vrf <vrf>] [<source-ip>] [<group-ip>] [summary]`

Displays software multicast routing (MRIB) information

Can also use traditional `show ip mroute` command

- `show forwarding [ipv4|ipv6] multicast route [source <ip>] [group <ip>] [vrf <vrf>] module <mod>`

Displays hardware multicast forwarding (FIB) information on per-module basis

Displaying Multicast Routing and Forwarding Information (Cont)



```
n7010# sh routing multicast 10.1.1.2 239.1.1.1
```

```
IP Multicast Routing Table for VRF "default"
```

```
(10.1.1.2/32, 239.1.1.1/32), uptime: 00:40:31, ip mrib pim
  Incoming interface: Ethernet9/1, RPF nbr: 10.1.1.2, internal
  Outgoing interface list: (count: 2)
    Ethernet9/17, uptime: 00:05:57, mrib
    Ethernet9/2, uptime: 00:06:12, mrib
```

```
n7010# sh routing multicast 239.1.1.1 summary
```

```
IP Multicast Routing Table for VRF "default"
```

```
Total number of routes: 202
Total number of (*,G) routes: 1
Total number of (S,G) routes: 200
Total number of (*,G-prefix) routes: 1
Group count: 1, average sources per group: 200.0
```

```
Group: 239.1.1.1/32, Source count: 200
```

Source	packets	bytes	aps	pps	bit-rate	oifs
(*,G)	767	84370	110	0	0 bps	2
10.1.1.2	9917158	1269395810	127	4227	4 mbps	2
10.1.1.3	9917143	1269393890	127	4227	4 mbps	2
10.1.1.4	9917127	1269391824	127	4227	4 mbps	2

```
<...>
```

Displaying Multicast Routing and Forwarding Information (Cont.)



```
n7010# sh forwarding ipv4 multicast route group 239.1.1.1 source 10.1.1.2 module 9
```

```
(10.1.1.2/32, 239.1.1.1/32), RPF Interface: Ethernet9/1, flags:
```

```
Received Packets: 10677845 Bytes: 1366764160
```

```
Number of Outgoing Interfaces: 2
```

```
Outgoing Interface List Index: 15
```


```
Ethernet9/2 Outgoing Packets:432490865 Bytes:55358830720
```

```
Ethernet9/17 Outgoing Packets:419538767 Bytes:53700962176
```

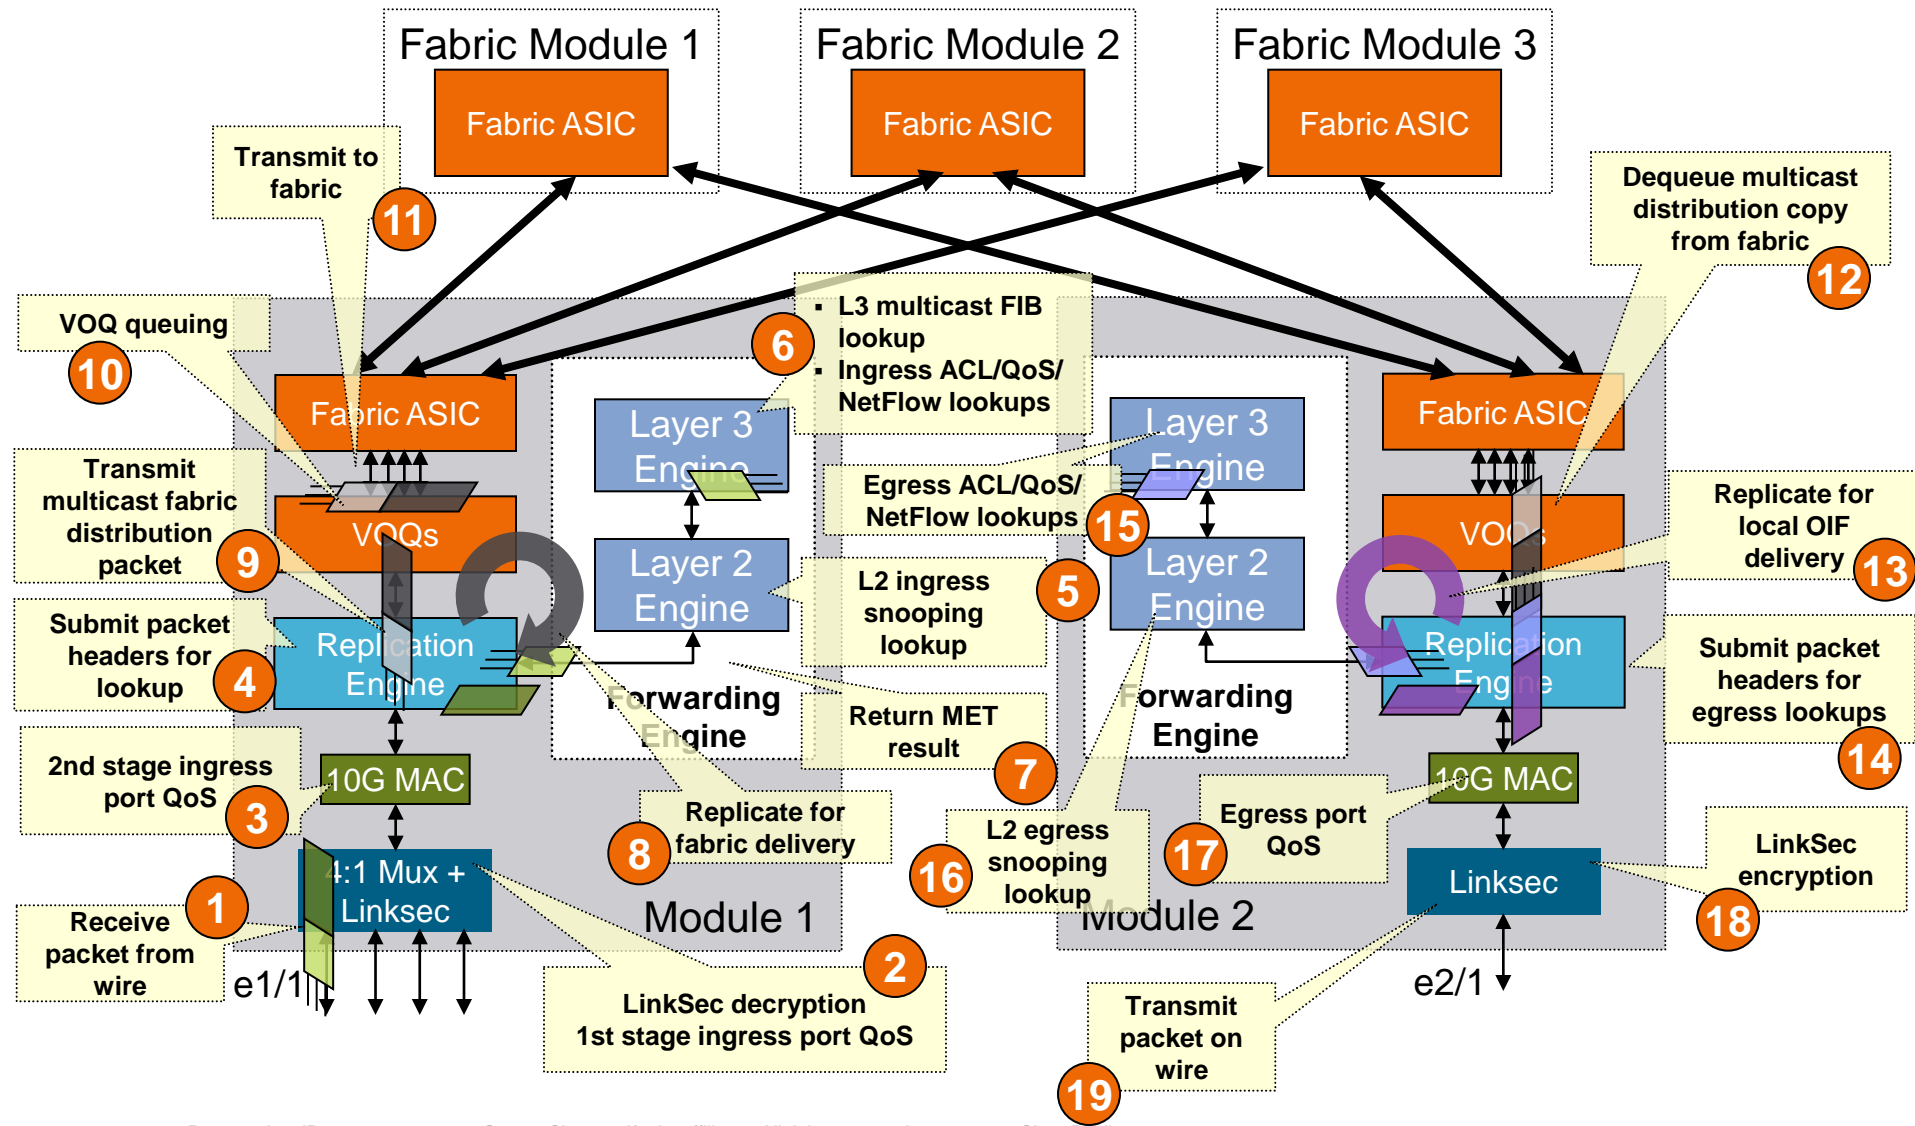
```
n7010#
```

Egress Replication

- Distributes multicast replication load among replication engines of all I/O modules with OIFs
- Input packets get lookup on ingress forwarding engine
- For OIFs on ingress module, ingress replication engine performs the replication
- For OIFs on other modules, ingress replication engine replicates a single copy of packet over fabric to those egress modules
- Each egress forwarding engine performs lookup to drive replication
- Replication engine on egress module performs replication for local OIFs



L3 Multicast Packet Flow



Unicast vs. Multicast on Fabric

Fabric consists of two parallel “fabric planes”:

- Unicast traffic

- Centrally arbitrated

- Round-robin load balanced over available fabric channels

- Multicast traffic

- Locally arbitrated

- Load balanced over available fabric channels using hash

Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- **ACLs**
- QoS
- NetFlow

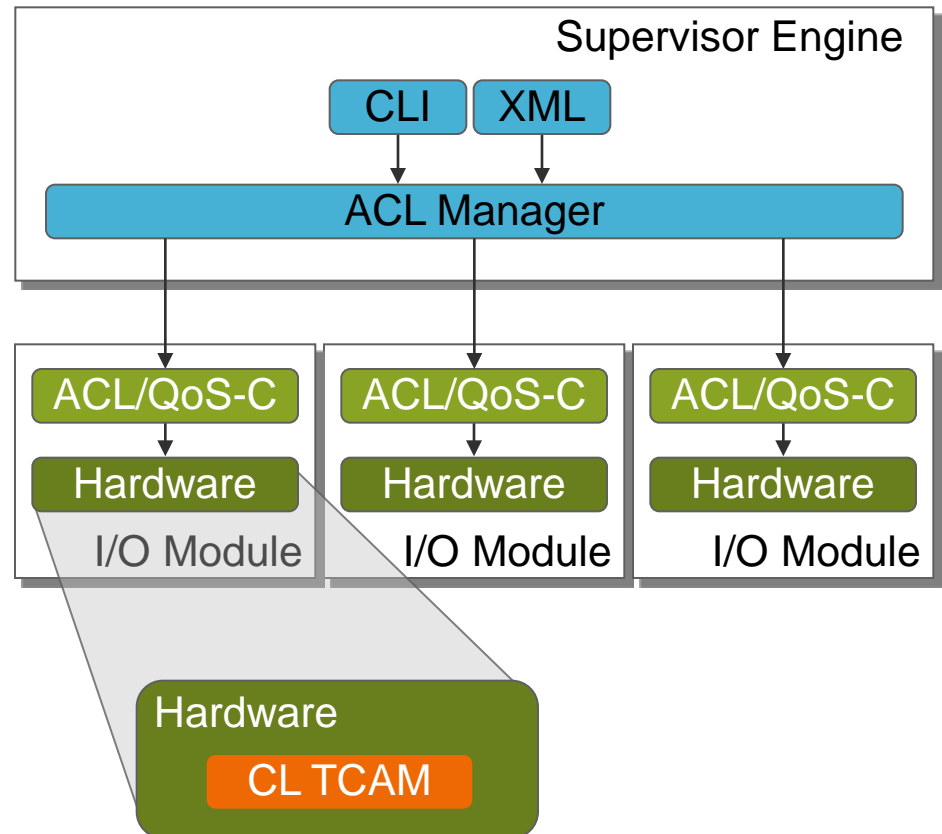
Security ACLs

- Enforce security policies based on Layer 2, Layer 3, and Layer 4 information
- Classification TCAM (CL TCAM) provides ACL lookups in forwarding engine
 - 64K hardware entries
- Router ACL (RACL)—Enforced for all traffic crossing a Layer 3 interface in a specified direction
 - IPv4, ARP RACLs supported
- VLAN ACLs (VACLs)—Enforced for all traffic in the VLAN
 - IPv4, MAC VACLs supported
- Port ACLs (PACLs)—Enforced for all traffic input on a Layer 2 interface
 - IPv4, MAC PACLs supported



ACL Architecture

- ACL manager receives policy via configuration
- ACL manager distributes policies to ACL/QoS Clients on I/O modules
- Clients perform ACL merge and program ACEs in Classification (CL) TCAM in forwarding engines



```
n7010# sh processes cpu | egrep aclmgr|PID
PID    Runtime(ms)   Invoked    uSecs  lSec  Process
3589           1662    516430000      0     0  aclmgr
```

```
module-9# sh processes cpu | egrep aclqos
1532           9885     671437      14    0.0  aclqos
module-9#
```

ACL CL TCAM Lookup

Packet header:
SIP: 10.1.1.1
DIP: 10.2.2.2
Protocol: TCP
SPORT: 33992
DPORT: 80

Generate TCAM lookup key (source/dest IPs, protocol, L4 ports, etc.)

Generate Lookup Key

SIP | DIP | Protocol | SPORT | DPORT

Compare lookup key

10.1.1.1 | 10.2.2.2 | 06 | 84C8 | 0050

Forwarding Engine

X="Mask"

X-XXXXXXX	10.2.2.2	XX	XXX	XXX	X
XXXXXXXX	10.1.68.44	XX	XXX	XXX	
XXXXXXXX	10.33.2.25	XX	XXX	XXX	
XXXXXXXX	XXXXXXXX	06	XXX	0050	6
XXXXXXXX	XXXXXXXX	06	XXX	0017	
XXXXXXXX	XXXXXXXX	11	XXX	0202	
XXXXXXXX	XXXXXXXX	06	XXX	0050	
XXXXXXXX	XXXXXXXX	11	XXX	00A1	

HIT!

CL TCAM

Hit in CL TCAM returns result in CL SRAM

Permit

Deny

Deny

Permit

Deny

Deny

Permit

Permit

Result

Return lookup result

Result affects final packet handling

Security ACL

ip access-list example

permit ip any host 10.1.2.100

deny ip any host 10.1.68.44

deny ip any host 10.33.2.25

permit tcp any any eq 22

deny tcp any any eq 23

deny udp any any eq 514

permit tcp any any eq 80

permit udp any any eq 161

Displaying Classification Resources

- `show hardware access-list resource utilization module <mod>`

```
n7010# sh hardware access-list resource utilization module 9
```

Hardware Modules	Used	Free	Percent Utilization

Tcam 0, Bank 0	1	16383	0.000
Tcam 0, Bank 1	4121	12263	25.000
Tcam 1, Bank 0	4013	12371	24.000
Tcam 1, Bank 1	4078	12306	24.000
LOU	2	102	1.000
Both LOU Operands	0		
Single LOU Operands	2		
TCP Flags	0	16	0.000
Protocol CAM	4	3	57.000
Mac Etype/Proto CAM	0	14	0.000
Non L4op labels, Tcam 0	3	6140	0.000
Non L4op labels, Tcam 1	3	6140	0.000
L4 op labels, Tcam 0	0	2047	0.000
L4 op labels, Tcam 1	1	2046	0.000

Agenda

- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- NetFlow

Quality of Service

- Comprehensive LAN QoS feature set
- Ingress and egress queuing and scheduling
Applied in I/O module port ASICs
- Ingress and egress mutation, classification, marking, policing
Applied in I/O module forwarding engines
- All configuration through Modular QoS CLI (MQC)
All QoS features applied using class-maps/policy-maps/service-policies

QoS Architecture

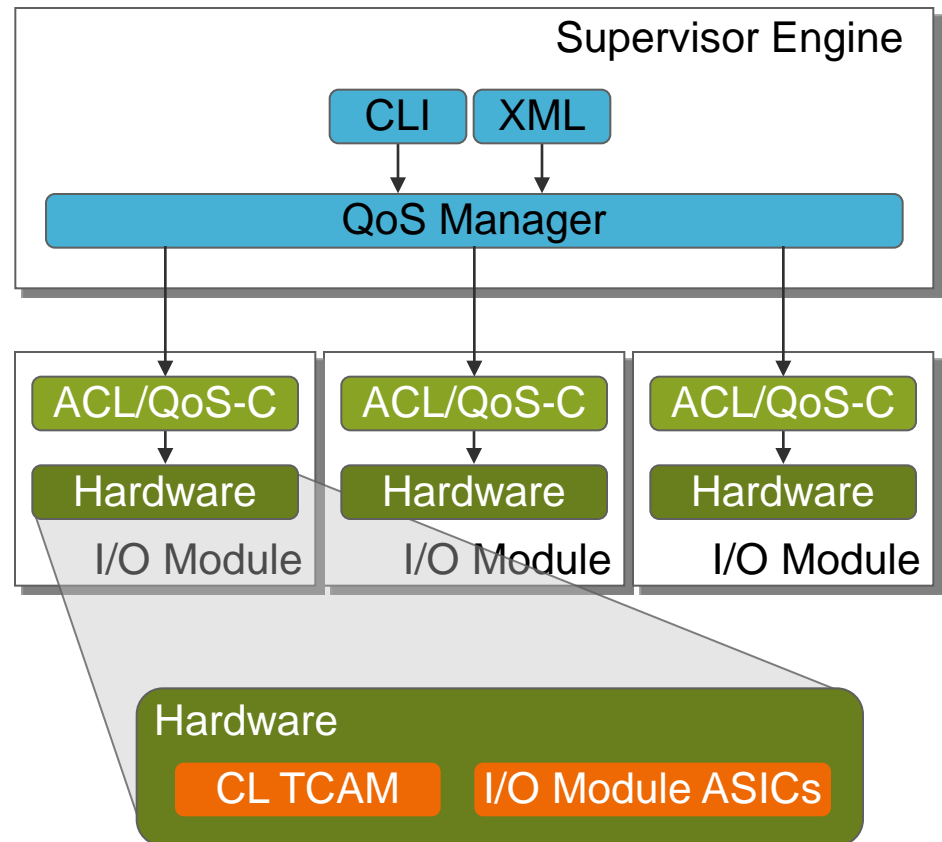
- QoS manager receives policy via configuration
- QoS manager distributes policies to ACL/QoS Clients on I/O modules
- Clients perform ACL merge and program hardware:

ACEs in Classification (CL)
TCAM in forwarding engines

Queuing policies in I/O module
port ASICs

```
n7010# sh processes cpu | egrep qos|PID
PID    Runtime(ms)   Invoked    uSecs   lSec    Process
 3849         1074    66946870         0      0    ipqosmgr

module-9# sh processes cpu | egrep aclqos
1532         9885    671437         14    0.0    aclqos
module-9#
```

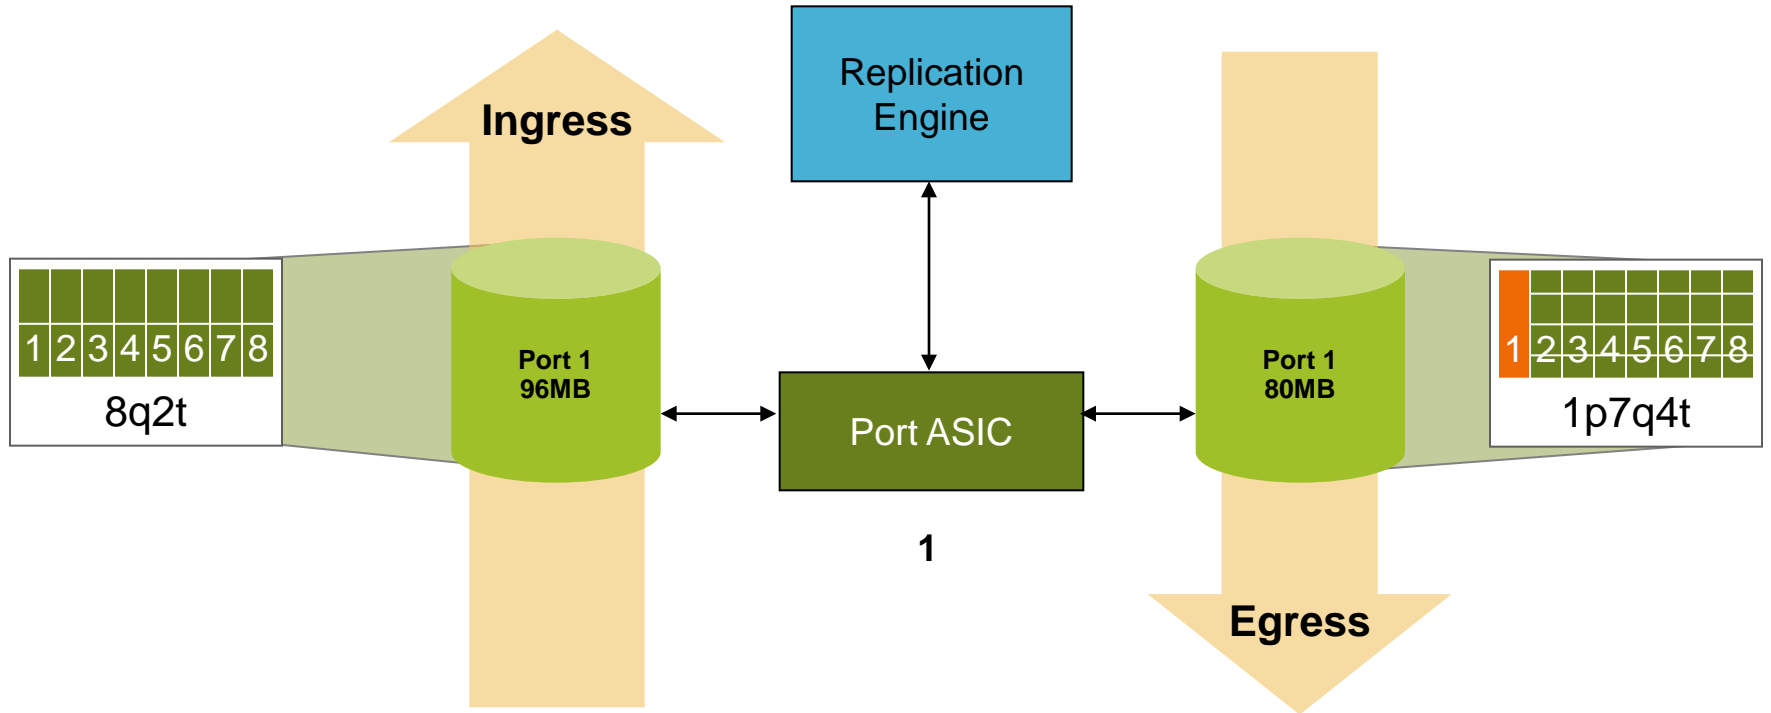


Port QoS – 8-Port 10G Module

- Buffers
 - 96MB ingress per port
 - 80MB egress per port
- Queue Structure
 - 8q2t ingress
 - 1p7q4t egress



8-Port 10G Module Buffering



Port QoS – 32-Port 10G Module

Buffers

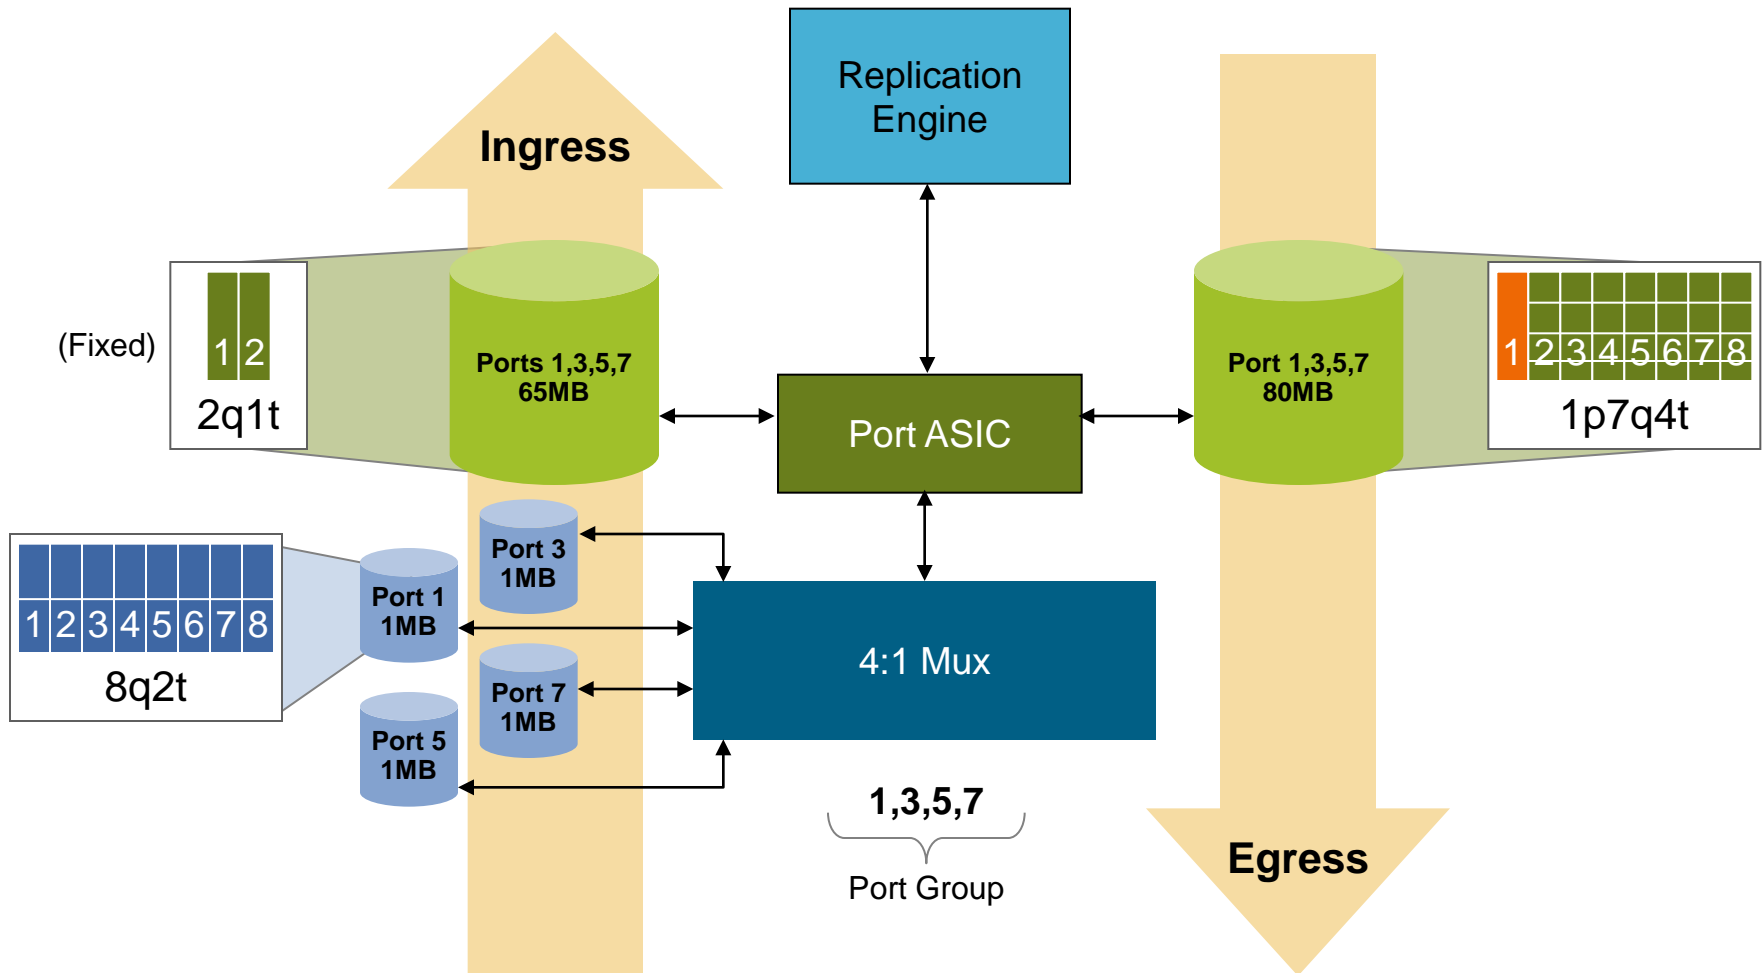
- Ingress (two-stage ingress buffering)
 - Dedicated mode: 1MB per port + 65MB per port
 - Shared mode: 1MB per port + 65MB per port group
- Egress
 - Dedicated mode: 80MB per port
 - Shared mode: 80MB per port-group

Queue Structure

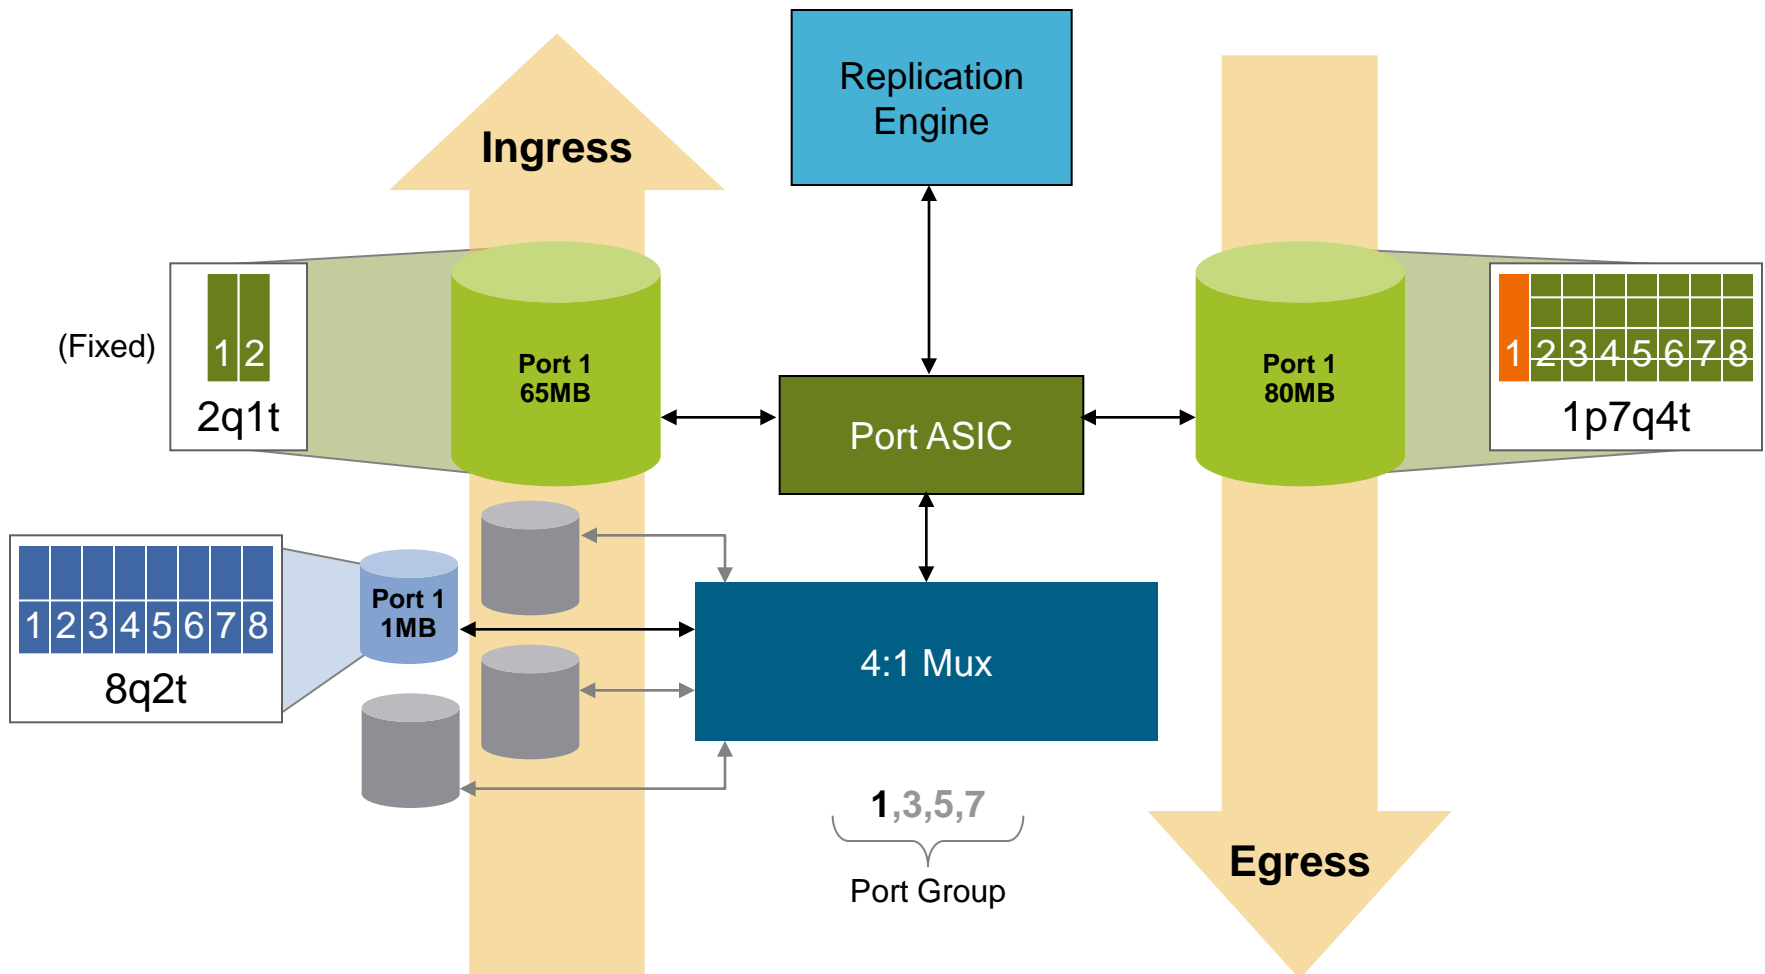
- 8q2t + 2q1t ingress
- 1p7q4t egress



32-Port 10G Module Buffering – Shared Mode



32-Port 10G Module Buffering – Dedicated Mode




Port QoS – 48-Port 1G Modules

Buffers

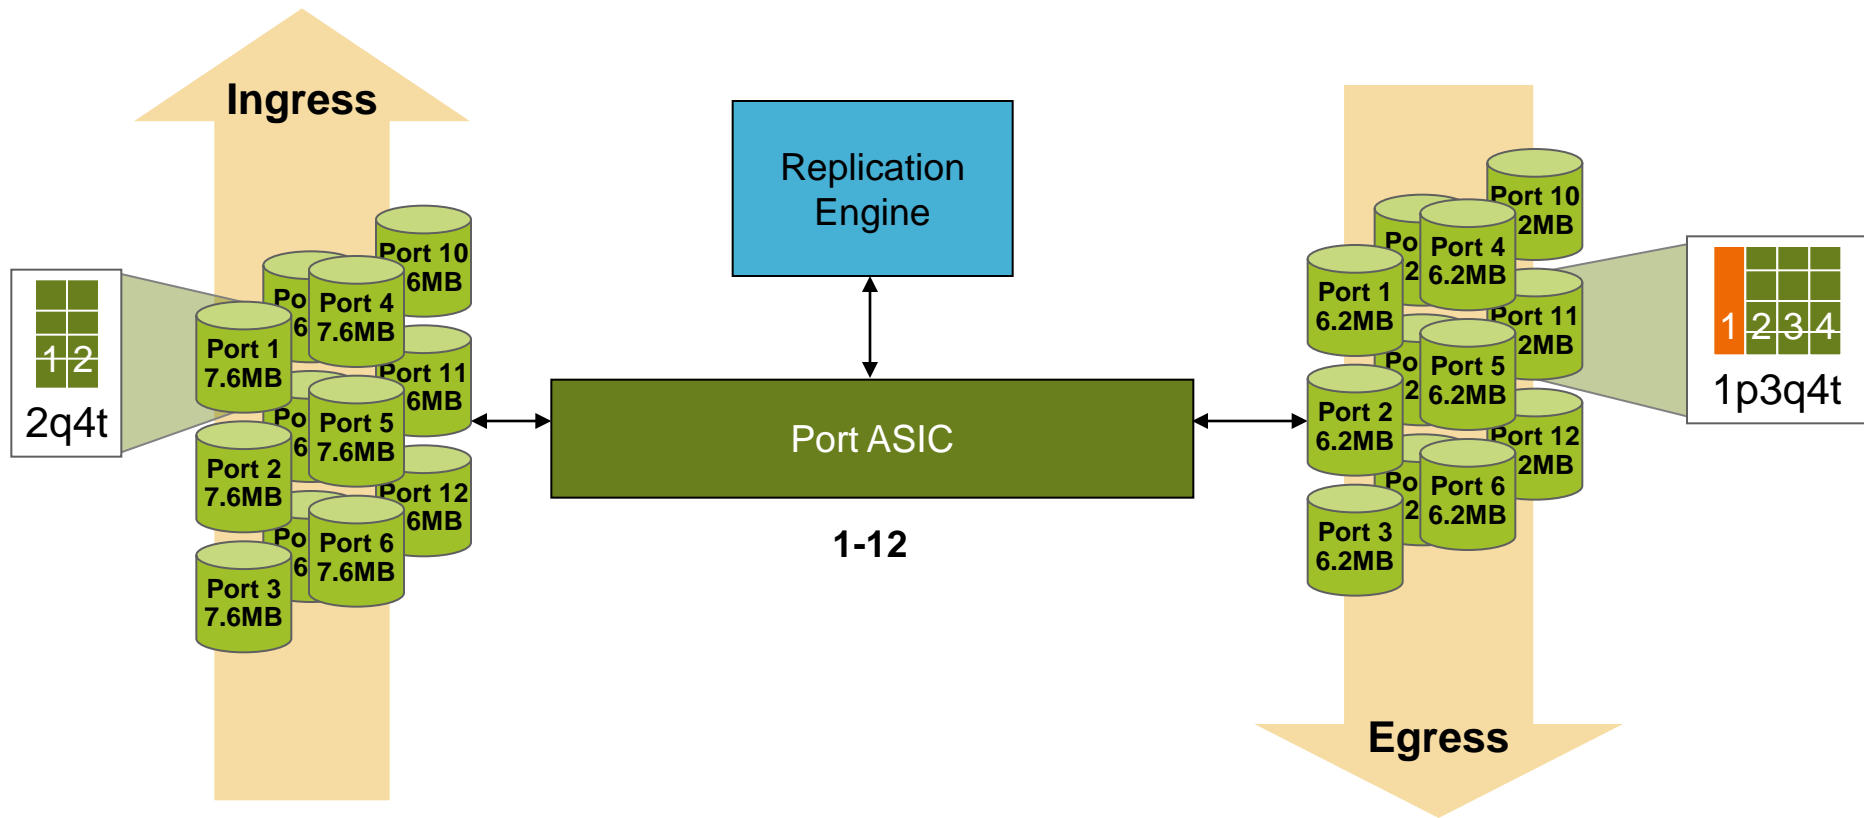
- 7.56MB ingress per port
- 6.15MB egress per port

Queue Structure

- 2q4t ingress
- 1p3q4t egress



48-Port 1G Modules Buffering



Marking and Policing

- Classification uses CL TCAM in forwarding engine to match traffic
- After classification, traffic can be marked or policed
- Marking policies statically set QoS values for each class
- Policing performs markdown and/or policing (drop)
- Policers use classic token-bucket scheme
 - Uses Layer 2 frame size when determining rate
- Note: policing performed on per-forwarding engine basis
 - Shared interfaces (such as SVI/EtherChannel) and egress policies could be policed at $\text{<policing rate> * <number of forwarding engines>}$

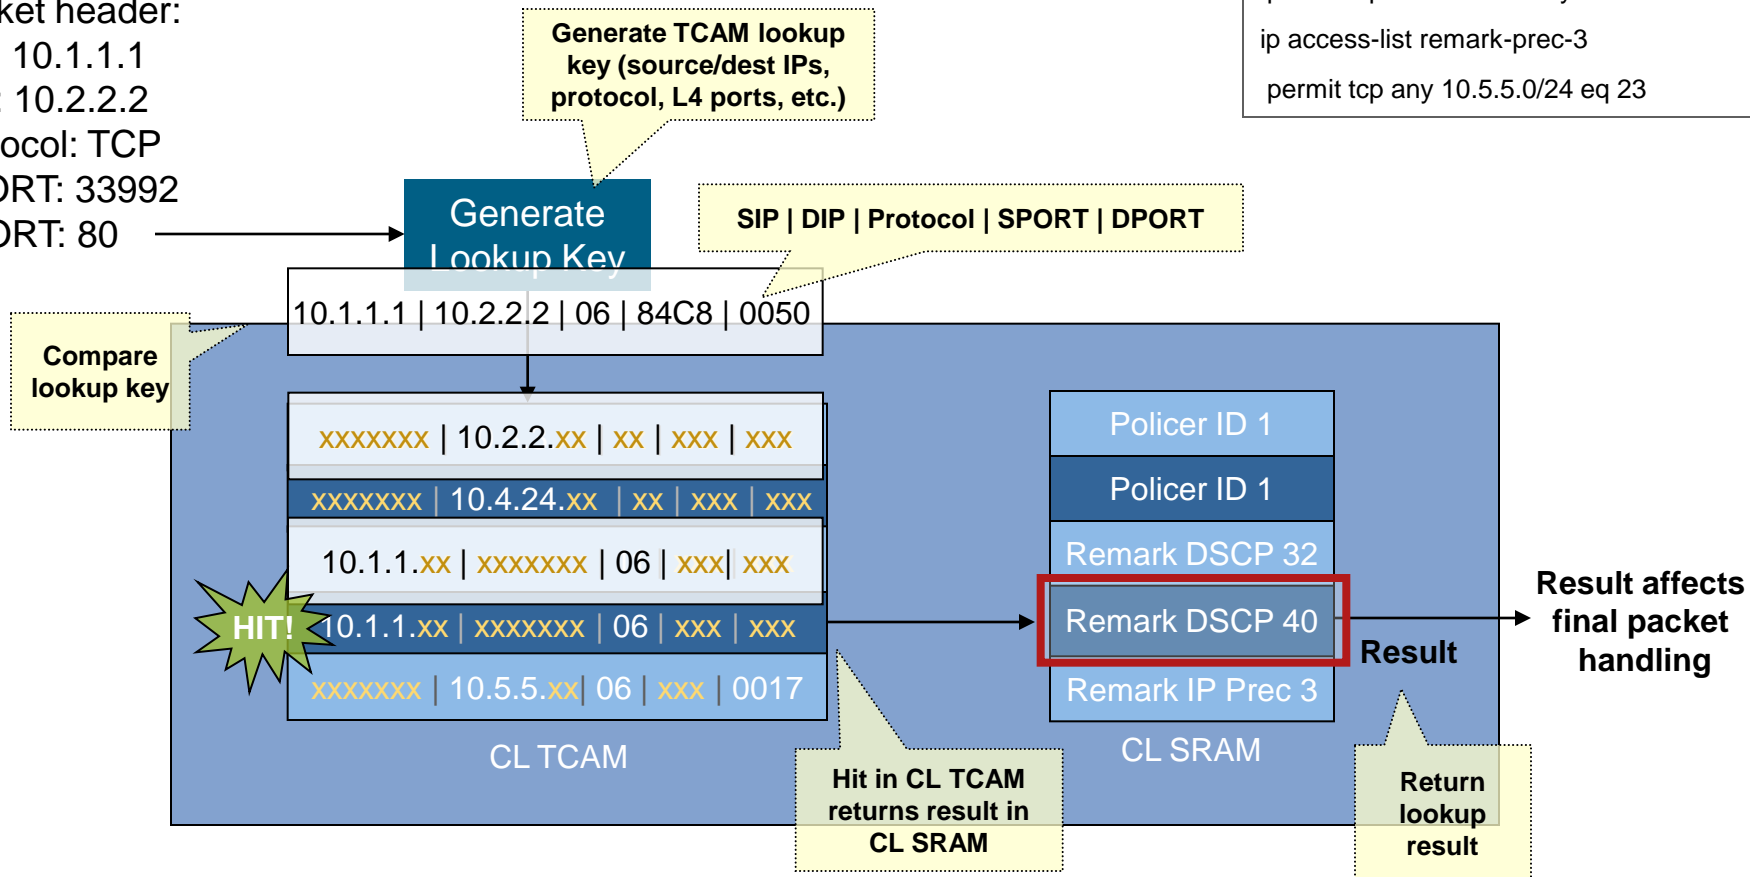


QoS CL TCAM Lookup

Packet header:
SIP: 10.1.1.1
DIP: 10.2.2.2
Protocol: TCP
SPORT: 33992
DPORT: 80

QoS Classification ACLs

```
ip access-list police
permit ip any 10.3.3.0/24
permit ip any 10.4.12.0/24
ip access-list remark-dscp-32
permit udp 10.1.1.0/24 any
ip access-list remark-dscp-40
permit tcp 10.1.1.0/24 any
ip access-list remark-prec-3
permit tcp any 10.5.5.0/24 eq 23
```



Monitoring QoS Service Policies

```
show policy-map interface [[<interface>] [type  
qos|queuing]]|brief]
```

```
n7010# show policy-map interface e9/1
```

```
Global statistics status :    enabled
```

```
Ethernet9/1
```

```
Service-policy (qos) input:    mark  
policy statistics status:    enabled
```

```
Class-map (qos):    udp-mcast (match-all)  
432117468 packets  
Match: access-group multicast  
set dscp cs4
```

```
Class-map (qos):    udp (match-all)  
76035663 packets  
Match: access-group other-udp  
police cir 2 mbps bc 1000 bytes pir 4 mbps be 1000 bytes  
conformed 587624064 bytes, 3999632 bps action: transmit  
exceeded 293811456 bytes, 1999812 bps action: set dscp dscp table cir-markdown-map  
violated 22511172352 bytes, 153221133 bps action: drop
```

```
n7010#
```

Agenda

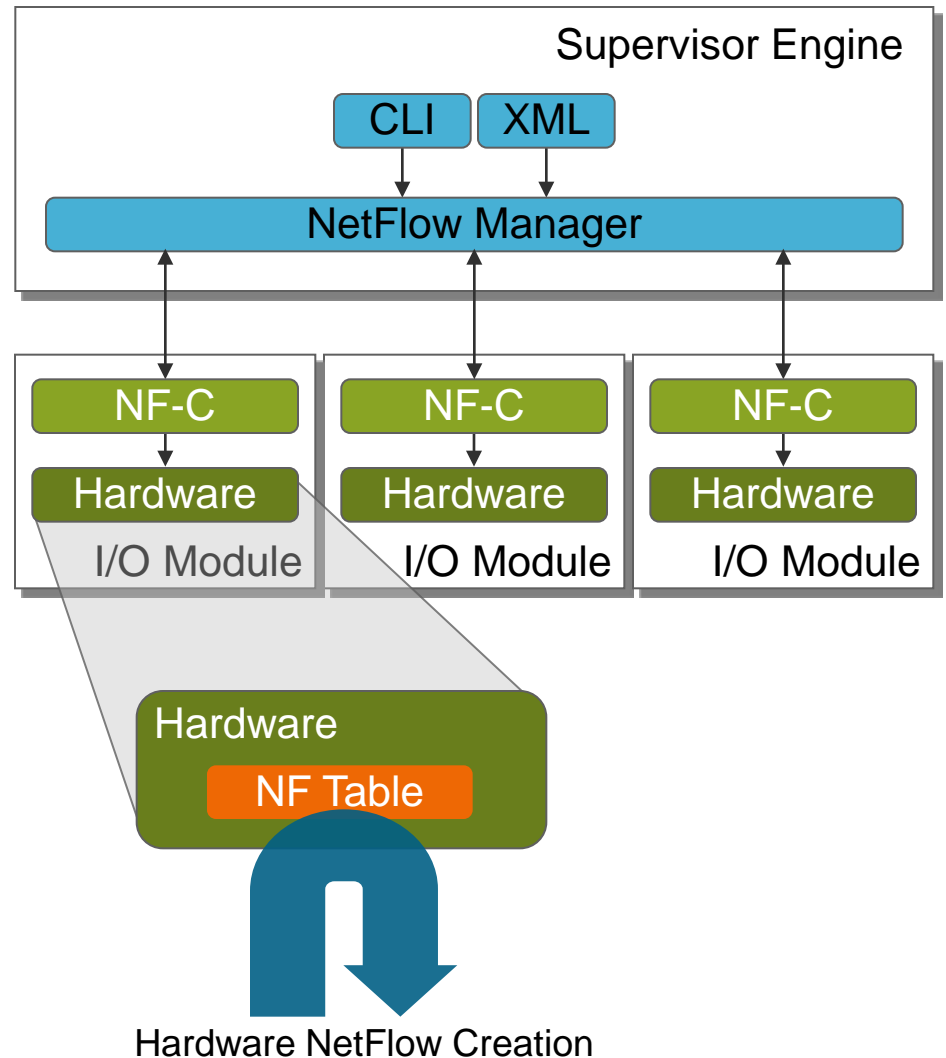
- Chassis Architecture
- Supervisor Engine Architecture
- I/O Module Architecture
- Forwarding Engine Architecture
- Fabric Architecture
- Layer 2 Forwarding
- IP Forwarding
- IP Multicast Forwarding
- ACLs
- QoS
- **NetFlow**

NetFlow

- NetFlow table is 512K entries (490K effective), shared between ingress/egress NetFlow
- Hardware NetFlow entry creation
 - CPU not involved in NetFlow entry creation/update
- All modules have independent NetFlow table
- Full and sampled NetFlow supported by hardware

NetFlow Architecture

- NetFlow manager receives configuration via CLI/XML
- NetFlow manager distributes configuration to NetFlow-Clients on I/O modules
- NetFlow-Clients apply policy to hardware



```
n7010# sh processes cpu | egrep nfm|PID
PID    Runtime(ms)   Invoked    uSecs   lSec    Process
24016          1463    735183570      0      0     nfm
```

```
module-9# sh processes cpu | egrep nfp
1538          68842    424290      162     0.0    nfp
module-9#
```

Full vs. Sampled NetFlow

- NetFlow configured per-direction and per-interface
Ingress and/or egress on per-interface basis
- Each interface can collect full or sampled flow data
- Full NetFlow: Accounts for every packet of every flow on interface, up to capacity of NetFlow table
- Sampled NetFlow: Accounts for M in N packets on interface, up to capacity of NetFlow table

Viewing NetFlow Records

show hardware flow ip [detail] module <mod>

n7010# **sh hardware flow ip interface e9/1 module 9**

D - Direction; IF - Intf/VLAN; L4 Info - Protocol:Source Port:Destination Port
TCP Flags: Ack, Flush, Push, Reset, Syn, Urgent

D	IF	SrcAddr	DstAddr	L4 Info	PktCnt	TCP Flags
I	9/1	010.001.001.002	010.001.002.002	006:01024:01024	0001403880	A . . . S .
I	9/1	010.001.001.003	010.001.002.003	006:01024:01024	0001403880	A . . . S .
I	9/1	010.001.001.004	010.001.002.004	006:01024:01024	0001403880 S .

<...>

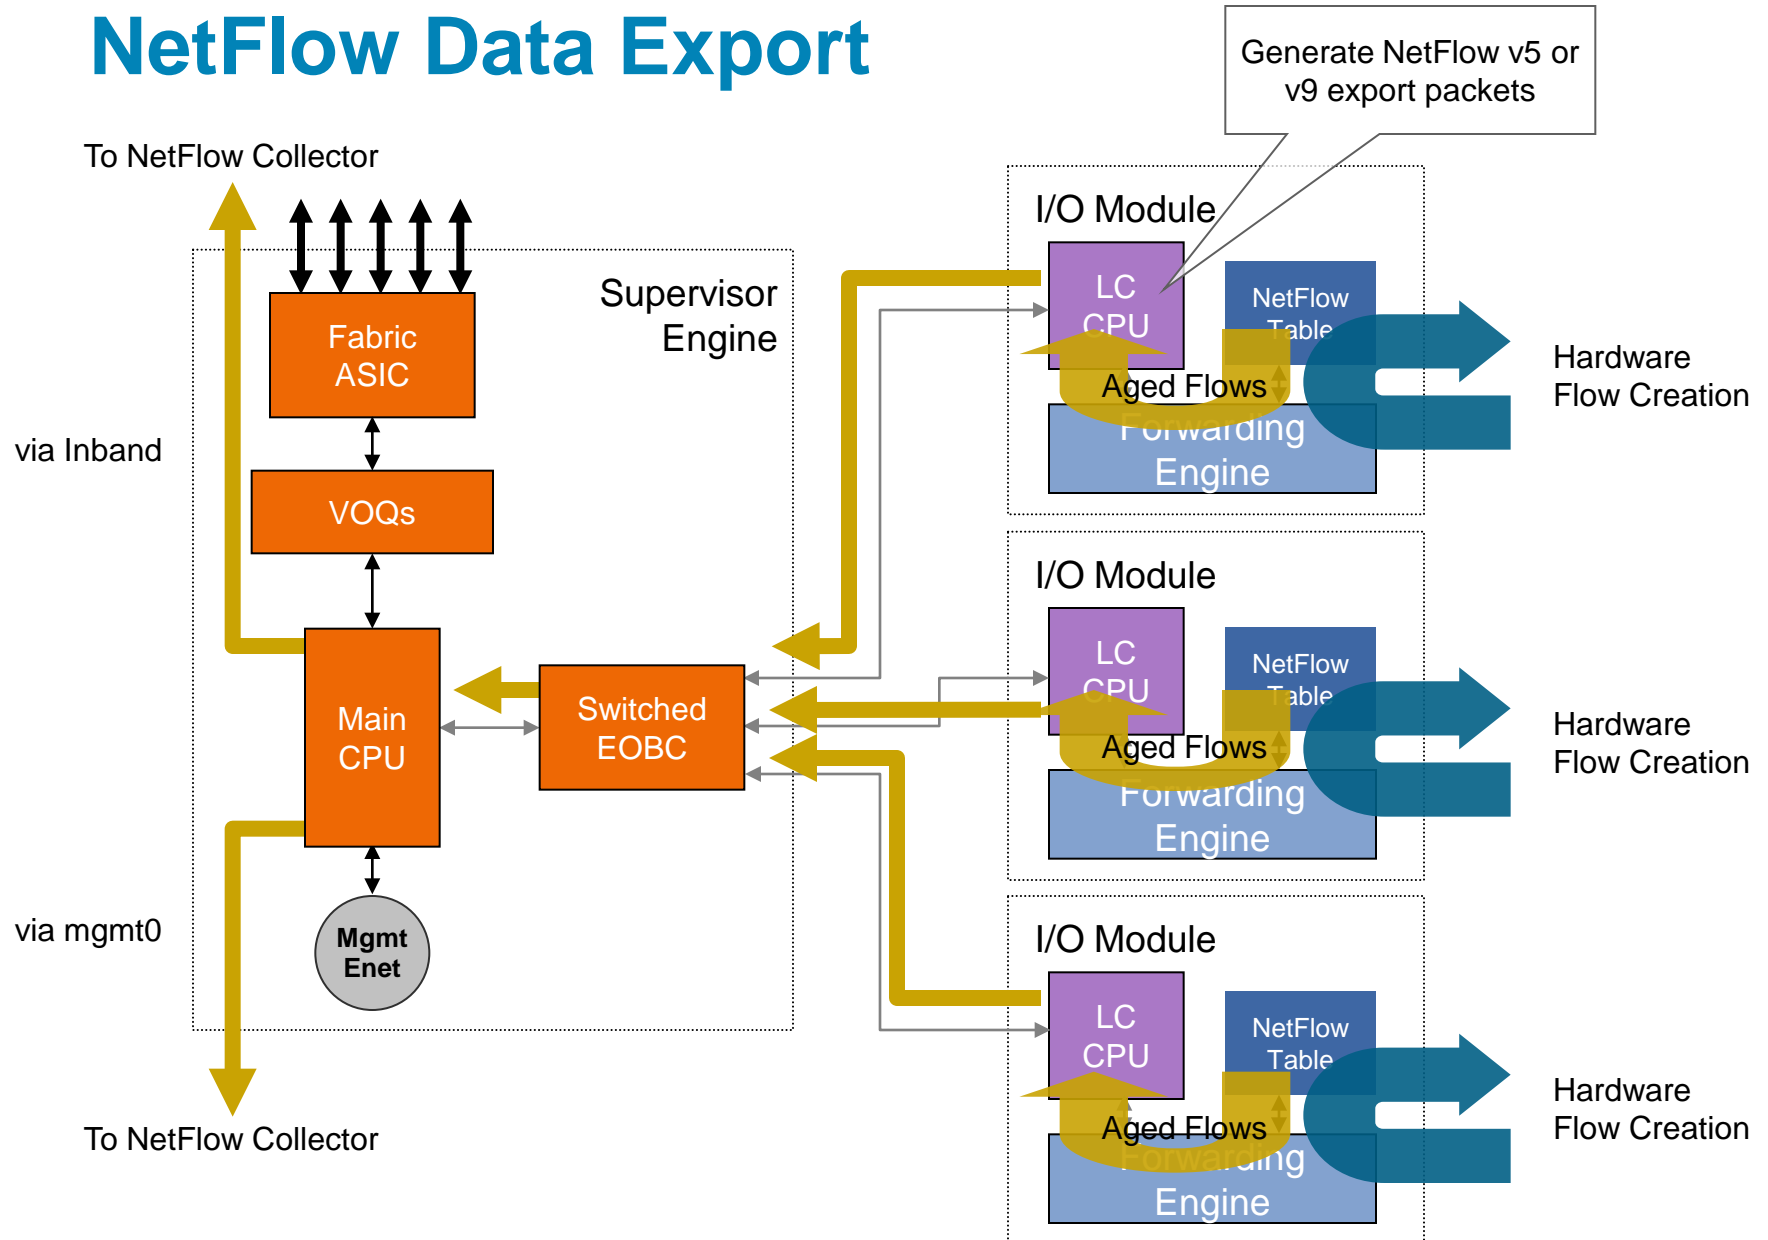
n7010# **sh hardware flow ip interface e9/1 detail module 9**

D - Direction; IF - Intf/VLAN; L4 Info - Protocol:Source Port:Destination Port
TCP Flags: Ack, Flush, Push, Reset, Syn, Urgent; FR - FRagment; FA - FastAging
SID - Sampler/Policer ID; AP - Adjacency/RIT Pointer
CRT - Creation Time; LUT - Last Used Time; NtAddr - NT Table Address

D	IF	SrcAddr	DstAddr	L4 Info	PktCnt	TCP Flags
I	9/1	010.001.001.002	010.001.002.002	006:01024:01024	0001706722	A . . . S .


ByteCnt	TOS	FR	FA	SID	AP	CRT	LUT	NtAddr
0000218460416	000	N	Y	0x000	0x000000	02168	02571	0x000331

NetFlow Data Export



Conclusion

- You should now have a thorough understanding of the Nexus 7000 switching architecture, I/O module design, packet flows, and key forwarding engine functions...
- **Any questions?**



Q and A



Complete Your Online Session Evaluation

- Give us your feedback and you could win fabulous prizes. Winners announced daily.
- Receive 20 Passport points for each session evaluation you complete.
- Complete your session evaluation online now (open a browser through our wireless network to access our portal) or visit one of the Internet stations throughout the Convention Center.



Don't forget to activate your Cisco Live Virtual account for access to all session material, communities, and on-demand and live activities throughout the year. Activate your account at the Cisco booth in the World of Solutions or visit www.ciscolive.com.



CISCO